

EECS 370 - Lecture 6

Function Calls



Live Poll + Q&A: [slido.com #eeecs370](https://slido.com/#eeecs370)

Poll and Q&A Link

Announcements

- Project 1a due tonight
- Project 1s+m due next Thursday
- HW 1 Posted
 - Due Monday 9/22
- Let us know about exam conflicts in the **next week**
 - Form on Ed



Datatype	size (bytes)
char	1
short	2
int	4
double	8

Reminder: Memory Alignment

```

short   a[100];
char    b;
int     c;
double  d;
short   e;
struct {
    char  f;
    int   g[1];
    char  h;
} i;

```

- *Problem:* Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed

$$a = 2 \text{ bytes} * 100 = 200$$

$$b = 1 \text{ byte}$$

$$c = 4 \text{ bytes}$$

$$d = 8 \text{ bytes}$$

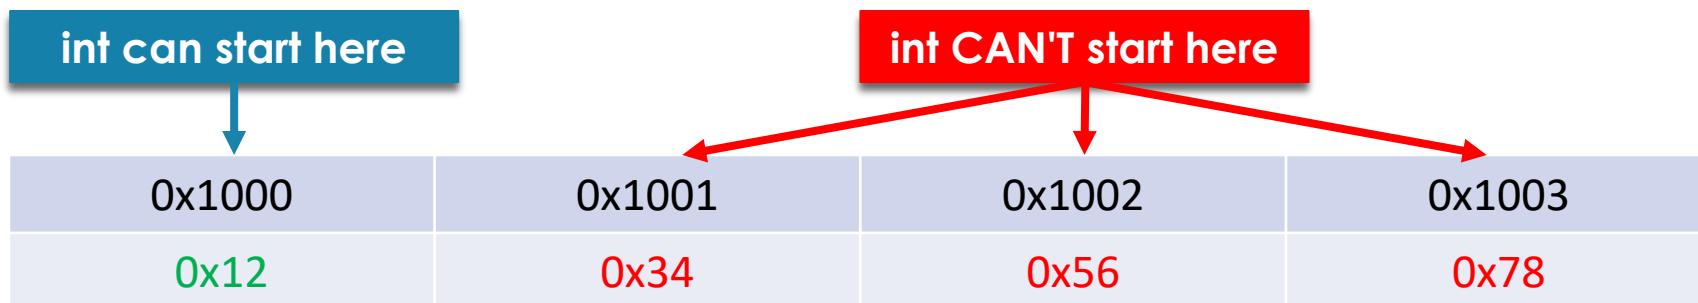
$$e = 2 \text{ bytes}$$

$$i = 1 + 4 + 1 = 6 \text{ bytes}$$

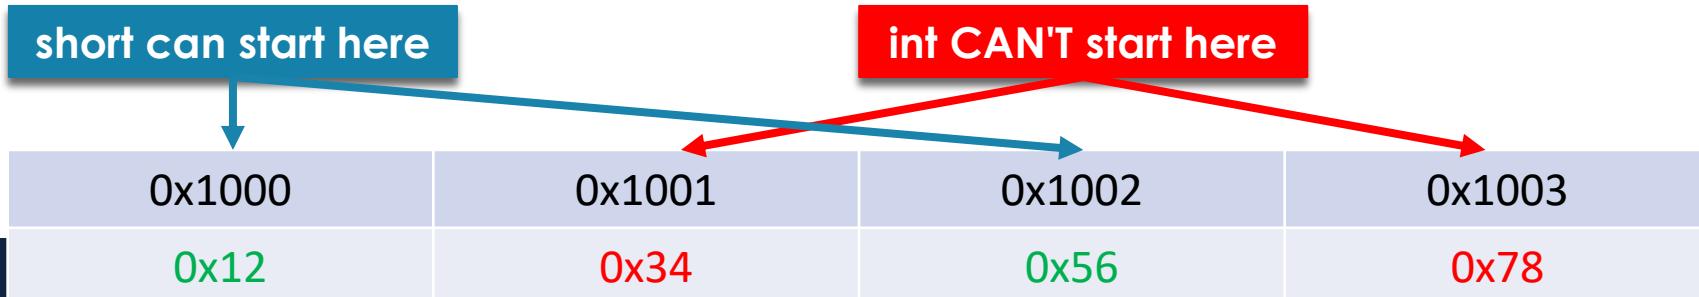
total = 221, right or wrong?

Reminder: Memory Alignment

- Most modern ISAs require that data be aligned
 - An N-byte variable must start at an address A, such that $(A\%N) == 0$
- For example, starting address of a 32 bit **int** must be divisible by 4



- Starting address of a 16 bit **short** must be divisible by 2



Reminder: Golden Rule of Alignment

- Every (primitive) object starts at an address divisible by its size
- "Padding" is placed in between objects if needed

```
char    c;  
short   s;  
int    i;
```

0x1000	0x1001	0x1002	0x1003	0x1004	0x1005	0x1006	0x1007
[c]	[padding]		[s]				[i]

- But what about non-primitive data types?
 - Arrays? Treat as independent objects
 - Structs? Trickier...

Structure Alignment

- In addition, for structs...
 - Identify largest (potentially nested) **primitive** component
 - Starting address of overall struct is aligned based on the largest component
 - Padded in the back so total size is a multiple of the largest component

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	[pad]	[pad]	[pad]	s[0].c	[pad]	[pad]	[pad]		s[0].i		s[1.c]	[pad]	[pad]	[pad]	

Guaranteed to lay
out each instance
identically



Structure Alignment

- Solution: in addition to laying out each field according to Golden Rule...
 - Identify largest (primitive) field
 - Starting address of overall struct is aligned based on the largest field
 - Padded in the back so total size is a multiple of the largest primitive

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	[pad]	[pad]	[pad]	s[0].c	[pad]	[pad]	[pad]		s[0].i		s[1.c]	[pad]	[pad]	[pad]	

Guaranteed to lay
out each instance
identically



Structure Example

```
struct {  
    char w;  
    int x[3];  
    char y;  
    short z;  
}
```

Poll: What boundary should this struct be aligned to?

- a) 1 byte
- b) 4 bytes
- c) 12 bytes
- d) 2 bytes
- e) 19 bytes

- Assume struct starts at location 1000,
 - char w → 1000
 - x[0] → 1004–1007, x[1] → 1008 – 1011, x[2] → 1012 – 1015
 - char y → 1016
 - short z → 1018 – 1019

Total size = 20 bytes!

Datatype	size (bytes)
char	1
short	2
int	4
double	8

Calculating Load/Store Addresses for Variables

```

short a[100];
char b;
int c;
double d;
short e;
struct {
    char f;
    int g[1];
    char h;
} i;

```

- *Problem:* Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed

a = 200 bytes (100-299)

b = 1 byte (300-300)

c = 4 bytes (304-307)

d = 8 bytes (312-319)

e = 2 bytes (320-321)

struct: largest field is 4 bytes, start at 324

f = 1 byte (324-324)

g = 4 bytes (328-331)

h = 1 byte (332-332)

i = 12 bytes (324-335)

236 bytes total!! (compared to 221, originally)

Data Layout – Why?

- Does gcc (or another compiler) reorder variables in memory to avoid padding?
- Only outside structs
- C99 forbids reordering elements inside a struct
- The programmer (i.e., you) are expected to manage data layout of variables for your program and structs.
- Two optimal strategies:
 - Order fields in struct by datatype size, smallest first
 - Or by largest first



Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- **Lecture 6 : Converting C to assembly – functions**
- Lecture 7 : Translation software; libraries, memory layout



LEGv8 Conditional Instructions

- Two varieties of conditional branches
 1. One type compares a register to see if it is equal to zero.
 2. Another type checks the condition codes set in the status register.

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Let's look at the first type: CBZ and CBNZ
 - CBZ: Conditional Branch if Zero
 - CBNZ: Conditional Branch if Not Zero

LEGv8 Conditional Instructions

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
 - The relative address is a 19 bit signed integer representing the number of instructions (not bytes) to branch. Recall instructions are 4 bytes

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Example: CBNZ X3, Again
 - If X3 doesn't equal 0, then branch to label "Again"
 - "Again" is an offset from the PC of the current instruction (CBNZ)
 - Why does "25" in the above table result in PC + 100?!

LEGv8 Conditional Instructions

- Motivation:
 - Some types of branches makes sense to check if a certain value is zero or not
 - `while(a)`
 - But not all:
 - `if(a > b)`
 - `if(a == b)`
 - Using an extra **program status register** to check for various conditions allows for a greater breadth of branching behavior



LEGv8 Conditional Instructions Using FLAGS

- FLAGS: NZVC record the results of (arithmetic) operations Negative, Zero, oVerflow, Carry—not present in LC2K
- We explicitly set them using the “set” modification to ADD/SUB etc.
- Example: ADDS causes the 4 flag bits to be set according as the outcome is negative, zero, overflows, or generates a carry

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	$X1 = X2 + X3$	Three register operands
	subtract	SUB X1, X2, X3	$X1 = X2 - X3$	Three register operands
	add immediate	ADDI X1, X2, 20	$X1 = X2 + 20$	Used to add constants
	subtract immediate	SUBI X1, X2, 20	$X1 = X2 - 20$	Used to subtract constants
	add and set flags	ADDS X1, X2, X3	$X1 = X2 + X3$	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	$X1 = X2 - X3$	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	$X1 = X2 + 20$	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	$X1 = X2 - 20$	Subtract constant, set condition codes



ARM Condition Codes Determine Direction of Branch--continued

Encoding	Name (& alias)	Meaning (integer)	Flags
0000	EQ	Equal	$Z==1$
0001	NE	Not equal	$Z==0$
0010	HS (CS)	Unsigned higher or same (Carry set)	$C==1$
0011	LO (CC)	Unsigned lower (Carry clear)	$C==0$
0100	MI	Minus (negative)	$N==1$
0101	PL	Plus (positive or zero)	$N==0$
0110	VS	Overflow set	$V==1$
0111	VC	Overflow clear	$V==0$
1000	HI	Unsigned higher	$C==1 \&& Z==0$
1001	LS	Unsigned lower or same	$! (C==1 \&& Z==0)$
1010	GE	Signed greater than or equal	$N==V$
1011	LT	Signed less than	$N!=V$
1100	GT	Signed greater than	$Z==0 \&& N==V$
1101	LE	Signed less than or equal	$! (Z==0 \&& N==V)$
1110	AL	Always	Any
1111	NV [†]		

Need to know the 7 with the red arrows

```
CMP X1, X2  
B.LE Label1
```

For this example,
we branch if X1 is
 \geq to X2

Conditional Branches: How to use

- CMP instruction lets you compare two registers.
 - Could also use SUBS etc.
 - That could save you an instruction.
- B.cond lets you branch based on that comparison.
- Example:

```
CMP    X1,  X2
B.GT  Label1
```

- Branches to Label1 if X1 is greater than X2.



Branch—Example

- Here's how we can convert an if/else statement to ARM (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

Using Labels

```
CMP X1, X2  
B.NE L1  
ADD X1, X1, #1  
B L2  
L1: ADD X2, X2, #1  
L2: ...
```

Note that conditions in assembly are often the inverse of the "if" condition. Why?

Without Labels

CMP	X1, X2
B.NE	3
ADD	X1, X1, #1
B	2
ADD	X2, X2, #1

Assemblers must deal with labels and assign displacements

Loop—Example

```
// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 800, sum is in X2
```

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

of branch instructions
= $3 \times 10 + 1 = 31$

a.k.a. while-do template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	CMPI	X1, #10
	B.EQ	endLoop
	LSL	X6, X1, #3
	LDUR	X5, [X6, #800]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endif:	ADDI	X1, X1, #1
	B	Loop1
endLoop:		

Agenda

- Using branches more generally
- **Function calls and the call stack**
- Assigning variables to memory locations
- Saving registers
- Caller/callee example



Implementing Functions

- Does this assembly code do what we need?

```
int mult_2(int x){  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main(){  
    int result = mult_2(GLOBAL+1);  
    other(result);  
}
```

**Poll: What's wrong
with this approach?**

	LDURSW X1, [XZR, GLOBAL] ADD X2, X1, #1 // Inc GLOBAL STURW X2, [XZR, X] // Pass arg B MULT_2 // Execute func
RETURN:	LDURSW X3, [XZR, TEMP] // load result STURW X3, [XZR, STRING] // Pass arg B OTHER // Execute func
...	
MULT_2:	LDURSW X1, [XZR, X] // load arg ADD X2, X1, X1 // mult by 2 STURW X2, [XZR, TEMP] // return result B RETURN // return

Problem 1: Returning from Functions

- Branches so far have hard-coded destination

```
B.NE L1  
ADD X1, X1, #1  
B L2  
L1: ADD X2, X2, #1  
L2: ...
```

```
B.NE      3  
ADD      X1, X1, #1  
B        2  
ADD      X2, X2, #1
```

```
int func(int x) {  
    printf(x * 10);  
    return;  
}  
int helper() {  
    func(7);  
}  
int main() {  
    helper();  
    func(13);  
}
```

Should this return to
"helper" or "main"?

- This is fine for if-statements, for-loops etc
- But functions can be called from multiple places
 - Meaning we'll return to different spots on each func call! Can't hardcode offset!

Solution: Indirect Jumps

- Indirect branches or "jumps" don't hardcode destination in instruction
- They index a register whose value holds destination

Unconditional branch	branch	B 2500	go to PC + 10000	Branch to target address; PC-relative
	branch to register	BR X30	go to X30	For switch, procedure return
	branch with link	BL 2500	X30 = PC + 4; PC + 10000	For procedure call PC-relative

- Use "**BL**" to call a function
 - Destination is hardcoded
 - PC +4 (return address) stored in X30
- Use "**BR**" to return from a function
 - X30 is read for return address
 - Allows us to return to different places



Solution: Indirect Jumps

```
int mult_2(int x){  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main(){  
    int result = mult_2(GLOBAL+1);  
    other(result);  
}
```

Also don't
need "return"
labels

```
LDURSW X1, [XZR, GLOBAL]  
ADD    X2, X1, #1           // Inc GLOBAL  
STURW  X2, [XZR, X]        // Pass arg  
BL     MULT_2              // Execute func  
  
RETURN:  
LDURSW X3, [XZR, TEMP]    // load result  
STURW  X3, [XZR, STRING]  // Pass arg  
BL     OTHER                // Execute func  
...  
MULT_2:  
LDURSW X1, [XZR, X]        // load arg  
ADD    X2, X1, X1          // mult by 2  
STURW  X2, [XZR, TEMP]    // return result  
BR                            // return
```

Now MULT_2
can return to
whatever
function
called it

Problem 2: Passing Parameters

For recursive functions,
global variables could be
overwritten

```
int mult_2(int x) {  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main() {  
    int result = mult_2(GLOBAL+1);  
    other(result);  
}
```

```
LDURSW X1, [XZR, GLOBAL]  
ADD X2, X1, #1 // Inc GLOBAL  
STURW X2, [XZR, X] // Pass arg  
BL MULT_2 // Execute func  
LDURSW X3, [XZR, TEMP] // load result  
STURW X3, [XZR, STRING] // Pass arg  
BL OTHER // Execute func  
...  
MULT_2:  
    LDURSW X1, [XZR, X] // load arg  
    ADD X2, X1, X1 // mult by 2  
    STURW X2, [XZR, TEMP] // return result  
    BR // return
```

How to pass in
parameters?

Task 1: Passing parameters

- Where should you put all of the parameters?
 - Registers?
 - Fast access but few in number and wrong size for some objects
 - Memory?
 - Where to put it? And, memory is slow
- ARMv8 solution—and the usual answer:
 - Both
 - Put the first few parameters in registers (if they fit) (X0 – X7)
 - Put the rest in memory on the call stack— **important concept**

Call stack

- ARM conventions (and most other architectures) allocate a region of memory for the “call” stack
 - This memory is used to manage all the storage requirements to simulate function call semantics
 - Parameters (that were not passed through registers)
 - Local variables
 - Temporary storage (when you run out of registers and need somewhere to save a value)
 - Return address
 - Etc.
- Sections of memory on the call stack [**stack frames**] are allocated when you make a function call, and de-allocated when you return from a function
- Grows “down” to lower addresses, usually (but not for LC2K projects)

The stack grows as functions are called

FUNCTION CALLS

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

inside foo

foo's stack frame

bar calls printf

foo's stack frame

bar's stack frame

printf's stack frame

foo calls bar

foo's stack frame

bar's stack frame

The stack shrinks as functions return

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

printf returns

foo's stack frame

bar's stack frame

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

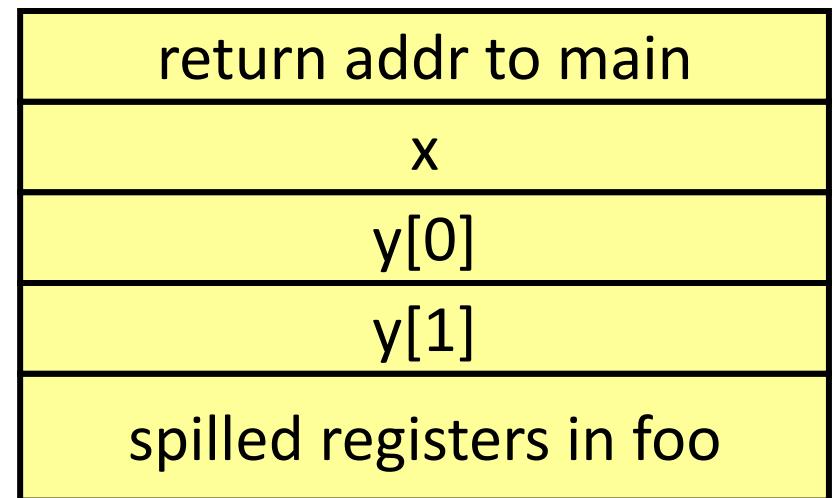
bar returns

foo's stack frame

Stack frame contents

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}  
  
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

foo's stack frame



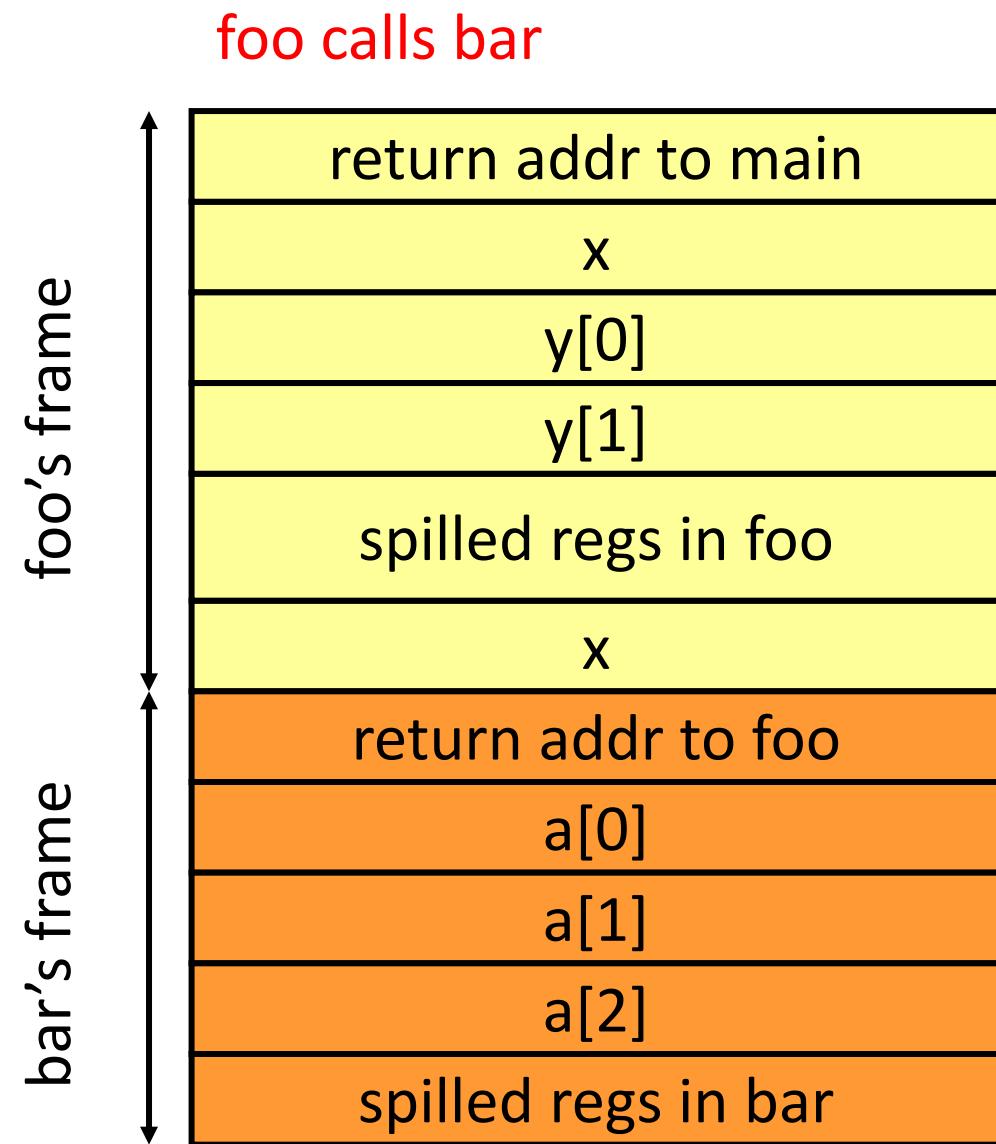
Stack frame contents (2)

FUNCTION CALLS

```
void foo()
{
    int x, y[2];
    bar(x);
}

void bar(int x)
{
    int a[3];
    printf();
}
```

Spill data—not enough room in x0-x7 for params and also caller and callee saves



Agenda

- Using branches more generally
- Function calls and the call stack
- **Assigning variables to memory locations**
- Saving registers
- Caller/callee example

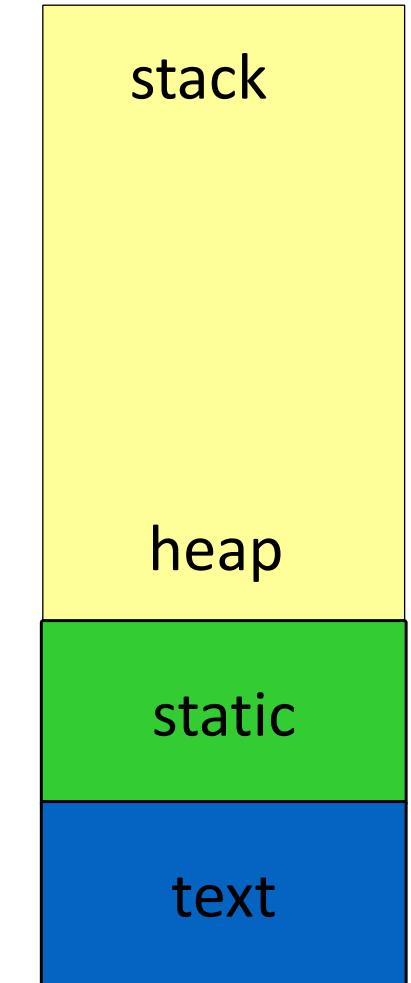


Review: Where do the variables go?

Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
    free(p);  
}
```



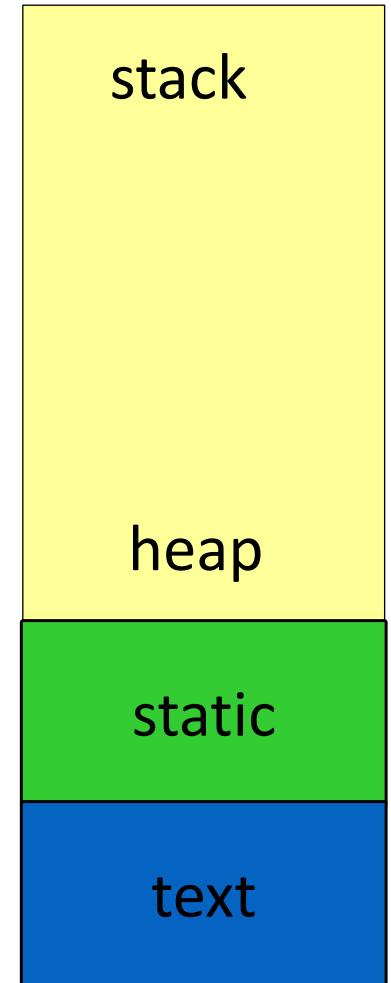
Assigning variables to memory spaces

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
    free(p);  
}
```

w goes in static, as it's a global
x goes on the stack, as it's a parameter

y goes in static, 1 copy of this!!
p goes on the stack
allocate 10 bytes on heap, ptr
set to the address
string literal "%s\n" goes in static,
implicit pointer to string on stack, p goes
on stack

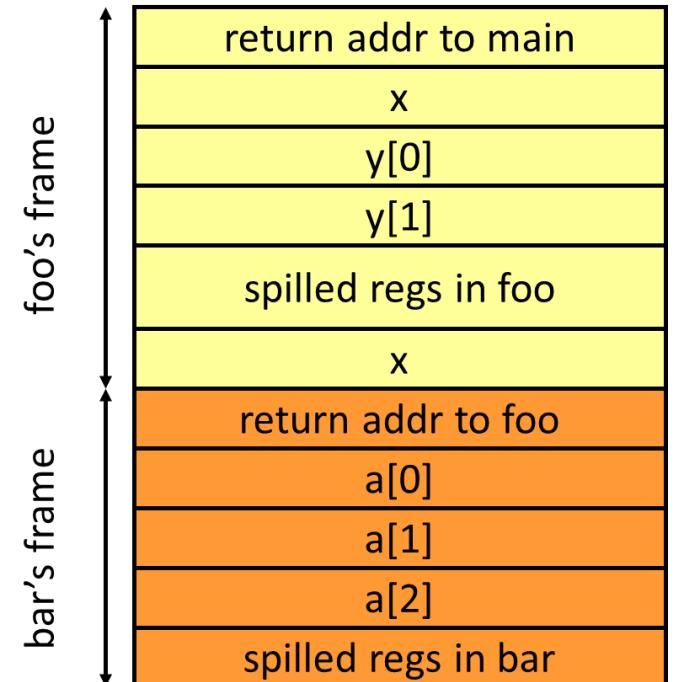
The addresses of local variables
will be different depending on
where we are in the call stack



Accessing Local Variables

- Stack pointer (SP):
 - register intended keeps track of current top of stack
 - Aliases to “X28” in LEGv8
- Compiler (or assembly writer) knows relative offsets of objects in stack
- Can modify SP or use load/store offsets
- **DOESN'T USE LABELS!**

```
sub    sp, sp, #16      // make room on stack
      mov    x0, #42        // put 42 in x0
      stur   x0, [sp, 0]    // store local var on stack
      ldur   x1, [sp, 0]    // load back into x1
```



Agenda

- Using branches more generally
- Function calls and the call stack
- Assigning variables to memory locations
- **Saving registers**
- Caller/callee example



Problem 3: Reusing registers

- Higher level languages (like C/C++) provide many abstractions that don't exist at the assembly level
- E.g. in C, each function has its own local variables
 - Even if different functions have local variables with the same name, they are independent and guaranteed not to interfere with each other!

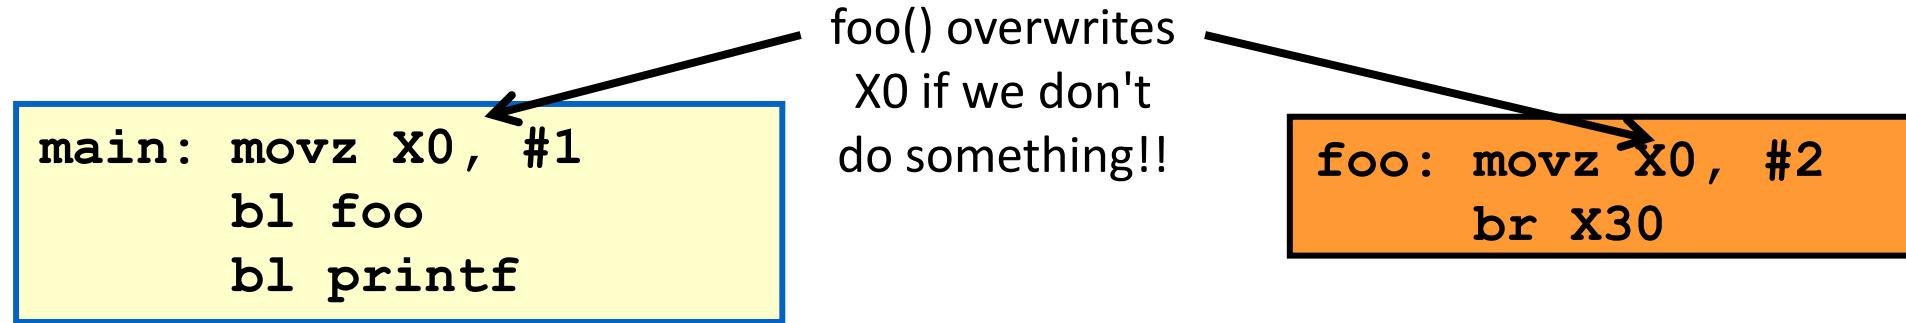
```
void foo() {  
    int a=1;  
    bar();  
    printf(a);  
}
```

Still prints "1"...
these don't
interfere

```
void bar() {  
    int a=2;  
    return;  
}
```

What about registers?

- But in assembly, all functions share a small set (e.g. 32) of registers
 - Called functions will overwrite registers needed by calling functions



- "Someone" needs to save/restore values when a function is called to ensure this doesn't happen

Two Possible Solutions

- Either the **called** function **saves** register values before it overwrites them and **restores** them before the function returns (**callee saved**)...

```
main: movz x0, #1
      bl foo
      bl printf
```

```
foo: stur x0, [stack]
      movz x0, #2
      ldur x0, [stack]
      br x30
```

- Or the **calling** function **saves** register values before the function call and **restores** them after the function call (**caller saved**)...

```
main: movz x0, #1
      stur x0, [stack]
      bl foo
      ldur x0, [stack]
      bl printf
```

```
foo: movz x0, #2
      br x30
```

Next Time

- Finish Up Function Calls
- Talks about linking – the final puzzle piece of software



Extra problems

Extra Example: Do-while Loop

```
// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2
```

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

of branch instructions
= $2 * 10 = 20$

a.k.a. do-while template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endIf:	ADDI	X1, X1, #1
	CMPI	X1, #10
	B.LT	Loop1
endLoop:		

Extra Example: Do-while Loop

```
// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2
```

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

of branch instructions

$$= 2 * 10 = 20$$

a.k.a. do-while template



Extra Problem – For Your Reference

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in x1
// struct {int val; struct node *next;} node;
// struct node *ptr;

if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```

Extra Problem

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in x1
// struct {int val; struct node *next;} node;
// struct node *ptr;

if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```

```
        cmp r1, #0
        beq Endif
        ldursw r2, [r1, #0]
        cmp r2, #0
        b.le Endif
        add r2, r2, #1
        str r2, [r1, #0]
Endif : ....
```

Extra Class Problem

- How much memory is required for the following data, assuming that the data starts at address 200 and is a 32 bit address space?

```
int a;  
struct {double b, char c, int d} e;  
char* f;  
short g[20];
```

Poll: How much memory?

- a) $x < 40$ bytes
- b) $40 < x < 50$ bytes
- c) $50 < x < 60$ bytes
- d) $60 < x$ bytes