

EECS 370 - Lecture 7

Linking



```
jbbeau@JON-PC:~$  
jbbeau@JON-PC:~$  
jbbeau@JON-PC:~$  
jbbeau@JON-PC:~$ gcc integration.c -o integration  
integration.c: In function 'main':  
integration.c:12:5: warning: argument 3 null where non-null expected [-Wnonnull]  
12 |     pthread_create(&t, NULL, NULL, NULL);  
|     ^~~~~~  
In file included from integration.c:2:  
/usr/include/pthread.h:202:12: note: in a call to function 'pthread_create' declared 'nonnull'  
202 | extern int pthread_create (pthread_t *__restrict __newthread,  
|     ^~~~~~  
/usr/bin/ld: /tmp/ccEp5qSZ.o: in function 'main':  
integration.c:(.text+0x65): undefined reference to 'deflateInit_'  
/usr/bin/ld: integration.c:(.text+0x74): undefined reference to 'OPENSSL_init_ssl'  
collect2: error: ld returned 1 exit status  
jbbeau@JON-PC:~$ echo "wut?"
```



Live Poll + Q&A: [slido.com #eeecs370](https://slido.com/#eeecs370)

Poll and Q&A Link

Announcements

- P1
 - Project 1 s + m due Thu
 - Instructor assembler available on the AG
- HW 1
 - Due Monday (9/22)
- Lab 4 meets Fr/M
 - Pre-Lab 4 quiz due Thursday
- Get exam conflicts sent to us **ASAP**
 - Forms listed on Ed
- My OH adjustments:
 - Today's OH: in **2733 BBB** (instead of 2901)
 - Thursday OH cancelled for this week
 - I'm adding Monday and Wednesday OH as well (see Google calendar)



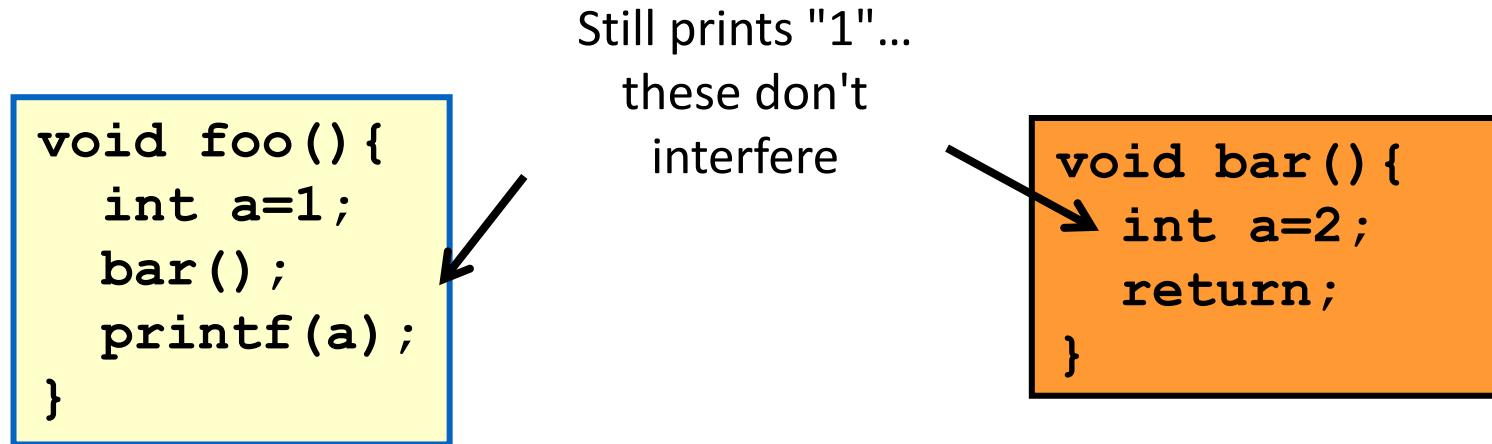
Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- **Lecture 7 : Translation software; libraries, memory layout**



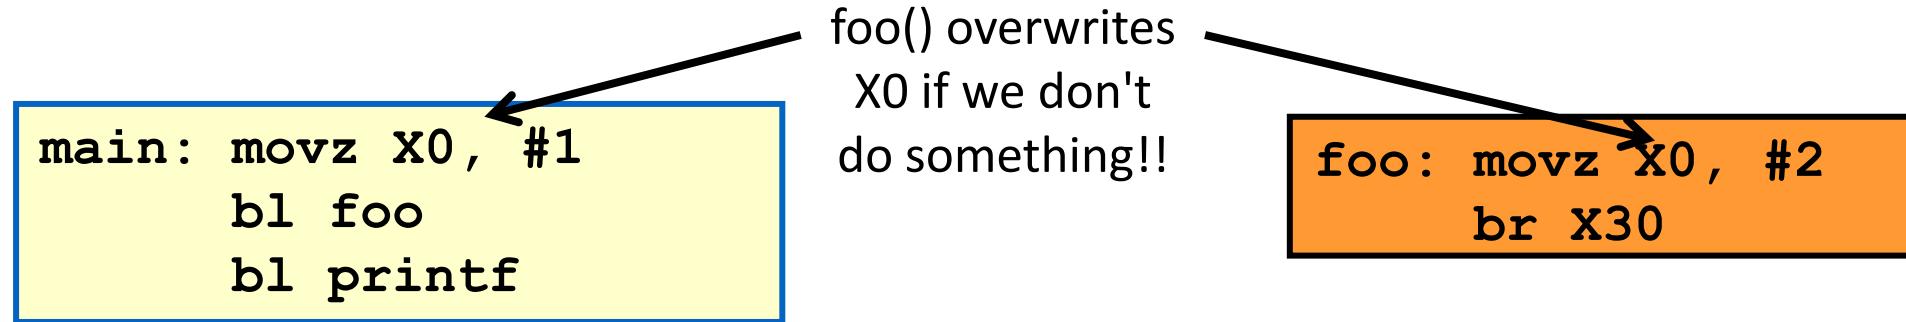
Review - Saving / Restoring Registers

- Higher level languages (like C/C++) provide many abstractions that don't exist at the assembly level
- E.g. in C, each function has its own local variables
 - Even if different functions have local variables with the same name, they are independent and guaranteed not to interfere with each other!



What about registers?

- But in assembly, all functions share a small set (e.g. 32) of registers
 - Called functions will overwrite registers needed by calling functions



- "Someone" needs to save/restore values when a function is called to ensure this doesn't happen

Two Possible Solutions

- Either the **called** function **saves** register values before it overwrites them and **restores** them before the function returns (**callee saved**)...

```
main: movz x0, #1
      bl foo
      bl printf
```

```
foo: stur x0, [stack]
      movz x0, #2
      ldur x0, [stack]
      br x30
```

- Or the **calling** function **saves** register values before the function call and **restores** them after the function call (**caller saved**)...

```
main: movz x0, #1
      stur x0, [stack]
      bl foo
      ldur x0, [stack]
      bl printf
```

```
foo: movz x0, #2
      br x30
```

Another Visualization

Original C
Code

```
void foo(){  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

No need to
save r2/r3.
Why?

Additions for
Caller-save

```
void foo(){  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
save r1 to stack  
save r4 to stack  
bar();  
restore r1  
restore r4  
d = a+d;  
return();  
}
```

Assume bar() will
overwrite all registers

Additions for
Callee-save

```
void foo(){  
    int a,b,c,d;  
save r1  
save r2  
save r3  
save r4  
a = 5; b = 6;  
c = a+1; d=c-1;  
bar();  
d = a+d;  
restore r1  
restore r2  
restore r3  
restore r4  
return();  
}
```

bar() will save a,b, but
now foo() must save
main's variables

Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- **Caller-saved**
 - Only needs saving if value is “**live**” across function call
 - **Live** = contains a useful value: Assign value before function call, use that value after the function call
 - In a leaf function (a function that calls no other function), caller saves can be used without saving/restoring

a, d are live

b, c are NOT
live

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Callee-saved
 - Only needs saving at beginning of function and restoring at end of function
 - Only save/restore it if function overwrites the register

Only use r1-
r4

No need to
save other
registers

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

Caller versus Callee

- Which is better??
- Let's look at some examples...
- Simplifying assumptions:
 - A function can be invoked by many different call sites in different functions.
 - Assume no inter-procedural analysis (hard problem)
 - A function has no knowledge about which registers are used in either its caller or callee
 - Assume main() is not invoked by another function
 - Implication
 - Any register allocation optimization is done using function local information

Caller-saved vs. callee saved – Multiple function case

```
void main() {
    int a,b,c,d;

    c = 5; d = 6;
    a = 2; b = 3;
    foo();
    d = a+b+c+d;

}
```

```
void foo() {
    int e,f;

    e = 2; f = 3;
    bar();
    e = e + f;
```

```
void bar() {
    int g,h,i,j;

    g = 0; h = 1;
    i = 2; j = 3;
    final();
    j = g+h+i;
```

```
void final() {
    int y,z;

    y = 2; z = 3;
    z = y+z;
```

Note: assume main does not have to save any callee registers

Question 1: Caller-save

```
void main() {
    int a,b,c,d;
    c = 5; d = 6;
    a = 2; b = 3;
    [4 STUR]
    foo();
    [4 LDUR]
    d = a+b+c+d;
}
```

```
void foo() {
    int e,f;
    e = 2; f = 3;
    [2 STUR]
    bar();
    [2 LDUR]
    e = e + f;
}
```

```
void bar() {
    int g,h,i,j;
    g = 0; h = 1;
    i = 2; j = 3;
    [3 STUR]
    final();
    [3 LDUR]
    j = g+h+i;
}
```

```
void final() {
    int y,z;
    y = 2; z = 3;
    z = y+z;
}
```

Total: 9 STUR / 9 LDUR

Question 2: Callee-save

Poll: How many Id/st pairs are needed?

```
void main() {
    int a,b,c,d;

    c = 5; d = 6;
    a = 2; b = 3;
    foo();
    d = a+b+c+d;

}
```

```
void foo() {
    [2 STUR]
    int e,f;

    e = 2; f = 3;
    bar();
    e = e + f;

}

[2 LDUR]
```

```
void bar() {
    [4 STUR]
    int g,h,i,j;
    g = 0; h = 1;
    i = 2; j = 3;
    final();
    j = g+h+i;

}

[4 LDUR]
```

```
void final() {
    [2 STUR]
    int y,z;

    y = 2; z = 3;
    z = y+z;

}

[2 LDUR]
```

Total: 8 STUR / 8 LDUR



Is one better?

- **Caller-save** works best when we don't have many live values across function call
- **Callee-save** works best when we don't use many registers overall
- We probably see functions of both kinds across an entire program
- Solution:
 - Use both!
 - E.g. if we have 6 registers, use some (say r0-r2) as **caller-save** and others (say r3-r5) as **callee-save**
 - Now each function can optimize for each situation to reduce saving/restoring
 - Not discussed further for this class

LEGv8 ABI- Application Binary Interface

- The ABI is an agreement about how to use the various registers
- Not enforced by hardware, just a convention by programmers / compilers
- If you want your code to work with other functions / libraries, **follow these**
- Some register conventions in ARMv8
 - X30 is the **link register** – used to hold return address
 - X28 is **stack pointer** – holds address of top of stack
 - X19-X27 are **callee-saved** – function must save these before writing to them
 - X0-15 are **caller-saved** – function must save live values before call
 - X0-X7 used for **arguments** (memory used if more space is needed)
 - X0 used for **return value**

Caller/Callee

- Still not clicking?
- Don't worry, this is a tricky concept for students to get
- Check out supplemental video
 - <https://www.youtube.com/watch?v=SMH5uL3HiiU>
- Come to office hours to go over examples



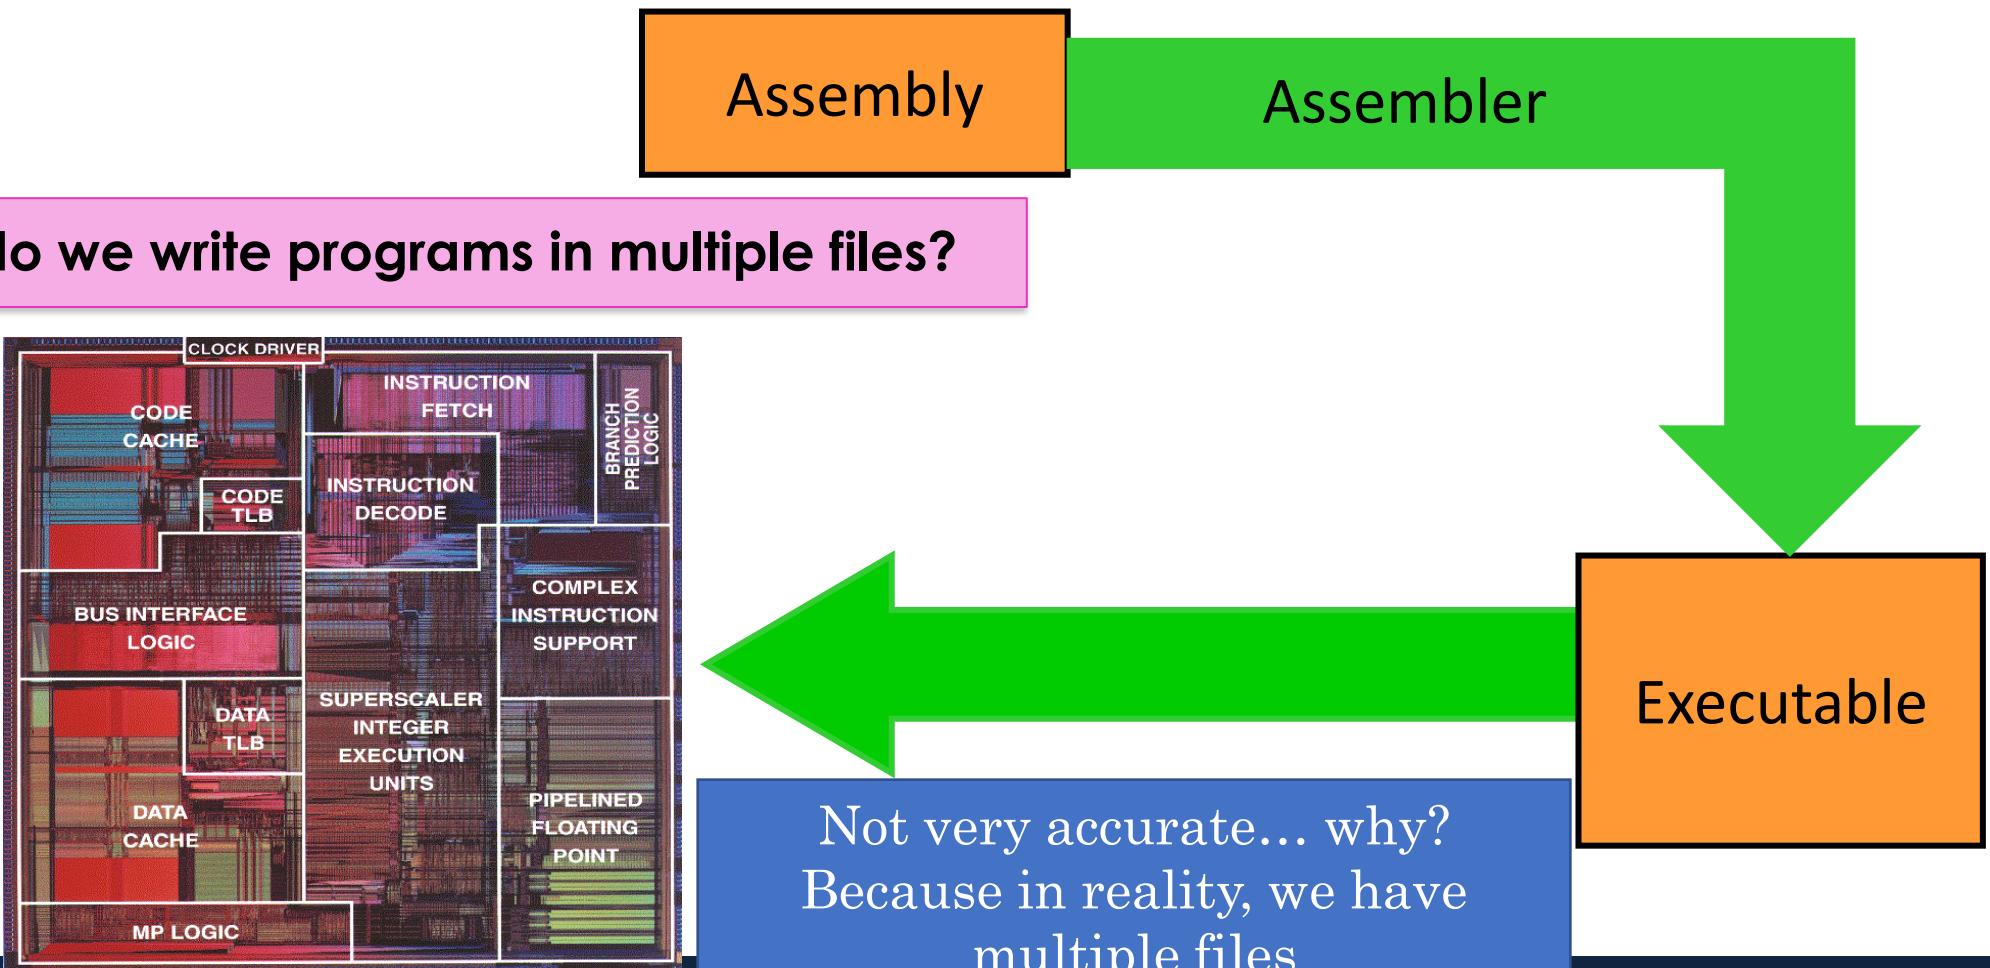
Today we'll finish up software

- Introduce linkers and loaders
 - Basic relationship of complier, assembler, linker and loader.
 - Object files
 - Symbol tables and relocation tables



Source Code to Execution

- In project 1a, our view is this:



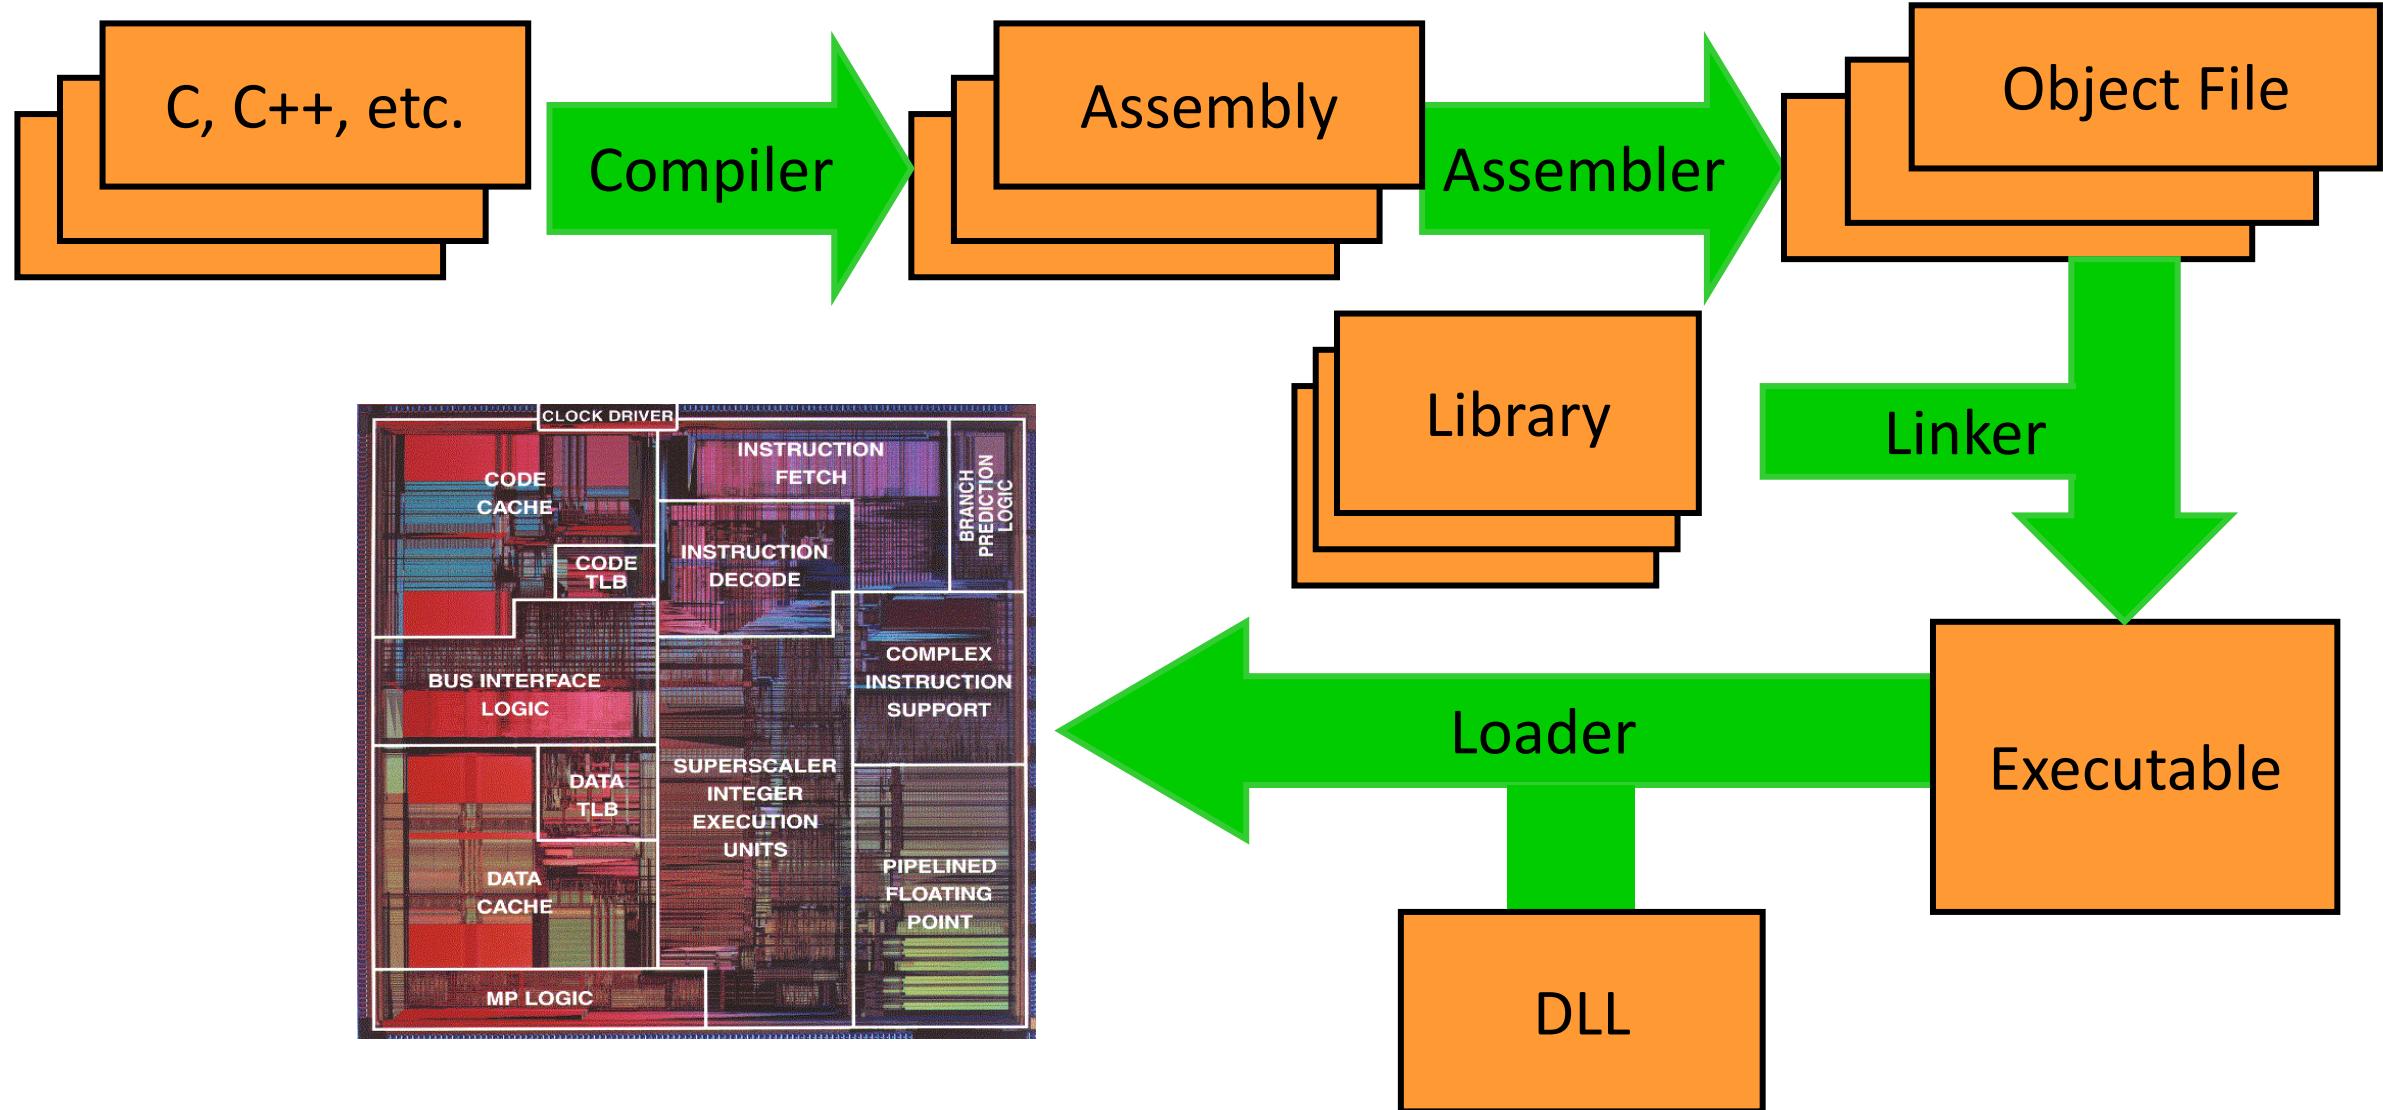
Multi-file programs

- In practice, programs are made from thousands or millions of lines of code
 - Use pre-existing libraries like stdlib
- If we change one line, do we need to recompile the whole thing?
 - No! If we compile each file into a separate **object file**, then we only need to recompile that one file and **link** it to the other, unchanged object files



What do object files look like?

Source Code to Execution



What do object files look like?

```
extern int X;  
extern void foo();  
int Y;  
  
void main() {  
    Y = X + 1;  
    foo();  
}
```

"extern" means
defined in another
file

```
extern int Y;  
int X;  
  
void foo() {  
    Y *= 2;  
}
```

.main:
LDUR X1, [XZR, X]
ADDI X9, X1, #1
STUR X9, [XZR, Y]
BL foo
HALT

Compile →

Compile →

Uh-oh!
Don't know
address of X, Y,
or foo!

.foo:
LDUR X1, [XZR, Y]
LSL X9, X1, #1
STUR X9, [XZR, Y]
BR X30

Linking

```
.main:  
LDUR X1, [XZR, X]  
ADDI X9, X1, #1  
STUR X9, [XZR, Y]  
BL foo  
HALT
```

```
.foo:  
LDUR X1, [XZR, Y]  
LSL X9, X1, #1  
STUR X9, [XZR, Y]  
BR X30
```

What needs to go
in this intermediate
"object file"?

Assemble

???

LINK

Assemble

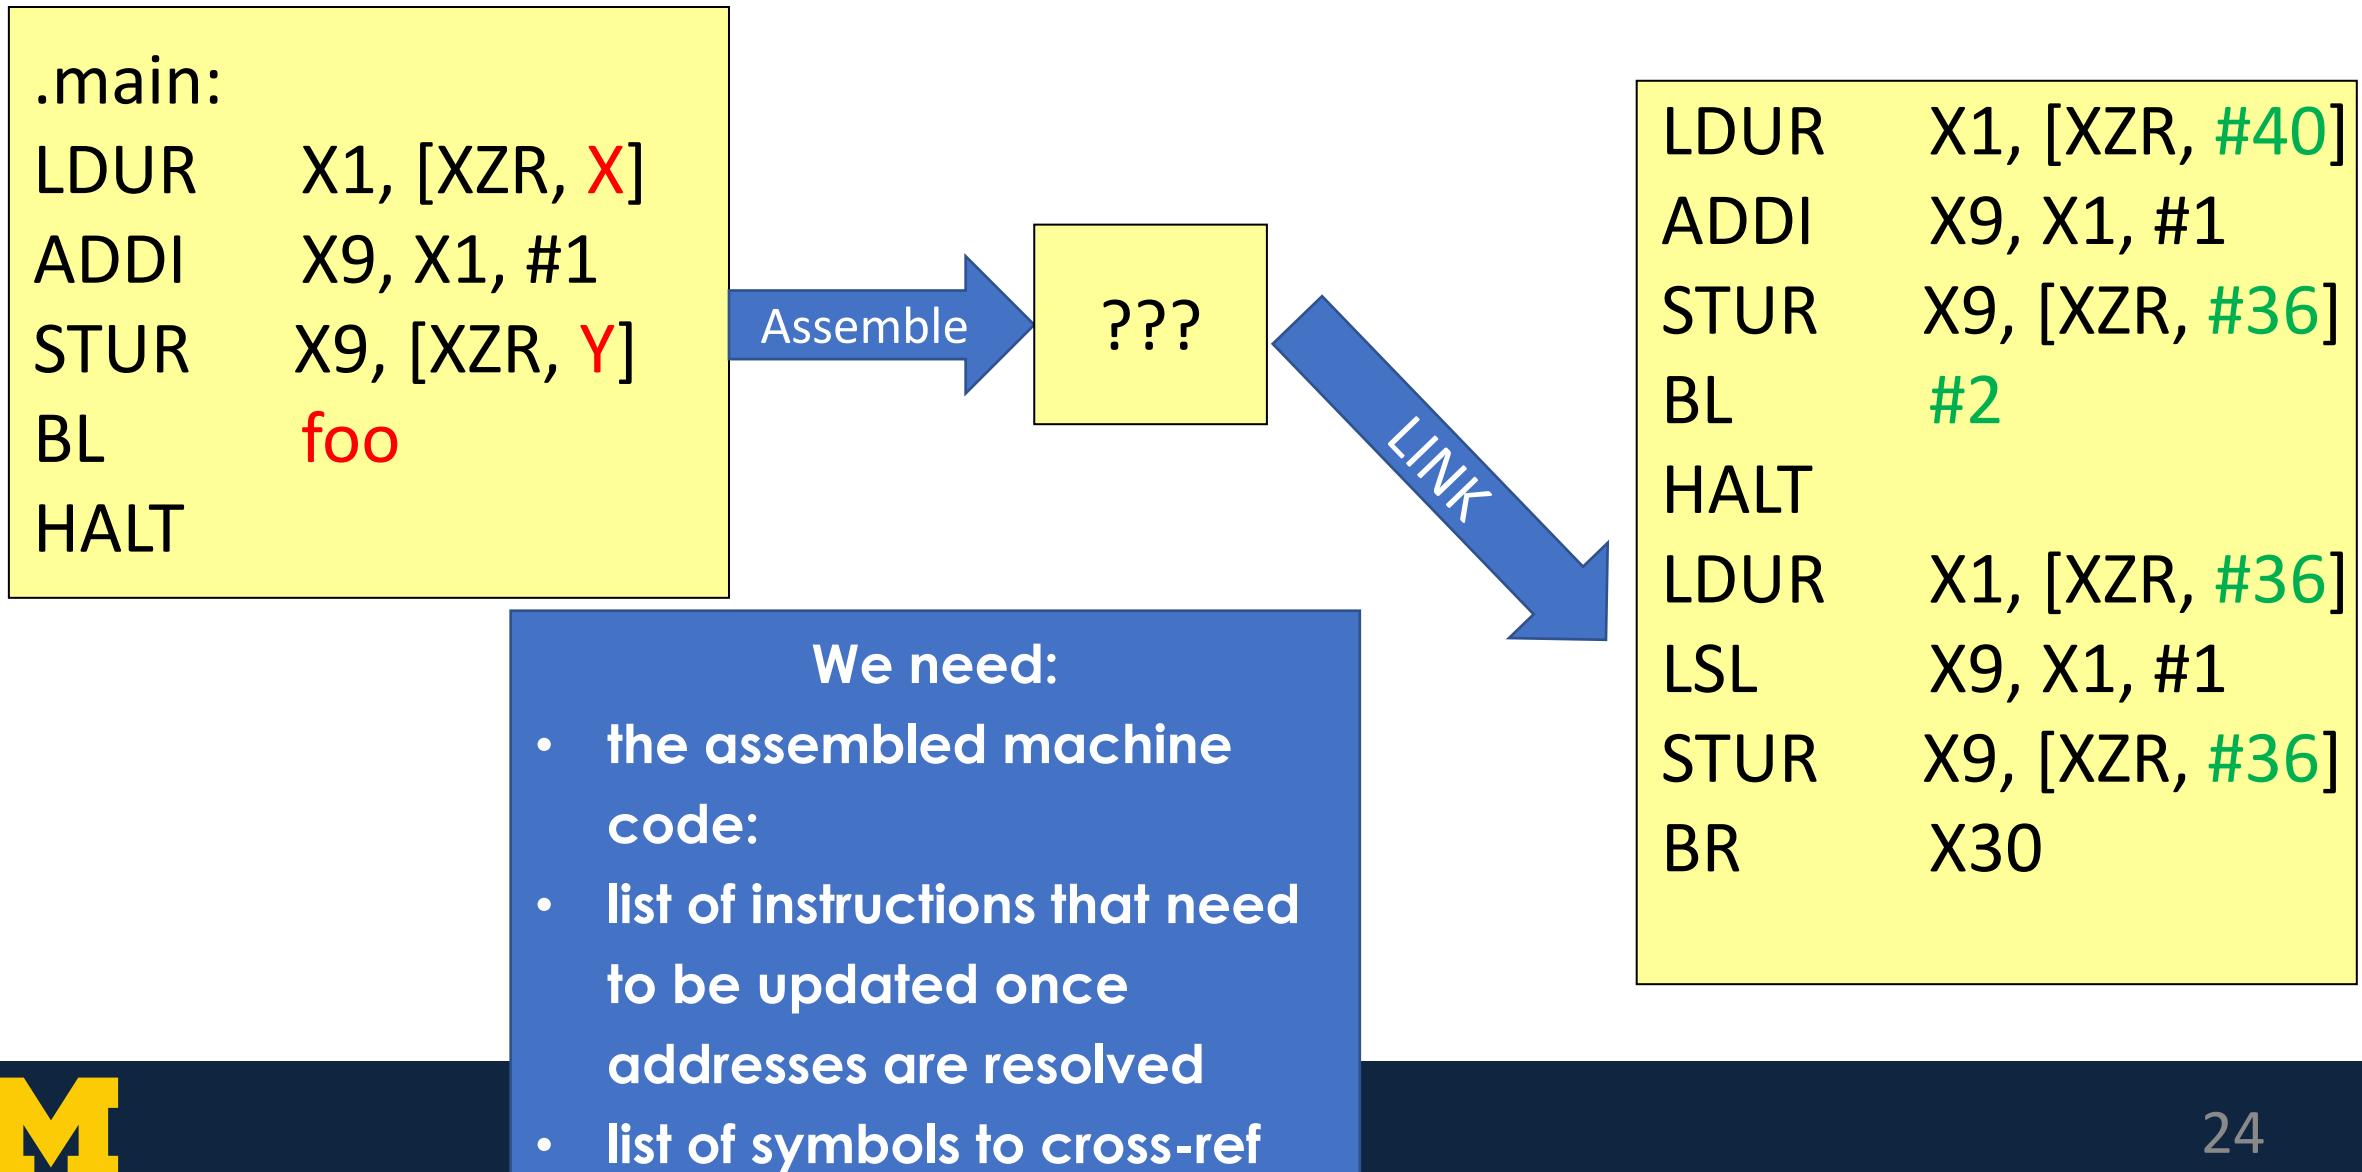
???

LINK

NOTE: this will
actually be in
machine code, not
assembly

```
LDUR X1, [XZR, #40]  
ADDI X9, X1, #1  
STUR X9, [XZR, #36]  
BL #2  
HALT  
LDUR X1, [XZR, #36]  
LSL X9, X1, #1  
STUR X9, [XZR, #36]  
BR X30  
// Addr #36 starts here
```

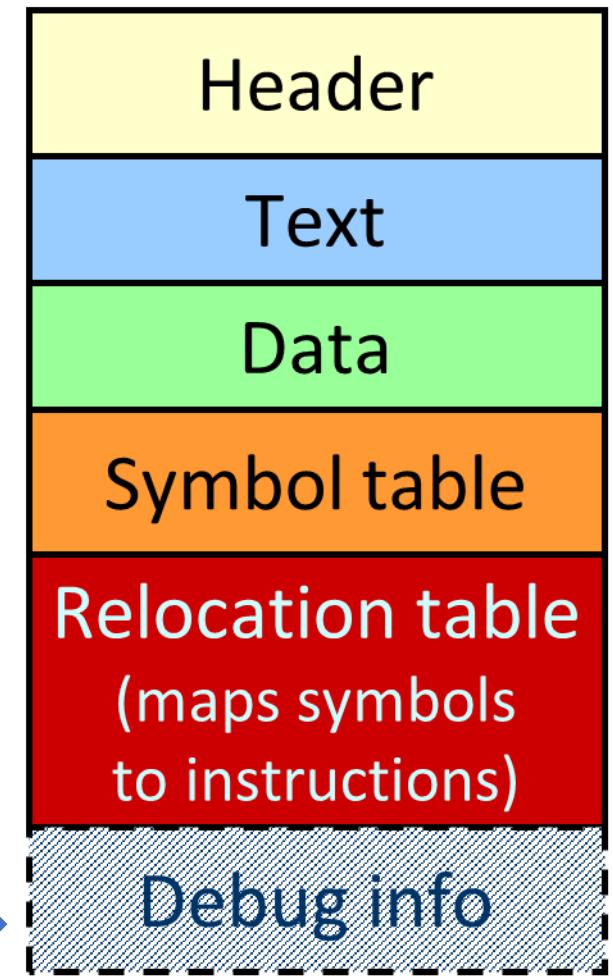
Linking



What do object files look like?

- Since we can't make executable, we make an object file
- Basically, includes the machine code that will go in the executable
 - Plus extra information on what we need to modify once we stitch all the other object files together
- Looks like this ->

Object code format



Assembly → Object file - example

```
extern int G;  
extern void B();  
int X = 3;  
main() {  
    Y = G + 1;  
    B();  
}
```

```
LDUR    X1, [XZR, G]  
ADDI    X9, X1, #1  
BL      B
```

| Header | | | |
|--------------|-----------|------------------------|------------|
| | Name | foo | |
| | Text size | 0x0C //probably bigger | |
| | Data size | 0x04 //probably bigger | |
| Text | | | |
| | Address | Instruction | |
| | 0 | LDUR X1, [XZR, G] | |
| | 4 | ADDI X9, X1, #1 | |
| | 8 | BL B | |
| Data | | | |
| | 0 | X | 3 |
| Symbol table | | | |
| | Label | Address | |
| | X | 0 | |
| | B | - | |
| | main | 0 | |
| | G | - | |
| Reloc table | | | |
| | Addr | Instruction type | Dependency |
| | 0 | LDUR | G |
| | 8 | BL | B |

Assembly → Object file - example

```
extern int X;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

Header:
keeps track of
size of each
section

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

| | | |
|---------------------|-----------|-----------------------------|
| Header | Name | foo |
| | Text size | 0x0C //probably bigger |
| | Data size | 0x04 //probably bigger |
| Text | Address | Instruction |
| | 0 | LDUR X1, [XZR, G] |
| | 4 | ADDI X9, X1, #1 |
| | 8 | BL B |
| Data | 0 | X 3 |
| Symbol table | Label | Address |
| | X | 0 |
| | B | - |
| | main | 0 |
| | G | - |
| Reloc table | Addr | Instruction type Dependency |
| | 0 | LDUR G |
| | 8 | BL B |



Assembly → Object file - example

```
extern int G;  
extern void B();  
int X = 3;  
main() {  
    Y = G + B();  
}
```

Text:
machine code

```
LDUR    X1, [XZR, G]  
ADDI    X9, X1, #1  
BL      B
```

| Header | Name | foo |
|--------------|-----------|-----------------------------|
| | Text size | 0x0C //probably bigger |
| | Data size | 0x04 //probably bigger |
| Text | Address | Instruction |
| | 0 | LDUR X1, [XZR, G] |
| | 4 | ADDI X9, X1, #1 |
| | 8 | BL B |
| Data | 0 | X 3 |
| Symbol table | Label | Address |
| | X | 0 |
| | B | - |
| | main | 0 |
| | G | - |
| Reloc table | Addr | Instruction type Dependency |
| | 0 | LDUR G |
| | 8 | BL B |

Simplifying Assumption for EECS370

All globals and static locals (initialized or not) go in the data segment

Assembly → Object file - example

```
extern int G;  
extern void B();  
int X = 3;  
main() {  
    Y = G + 1;  
    B();  
}
```

Data:
initialized globals
and static locals

```
LDUR    X1, [XZR, G]  
ADDI    X9, X1, #1  
BL      B
```

| | | |
|---------------------|-----------|-----------------------------|
| Header | Name | foo |
| | Text size | 0x0C //probably bigger |
| | Data size | 0x04 //probably bigger |
| Text | Address | Instruction |
| | 0 | LDUR X1, [XZR, G] |
| | 4 | ADDI X9, X1, #1 |
| | 8 | BL B |
| Data | 0 | X 3 |
| Symbol table | Label | Address |
| | X | 0 |
| | B | - |
| | main | 0 |
| | G | - |
| Reloc table | Addr | Instruction type Dependency |
| | 0 | LDUR G |
| | 8 | BL B |



Assembly → Object file - example

```
extern int G;  
extern void B();  
int X = 3;  
main() {  
    Y = G + 1;  
    B();  
}
```

Symbol table:
Lists all labels
visible outside this file
(i.e. function names
and global variables)

LDUR
ADDI
BL

| | | |
|---------------------|-----------|------------------------|
| Header | Name | foo |
| | Text size | 0x0C //probably bigger |
| | Data size | 0x04 //probably bigger |
| Text | Address | Instruction |
| | 0 | LDUR X1, [XZR, G] |
| | 4 | ADDI X9, X1, #1 |
| | 8 | BL B |
| Data | 0 | X |
| | | 3 |
| Symbol table | Label | Address |
| Symbol table | X | 0 |
| | B | - |
| | main | 0 |
| | G | - |
| Reloc table | Addr | Instruction type |
| | 0 | LDUR |
| | 8 | BL |
| | | Dependency |
| | | G |
| | | B |

Assembly → Object file - example

```
extern int G;  
extern void B();  
int X = 3;  
main() {  
    Y = G + 1;  
    B();  
}
```

LDUR X1 [XZR, G]

Relocation Table:

list of instructions and data words that must be updated if things are moved in memory

| | | |
|---------------------|-----------|-----------------------------|
| Header | Name | foo |
| | Text size | 0x0C //probably bigger |
| | Data size | 0x04 //probably bigger |
| Text | Address | Instruction |
| | 0 | LDUR X1, [XZR, G] |
| | 4 | ADDI X9, X1, #1 |
| | 8 | BL B |
| Data | 0 | X |
| | | 3 |
| Symbol table | Label | Address |
| | X | 0 |
| | B | - |
| | main | 0 |
| | G | - |
| Reloc table | Addr | Instruction type Dependency |
| | 0 | LDUR G |
| | 8 | BL B |

Class Problem 1

Poll: Which symbols will be put in the symbol table? (i.e. which "things" should be visible to all files?)

file1.c

```
extern void bar(int);
extern char c[];
int a;
int foo (int x) {
    int b;
    a = c[3] + 1;
    bar(x);
    b = 27;
}
```

file 1 – symbol table

| sym | loc |
|------------|------------|
| a | data |
| foo | text |
| c | - |
| bar | - |

file2.c

```
extern int a;
char c[100];
void bar (int y) {
    char e[100];
    a = y;
    c[20] = e[7];
}
```

file 2 – symbol table

| sym | loc |
|------------|------------|
| c | data |
| bar | text |
| a | - |



Poll: Which lines / instructions are in the relocation table? (i.e. which "things" need to be updated after linking?)

Class Problem 2

file1.c

```
1 extern void bar(int);
2 extern char c[];
3 int a;
4 int foo (int x) {
5     int b;
6     a = c[3] + 1;
7     bar(x);
8     b = 27;
9 }
```

file 1 - relocation table

| line | type | dep |
|------|------|-----|
| 6 | ldur | c |
| 6 | stur | a |
| 7 | bl | bar |

file2.c

```
1 extern int a;
2 char c[100];
3 void bar (int y) {
4     char e[100];
5     a = y;
6     c[20] = e[7];
7 }
```

file 2 - relocation table

| line | type | dep |
|------|------|-----|
| 5 | stur | a |
| 6 | stur | c |

Note: in a real relocation table, the “line” would really be the address in “text” section of the instruction we need to update.



Linker

- Stitches independently created object files into a single executable file (i.e., `a.out`)
 - Step 1: Take text segment from each `.o` file and put them together.
 - Step 2: Take data segment from each `.o` file, put them together, and concatenate this onto end of text segments.
- What about libraries?
 - Libraries are just special object files.
 - You create new libraries by making lots of object files (for the components of the library) and combining them (see `ar` and `ranlib` on Unix machines).
- Step 3: Resolve cross-file references to labels
 - Make sure there are no undefined labels



Linker - Continued

- What kind of instructions get relocated?
- PC-relative branches? (if/else or loops)
 - No – amount we're branching forwards/backwards doesn't change after sliding instructions around
- Local variable accesses?
 - No – memory addresses are relative to stack pointer (SP) value
 - Relative offsets won't change
 - SP value will – but that's a dynamic value anyway, isn't encoded in instruction
- Global / static local variable access?
 - Yes – these use hardcoded addresses which change after linking

```
B.EQ    endLoop
```

```
sub    sp, sp, #16
mov    x0, #42
stur   x0, [sp, 0]
ldur   x1, [sp, 0]
```

```
ldur   x2, [0, MY_VAR]
```

Loader

- Executable file is sitting on the disk
- Puts the executable file code image into memory and asks the operating system to schedule it as a new process
 - Creates new address space for program large enough to hold text and data segments, along with a stack segment
 - Copies instructions and data from executable file into the new address space
 - Initializes registers (PC and SP most important)
- Take operating systems class (EECS 482) to learn more!



Summary

- Compiler converts a single source code file into a single assembly language file
- Assembler handles directives (.fill), converts what it can to machine language, and creates a checklist for the linker (relocation table). This changes each .s file into a .o file
- Assembler does 2 passes to resolve addresses, handling internal forward references
- Linker combines several .o files and resolves absolute addresses
- Linker enables separate compilation: Thus unchanged files, including libraries need not be recompiled.
- Linker resolves remaining addresses.
- Loader loads executable into memory and begins execution

Next Time

- Floating Point Arithmetic
- And... hardware time, baby!

