

Poll: How was the midterm
compared to practice exams?

EECS 370 - Lecture 13

Pipelining and Data Hazards



Live Poll + Q&A: [slido.com #eeecs370](https://slido.com/#eeecs370)

Poll and Q&A Link

Announcements

- P2
 - P2L+P2R due next Thursday
- Lab 7
 - No prelab this week
 - Will just be providing help on P2R
 - **Attendance is required unless you've completed P2R by the time your lab meets**



Single vs Multi-cycle Performance

1 ns – Register File read/write time

2 ns – ALU/adder

2 ns – memory access

0 ns – MUX, PC access, sign extend,
ROM

- Assuming the above delays, what is the best cycle time that the LC2k multi-cycle datapath could achieve? Single cycle?

$$MC: \text{MAX}(2, 1, 2, 2, 1) = 2\text{ns}$$

$$SC: 2 + 1 + 2 + 2 + 1 = 8 \text{ ns}$$

- Assuming the above delays, for a program consisting of 25 LW, 10 SW, 45 ADD, and 20 BEQ, which is faster?

$$SC: 100 \text{ cycles} * 8 \text{ ns} = 800 \text{ ns}$$

$$MC: (25*5 + 10*4 + 45*4 + 20*4)\text{cycles} * 2\text{ns} = 850 \text{ ns}$$



Single and Multi-cycle performance

- Wait, multi-cycle is worse??
- For our ISA, most instructions take about the same time
- Multi-cycle shines when some instructions take much longer
- E.g. if we add a long latency instruction like multiply:
 - Let's say operation takes 10 ns, but could be split into 5 stages of 2 ns
 - SC: clock period = 16 ns, performance is 1600 ns
 - MC: clock period = 2 ns, performance is 850 ns



Performance Metrics – Execution time

- What we really care about in a program is **execution time**
 - **Execution time** = total instructions executed X CPI x clock period
 - The "Iron Law" of performance
- CPI = **average** number of clock **cycles per instruction for an application**
- To calculate multi-cycle CPI we need:
 - Cycles necessary for each type of instruction
 - Mix of instructions executed in the application (dynamic instruction execution profile)

Poll: What are the units of
(instructions executed x CPI x
clock period)?

Datapath Summary

- Single-cycle processor
 - CPI = 1 (by definition)
 - clock period = ~10 ns
- Multi-cycle processor
 - CPI = ~4.25
 - clock period = ~2 ns
- Better design:
 - CPI = 1
 - clock period = ~2ns
- How??
 - Work on multiple instructions at the same time




Pipelining


- Want to execute an instruction?
 - Build a processor (multi-cycle)
 - Find instructions
 - Line up instructions (1, 2, 3, ...)
 - Overlap execution
 - Cycle #1: Fetch 1
 - Cycle #2: Decode 1 Fetch 2
 - Cycle #3: ALU 1 Decode 2 Fetch 3
 -
 - This is called pipelining instruction execution.
 - Used extensively for the first time on IBM 360 (1960s).
 - CPI approaches 1.



Pipelining



Pipelining




Pipelined implementation of LC2K

- Break the execution of the instruction into cycles.
 - Similar to the multi-cycle datapath
- Design a separate datapath **stage** for the execution performed during each cycle.
 - Build **pipeline registers** to communicate between the stages.
 - Whatever is on the left gets written onto the right during the next cycle
 - Kinda like the **Instruction Register** in our multi-cycle design, but we'll need one for each stage



Our new pipelined datapath




Sample Code (Simple)


Let's run the following code on pipelined LC2K:

- add 1 2 3 ; reg 3 = reg 1 + reg 2
- nor 4 5 6 ; reg 6 = reg 4 nor reg 5
- lw 2 4 20 ; reg 4 = Mem[reg2+20]
- add 2 5 5 ; reg 5 = reg 2 + reg 5
- sw 3 7 10 ; Mem[reg3+10] =reg 7


Pipeline datapath



Time 0 - Initial state




Time 1 - Fetch: add 1 2 3




add 1 2 3


Time 2 - Fetch: nor 4 5 6




Time 3 - Fetch: lw 2 4 20




Time 4 - Fetch: add 2 5 5



Time 5 - Fetch: sw 3 7 10




Time 6 – no more instructions




Time 7 – no more instructions


Poll: What all happens
on the next clock cycle?



Time 8 – no more instructions



Time 9 – no more instructions



SW

Pipelining - What can go wrong?



- **Data hazards**: since register reads occur in stage 2 and register writes occur in stage 5 it is possible to read an old / stale value before the correct value is written back.
- **Control hazards**: A branch instruction may change the PC, but not until stage 4. What do we fetch before that?
- **Exceptions**: Sometimes we need to pause execution, switch to another task (maybe the OS), and then resume execution... how do we make sure we resume at the right spot
- **Now - Data hazards**
 - What are they?
 - How do you detect them?
 - How do you deal with them?

Pipeline function for ADD

- Fetch: read instruction from memory
- Decode: **read source operands from reg**
- Execute: calculate sum
- Memory: pass results to next stage
- Writeback: **write sum into register file**


Data Hazards

Recall: registers
are read /sourced
In the “decode” stage



If not careful, nor will read a stale value of register 3

Data Hazards



Assume Register File gives the right value of **register 3** when read/written during same cycle. This is consistent with most processors (ARM/x86), but not Project 3.

Definitions

- Data Dependency: *one instruction uses the result of a previous one*
 - Doesn't necessarily cause a problem
- Data Hazard: *one instruction has a data dependency that will cause a problem if we don't "deal with it"*

Class Problem 1

Poll: Which of these instructions has a data dependency on an earlier one? Which of those are data hazards in our 5-stage pipeline?

1. add 1 2 3
2. nor 3 4 5
3. add 6 3 7
4. lw 3 6 10
5. sw 6 2 12

What about here?

1. add 1 2 3
2. beq 3 4 1
3. add 3 5 6
4. add 3 6 7

Class Problem 1


Which read-after-write (RAW) dependences do you see?

Which of those are data hazards?

1. add 1 2 3
2. nor 3 4 5
3. add 6 3 7
4. lw 3 6 10
5. sw 6 2 12

What about here?

1. add 1 2 3
2. beq 3 4 1
3. add 3 5 6
4. add 3 6 7



Three approaches to handling data hazards

- Avoid
 - Make sure there are no hazards in the code
- Detect and Stall
 - If hazards exist, stall the processor until they go away.
- Detect and Forward
 - If hazards exist, fix up the pipeline to get the correct value (if possible)