

EECS 370 - Lecture 14

Pipelining and Data Hazards II




Live Poll + Q&A: [slido.com #eeecs370](https://slido.com/#eeecs370)

Poll and Q&A Link

Announcements

- P2
 - P2L+P2R due Thursday
- Lab 8
 - **Pre-lab quiz due Thursday at 11:55 pm**
- Midterm scores out
 - Regrade requests by end of day Friday
 - Will announce any grade threshold changes after

Review Grades for Midterm



My OH

- My office hours today will be in 2901 BBB
- I'm holding extra OH this Wednesday @ 4pm over Zoom




Review: Pipelining

- Goal:
 - Achieve low clock period of multi-cycle processor...
 - ... while maintaining low cycles-per-instruction (CPI) of single cycle processor (close to 1)
 - Achieve this by overlapping execution of multiple instructions simultaneously





Our new pipelined datapath



Data Hazards

Recall: registers
are read /sourced
In the “decode” stage



If not careful, nor will read a stale value of register 3

Definitions

- Data Dependency: *one instruction uses the result of a previous one*
 - Doesn't necessarily cause a problem
- Data Hazard: *one instruction has a data dependency that will cause a problem if we don't "deal with it"*



Class Problem 1

Poll: Which of these instructions has a data dependency on an earlier one? Which of those are data hazards in our 5-stage pipeline?

1. add 1 2 3
2. nor 3 4 5
3. add 6 3 7
4. lw 3 6 10
5. sw 6 2 12

What about here?

1. add 1 2 3
2. beq 3 4 1
3. add 3 5 6
4. add 3 6 7

Class Problem 1

Which read-after-write (RAW) dependences do you see?

Which of those are data hazards?

1. add 1 2 3

2. nor 3 4 5

3. add 6 3 7

4. lw 3 6 10

5. sw 6 2 12

What about here?


1. add 1 2 3
2. beq 3 4 1
3. add 3 5 6
4. add 3 6 7

Three approaches to handling data hazards

- Avoid
 - Make sure there are no hazards in the code
- Detect and Stall
 - If hazards exist, stall the processor until they go away.
- Detect and Forward
 - If hazards exist, fix up the pipeline to get the correct value (if possible)

Handling data hazards I: Avoid all hazards

- Assume the programmer (or the compiler) knows about the processor implementation.
 - Make sure no hazards exist.
 - Put noops between any dependent instructions.




Problems with this solution


- Old programs (legacy code) may not run correctly on new implementations
 - Longer pipelines need more noops
- Programs get larger as noops are included
 - Especially a problem for machines that try to execute more than one instruction every cycle
 - Intel EPIC: Often 25% - 40% of instructions are noops
- Program execution is slower
 - CPI is 1, but some instructions are noops


Handling data hazards II: Detect and stall until ready


- Detect:
 - Compare regA with previous DestRegs
 - 3 bit operand fields
 - Compare regB with previous DestRegs
 - 3 bit operand fields
- Stall:
 - Keep current instructions in fetch and decode
 - Pass a noop to execute
- How do we modify the pipeline to do this?

Our pipeline currently does not handle hazards—let's fix it









0 1 1

regA

regB

0 1 1

3

Example

- Let's run this program with a data hazard through our 5-stage pipeline

add 1 2 3


nor 3 4 5

- We will start at the beginning of cycle 3, where add is in the EX stage, and nor is in the ID stage, about to read a register value


Time:	1	2	3
add 1 2 3	IF	ID	EX
nor 3 4 5		IF	ID

Hazard!


First half of cycle 3




First half of cycle 4




End of cycle 4



First half of cycle 5



End of cycle 5



Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF												
nor 3 4 5													

Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nor 3 4 5		IF	ID*	ID*	ID	EX	ME	WB					
add 6 3 7													
lw 3 6 10													
sw 6 2 12													

Solution

Poll: Which problems does "detect and stall" fix over "avoid hazards"? (select all)

1. Breaking backwards compatibility
2. Larger programs
3. Slower programs

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nor 3 4 5		IF	ID*	ID*	ID	EX	ME	WB					
add 6 3 7					IF	ID	EX	ME	WB				
lw 3 6 10						IF	ID	EX	ME	WB			
sw 6 2 12							IF	ID*	ID*	ID	EX	ME	WB



Problems with detect and stall

- CPI increases every time a hazard is detected!
- Is that necessary? Not always!
 - Re-route the result of the **add** to the **nor**
 - **nor** no longer needs to read R3 from reg file
 - It can get the data later (when it is ready)
 - This lets us complete the decode this cycle
 - But we need more control logic



Handling data hazards III: Detect and forward


- Detect: same as detect and stall
 - Except that all 4 hazards have to be treated differently
 - i.e., you can't logical-OR the 4 hazard signals
- Forward:
 - New **bypass datapaths** route computed data to where it is needed
 - New MUX and control to pick the right data
- **Beware:** Stalling may still be required even in the presence of forwarding

Forwarding example


- We will use this program for the next example
(same as last pipeline diagram example)

1. add	1 2 3
2. nor	3 4 5
3. add	6 3 7
4. lw	3 6 10
5. sw	6 2 12


First half of cycle 3




End of cycle 3




First half of cycle 4




End of cycle 4




First half of cycle 5




End of cycle 5




First half of cycle 6




End of cycle 6




First half of cycle 7




End of cycle 7



First half of cycle 8



End of cycle 8



Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nor 3 4 5		IF	ID	EX	ME								
add 6 3 7			IF	ID	EX								
lw 3 6 10				IF	ID								
sw 6 2 12					IF								

Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nor 3 4 5		IF	ID	EX	ME	WB							
add 6 3 7			IF	ID	EX	ME	WB						
lw 3 6 10				IF	ID	EX	ME	WB					
sw 6 2 12					IF	ID*	ID	EX	ME	WB			

Extra Slides




Other issues

- What other instruction(s) have we been ignoring so far??
- Branches!! (Let's not worry about jumps yet)
- Sequence for BEQ:
 - Fetch: read instruction from memory
 - Decode: read source operands from registers
 - Execute: calculate target address and test for equality
 - Memory: Send target to PC if test is equal
 - Writeback: nothing
 - Branch Outcomes
 - Not Taken
 - $PC = PC + 1$
 - Taken
 - $PC = \text{Branch Target Address}$



Control Hazards

beq	1	1	10
add	3	4	5



Approaches to handling control hazards

- 3 strategies – similar to handling data hazards
 - 1. Avoid
 - Make sure there are no hazards in code
 - 2. Detect and stall
 - Delay fetch until branch resolved
 - 3. Speculate and squash-if-wrong
 - Guess outcome of branch
 - Fetch instructions assuming we're right
 - Stop them if they shouldn't have been executed

Avoiding Control Hazards


- Don't have branch instructions!
 - Possible, but not practical
 - ARM offers **predicated** instructions (instructions that throw away result if some condition is not met)
 - Allows replacement of if/else conditions
 - Hard to use for everything
 - Not covered more in this class





Detect and Stall

- Detection
 - Wait until decode
 - Check if opcode == beq or jalr
- Stall
 - Keep current instruction in fetch
 - Insert noops
 - Pass noop to decode stage, not execute!











Target PC is now sent to memory, everything moves forward as normal





Control Hazards

beq	1	1	10
add	3	4	5



beq fetch decode execute memory writeback

add fetch fetch fetch fetch

or

Target: fetch

Problems with Detect and Stall

- CPI increases every time a branch is detected!
- Is that necessary? Not always!
 - Branch not always taken
 - Let's assume it is NOT taken...
 - In this case, we can ignore the beq (treat it like a noop)
 - Keep fetching PC + 1
 - What if we're wrong?
 - OK, as long as we do not COMPLETE any instruction we mistakenly execute
 - I.e. DON'T write values to register file or memory



Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- **Speculate-and-Squash**
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor




Speculate and Squash


- Speculate: assume not equal
 - Keep fetching from PC+1 until we know that the branch is really taken
- Squash: stop bad instructions if taken
 - Send a noop to Decode, Execute, and Memory
 - Sent target address to PC




Resolve branch in
execute stage



Send target to PC
in Mem stage,
squash previous 3
instructions



3 noops inserted,
correct PC being
sent to memory



Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- **Exceptions**
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor



What can go wrong?

- ❑ **Data hazards**: since register reads occur in stage 2 and register writes occur in stage 5 it is possible to read the wrong value if is about to be written.
- ❑ **Control hazards**: A branch instruction may change the PC, but not until stage 4. What do we fetch before that?
- ❑ **Exceptions**: How do you handle exceptions in a pipelined processor with 5 instructions in flight?

Exceptions

- ❑ Exception: when something unexpected happens during program execution.
 - Example: divide by zero.
 - The situation is more complex than the hardware can handle
 - So the hardware branches to a function, an “exception handler” which is code to try to deal with the problem.
- ❑ The exact way to set up such an exception handler will vary by ISA.
 - With C on x86 you would use <signal.h> functions to handle the “SIGFPE” signal.
 - There is a pretty good [Hackaday article](#) on this if you want to learn more.

Exceptions and Pipelining

- ❑ The hardware branches to the “exception handler”
 - This means that any instruction which can “throw” an exception *could* be a branch.
 - Throwing an exception should be rare (“exceptional”)

- ❑ So we would treat it much like a branch we predicted as “not taken”
 - Squash instructions behind it and then branch
 - It will introduce stalls, but since it should be rare, we don’t worry about it.
 - “Make the common case fast”.

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- **Practice Performance Problems**
 - **Problem 1**
 - Problem 2
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor



Classic performance problem

- ❑ Program with following instruction breakdown:

Iw	10%
sw	15%
beq	25%
R-type	50%
- ❑ Speculate “always not-taken” and squash. 80% of branches not-taken
- ❑ Full forwarding to execute stage. 20% of loads stall for 1 cycle
- ❑ What is the CPI of the program?
- ❑ What is the total execution time per instruction if clock frequency is 100MHz?



Classic performance problem

- ❑ Program with following instruction breakdown:

Iw	10%
sw	15%
beq	25%
R-type	50%
- ❑ Speculate “always not-taken” and squash. 80% of branches not-taken
- ❑ Full forwarding to execute stage. 20% of loads stall for 1 cycle
- ❑ What is the CPI of the program?
- ❑ What is the total execution time per instruction if clock frequency is 100MHz?

$$\begin{aligned}\text{CPI} = & 1 + 0.10 \text{ (loads)} * 0.20 \text{ (load use stall)} * 1 \\ & + 0.25 \text{ (branch)} * 0.20 \text{ (miss rate)} * 3\end{aligned}$$

$$\text{CPI} = 1 + 0.02 + 0.15 = 1.17$$

$$\text{Time} = 1.17 * 10\text{ns} = 11.7\text{ns per instruction}$$



Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - **Problem 2**
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor



Classic performance problem (cont.)

- ❑ Assume branches are resolved at Execute?
 - What is the CPI?
 - What happens to cycle time?

Classic performance problem (cont.)

- ❑ Assume branches are resolved at Execute?
 - What is the CPI?
 - What happens to cycle time?

$$\begin{aligned} \text{CPI} = & 1 + 0.10 \text{ (loads)} * 0.20 \text{ (load use stall)} * 1 \\ & + 0.25 \text{ (branch)} * 0.20 \text{ (miss rate)} * 2 \end{aligned}$$

$$\text{CPI} = 1 + 0.02 + 0.1 = 1.12$$

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - **Problem 3**
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor



Performance with deeper pipelines

- ❑ Assume the setup of the previous problem.
- ❑ What if we have a 10 stage pipeline?
 - Instructions are fetched at stage 1.
 - Register file is read at stage 3.
 - Execution begins at stage 5.
 - Branches are resolved at stage 7.
 - Memory access is complete in stage 9.
- ❑ What's the CPI of the program?
- ❑ If the clock rate was doubled by doubling the pipeline depth, is performance also doubled?

Performance with deeper pipelines

- ❑ Assume the setup of the previous problem.
- ❑ What if we have a 10 stage pipeline?
 - Instructions are fetched at stage 1.
 - Register file is read at stage 3.
 - Execution begins at stage 5.
 - Branches are resolved at stage 7.
 - Memory access is complete in stage 9.
- ❑ What's the CPI of the program?
- ❑ If the clock rate was doubled by doubling the pipeline depth, is performance also doubled?

$$\text{CPI} = 1 + 0.10 \text{ (loads)} * 0.20 \text{ (load use stall)} * 4 + 0.25 \text{ (branch)} * 0.20 \text{ (N stalls)} * 6$$

$$\text{CPI} = 1 + 0.08 + 0.30 = 1.38$$

$$\text{Time} = 1.38 * 5\text{ns} = 6.9 \text{ ns per instruction}$$

