

EECS 370

Control Hazards and Performance



Live Poll + Q&A: [slido.com #eeecs370](https://slido.com/#eeecs370)

Reminders

- Lab meets Friday
 - Don't forget pre-lab quiz (due today)!
- P2R + P2L due tonight
- Midterm scores published
 - Submit regrade requests by Friday 11:55 pm



Pipelining - What can go wrong?

- **Data hazards**: since register reads occur in stage 2 and register writes occur in stage 5 it is possible to read an old / stale value before the correct value is written back.
- **Control hazards**: A branch instruction may change the PC, but not until stage 4. What do we fetch before that?
- **Exceptions**: Sometimes we need to pause execution, switch to another task (maybe the OS), and then resume execution... how do we make sure we resume at the right spot
- Now – control hazards


Other issues

- What other instruction(s) have we been ignoring so far??
- Branches!! (Let's not worry about jumps yet)
- Sequence for BEQ:
 - Fetch: read instruction from memory
 - Decode: read source operands from registers
 - Execute: calculate target address and test for equality
 - Memory: Send target to PC if test is equal
 - Writeback: nothing
 - Branch Outcomes
 - Not Taken
 - $PC = PC + 1$
 - Taken
 - $PC = \text{Branch Target Address}$



Control Hazards

beq	1	1	10
add	3	4	5



Approaches to handling control hazards

- 3 strategies – similar to handling data hazards
 - 1. Avoid
 - Make sure there are no hazards in code
 - 2. Detect and stall
 - Delay fetch until branch resolved
 - 3. Speculate and squash-if-wrong
 - Guess outcome of branch
 - Fetch instructions assuming we're right
 - Stop them if they shouldn't have been executed

Avoiding Control Hazards


- Don't have branch instructions!
 - Possible, but not practical
 - ARM offers **predicated** instructions (instructions that throw away result if some condition is not met)
 - Allows replacement of if/else conditions
 - Hard to use for everything
 - Not covered more in this class





Detect and Stall

- Detection
 - Wait until decode
 - Check if opcode == beq or jalr
- Stall
 - Keep current instruction in fetch
 - Insert noops
 - Pass noop to decode stage, not execute!











beq DOESN'T
disable PC and
sends target, but
still inserts NOOP


Target PC is now sent to memory, everything moves forward as normal





Control Hazards

beq	1	1	10
add	3	4	5



beq fetch decode execute memory writeback

add fetch fetch fetch fetch

or

Target: fetch

Problems with Detect and Stall

- CPI increases every time a branch is detected!
- Is that necessary? Not always!
 - Branch not always taken
 - Let's assume it is NOT taken...
 - In this case, we can ignore the beq (treat it like a noop)
 - Keep fetching PC + 1
 - What if we're wrong?
 - OK, as long as we do not COMPLETE any instruction we mistakenly execute
 - I.e. DON'T write values to register file or memory



Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- **Speculate-and-Squash**
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor




Speculate and Squash


- Speculate: assume not equal
 - Keep fetching from PC+1 until we know that the branch is really taken
- Squash: stop bad instructions if taken
 - Send a noop to Decode, Execute, and Memory
 - Sent target address to PC




Resolve branch in
execute stage



Send target to PC
in Mem stage,
squash previous 3
instructions



3 noops inserted,
correct PC being
sent to memory



Classic performance problem

- ❑ Program with following instruction breakdown:

Iw	10%
sw	15%
beq	25%
R-type	50%
- ❑ Speculate “always not-taken” and squash. 80% of branches not-taken
- ❑ Full forwarding to execute stage. 20% of loads stall for 1 cycle
- ❑ What is the CPI of the program?
- ❑ What is the total execution time per instruction if clock frequency is 100MHz?



Classic performance problem

- ❑ Program with following instruction breakdown:

Iw 10%

sw 15%

beq 25%

R-type 50%

- ❑ Speculate “always not-taken” and squash. 80% of branches not-taken
- ❑ Full forwarding to execute stage. 20% of loads stall for 1 cycle
- ❑ What is the CPI of the program?
- ❑ What is the total execution time per instruction if clock frequency is 100MHz?

$$\begin{aligned} \text{CPI} = & 1 + 0.10 \text{ (loads)} * 0.20 \text{ (load use stall)} * 1 \\ & + 0.25 \text{ (branch)} * 0.20 \text{ (miss rate)} * 3 \end{aligned}$$

$$\text{CPI} = 1 + 0.02 + 0.15 = 1.17$$

$$\text{Time} = 1.17 * 10\text{ns} = 11.7\text{ns per instruction}$$



Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - **Problem 2**
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor



Classic performance problem (cont.)

- ❑ Assume branches are resolved in EX instead of MEM
 - What is the CPI?
 - What happens to cycle time?



Classic performance problem (cont.)

- ❑ Assume branches are resolved in EX instead of MEM
 - What is the CPI?
 - What happens to cycle time?

$$\begin{aligned} \text{CPI} = & 1 + 0.10 \text{ (loads)} * 0.20 \text{ (load use stall)} * 1 \\ & + 0.25 \text{ (branch)} * 0.20 \text{ (miss rate)} * 2 \end{aligned}$$

$$\text{CPI} = 1 + 0.02 + 0.1 = 1.12$$

Cycle time might go up as we are doing more work at the end of EX

Is it worth it? Depends on whether CPI * clock period is overall lower

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - **Problem 3**
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor



Performance with deeper pipelines

- ❑ Assume the setup of the previous problem.
- ❑ What if we have a 10 stage pipeline?
 - Instructions are fetched at stage 1.
 - Register file is read at stage 3.
 - Execution begins at stage 5.
 - Branches are resolved at stage 7.
 - Memory access is complete at end of stage 9.
- ❑ What's the CPI of the program?
- ❑ If the clock rate was doubled by doubling the pipeline depth, is performance also doubled?

Performance with deeper pipelines

- ❑ Assume the setup of the previous problem.
- ❑ What if we have a 10 stage pipeline?
 - Instructions are fetched at stage 1.
 - Register file is read at stage 3.
 - Execution begins at stage 5.
 - Branches are resolved at stage 7.
 - Memory access is complete at end of stage 9.
- ❑ What's the CPI of the program?
- ❑ If the clock rate was doubled by doubling the pipeline depth, is performance also doubled?

CPI = 1 + 0.10 (loads) *0.20 (load use stall)*4 + 0.25 (branch) * 0.20 (N stalls)*6

CPI = 1 + 0.08 + 0.30 = 1.38

Time = 1.38 * 5ns = 6.9 ns per instruction

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- **Improving Performance with Branch Predicting**
- Simple Direction Predictor
- Improving Direction Predictor



Can We Improve Branch Performance?

- CPI increases every time a branch is taken!
 - About 50%-66% of time
- Is that necessary?
- **No!** We can try to predict when branch is taken
 - But we would need to send target PC to memory before decoding branch
 - How do we:
 1. Know an instruction is a branch before decoding?
 2. Reliably guess whether it should be taken?
 3. Figure out the target PC before executing the branch?

Poll: If you had to guess, in real programs, what's the ratio of taken to not-taken branches?


- a) Very rarely taken
- b) Slightly biased towards not taken
- c) Slightly biased towards taken
- d) Very rarely not taken




Sometimes predict taken?

- When fetching an instruction, need to predict 3 things:
 1. Whether the fetched instruction is a branch
 2. Branch direction (if conditional)
 3. Branch target address (if direction is taken)
- Observation: Target address remains the same for conditional branch across multiple executions
 - Idea: store the target address of branch once we execute it, along with PC of instruction
 - Called Branch Target Buffer (BTB)





Sometimes predict taken?



Poll: Overall, which do you expect to be more accurate?

Branch Direction Prediction

- "Branch direction" refers to whether the branch was taken or not
- Two methods for predicting direction:
 - Static - We predict once during compilation, and that prediction never changes
 - Dynamic - We predict (potentially) many times during execution, and the prediction may change over time
- *Static vs dynamic strategies are a very common topic in computer architecture*




Branch Direction Prediction (Static)

- Always not-taken
 - Simple to implement: no need for BTB, no direction prediction
 - Low accuracy: ~30-40%
 - Compiler can layout code such that the likely path is the “not-taken” path
- Always taken
 - No direction prediction
 - Better accuracy: ~60-70%
 - Backward branches (i.e. loop branches) are usually taken
 - Backward branch: target address lower than branch PC
- Backward taken, forward not taken (BTFN)
 - Predict backward (loop) branches as taken, others not-taken

Branch Direction Prediction (Dynamic)

- Last time predictor
 - Single bit per branch (stored in BTB)
 - Indicates which direction branch went last time it executed
TTTTTTTTNNNNNNNNNN → 90% accuracy
- Always mispredicts the last iteration and the first iteration of a loop branch
 - Accuracy for a loop with N iterations = $(N-2)/N$
 - + Loop branches for loops with large number of iterations
 - Loop branches for loops with small number of iterations
TNTNTNTNTNTNTNTNTN → 0% accuracy


State Machine for Last-Time Prediction



Improving the Last Time Predictor

- Problem: A last-time predictor changes its prediction from T→NT or NT→T too quickly
 - Even though the branch may be mostly taken or mostly not taken
- Solution Idea: Add hysteresis to the predictor so that prediction does not change on a single different outcome
 - Use two bits to track the history of predictions for a branch instead of a single bit
 - Can have 2 states for T or NT instead of 1 state for each

State Machine for 2-bit Saturating Counter



Poll: How many branches do we get wrong?

Two-Bit Counter Based Prediction

- What's the prediction accuracy of a branch with the following sequence of taken/not taken outcomes:

- T T T T N T T N N N T N T N N

Br	T	T	T	T	N	T	T	N	N	N	T	N	T	N	N
State	10	11	11	11	X	10	11	X	X	01	X	01	X	01	00
Pred	T	T	T	T	T	T	T	T	T	N	N	N	N	N	N

Can We Do Better?

- Absolutely... take 470
 - Tons of sophisticated branch predictor designs
- I've worked on a few that found their way into some Chromebooks!



Branch Prediction

- Predict not taken: ~50% accurate
- Predict backward taken: ~65% accurate
- Predict same as last time: ~80% accurate
- Realistic designs: ~96% accurate



Remember this Example from Lecture 1?

- We know understand why sorting improves the inner-loop so much
 - The branch predictor is better at guessing what's gonna happen when data is sorted!

```
for (unsigned c = 0; c < arraySize; ++c)
|   data[c] = std::rand() % 256;
std::sort(data, data + arraySize);

// Test
clock_t start = clock();
long long sum = 0;
// Primary loop
for (unsigned c = 0; c < arraySize; ++c)
{
    if (data[c] >= 128)
        sum += data[c];
}

double elapsedTime =
    static_cast<double>(clock() - start);
```

Next time

- Into to caches


Advanced Pipelining

Not on the exam.

Creating more pipelines

- ❑ Instruction Level Parallelism – Superscalar Pipeline
 - Have two or more pipelines in same processor
 - pipelines need to work in tandem to improve single program performance
- ❑ Thread Level Parallelism – Multi-core
 - Have two or more processors (Independent Pipelines)
 - Need more programs or a parallel program
 - does not improve single program performance
- ❑ Data Level Parallelism – Single Inst. Multiple Data (SIMD)
 - Have two or more execution pipelines (ID->WB)
 - Share the same fetch and control pipeline to save power (IF+cont.)
 - Similar to GPU's

ILP Techniques: Superscalar




Other Techniques for ILP: Out of Order Execution

- ❑ Eliminating stall conditions decreases CPI
- ❑ Reorder instructions to avoid stalls
- ❑ Example (5-stage LC2K pipeline):

add	1	2	3
lw	3	2	16
nor	2	6	7
add	4	5	1

add	1	2	3
lw	3	2	16
add	4	5	1
nor	2	6	7




Why Use Out of Order Execution?


- ❑ Some instructions take a long time to execute
 - Floating point operations
 - Some loads and stores (more when we talk about memory hierarchy)
- ❑ Options:
 - Increase cycle time
 - Increase number of pipeline stages
 - Execute other instructions while you wait

Advanced Pipelining

TLP Techniques: Multiprocessors




Other Techniques for TLP: Multi-Threading




- ❑ Virtual Multiprocessor (Multi-Threading or HyperThreading)
 - Duplicate the state (PC, Registers) but time share hardware
 - User/Operating system see 2 cores, but only one execution
 - Used to hide long latencies (i.e. memory access to disk)

DLP Techniques: Single Instr. Multiple Data (SIMD)



Building a GPU



- Add special functional units in EX
- Combine Techniques
 - SIMD + MP + MT = SIMT
- MT used to hide memory latencies
- SIMD used to decrease power of fetch/control
- MP used to improve throughput