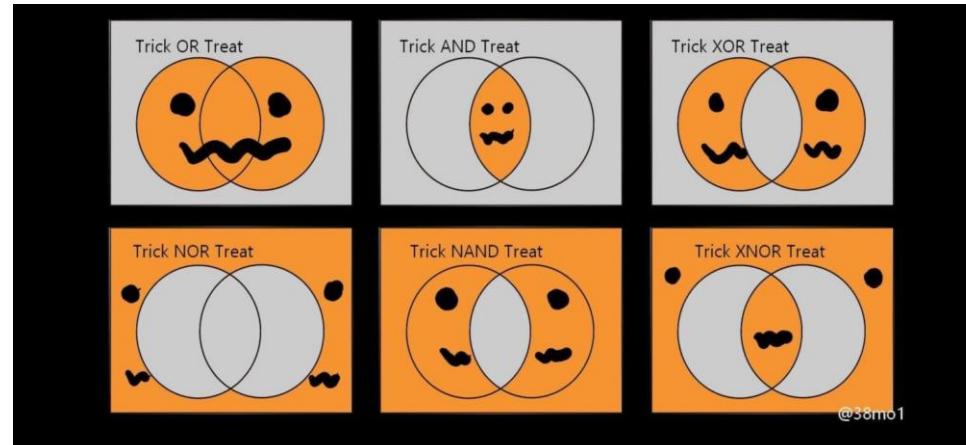


EECS 370

Improving Caches



Live Poll + Q&A: [slido.com #eeecs370](https://slido.com/#eeecs370)

Announcements

- Lab
 - Gradescope quiz due tonight
- P3 checkpoint due next Thursday
 - Get pipeline simulator working without hazards
- My Monday OH next week are cancelled
 - I'll add extra later in the week



Memory Hierarchy

- Key observation: we only need to access a small amount of data at a time
- Let's use a small array of SRAM to hold data we need now
 - Call this the **cache**
 - Able to keep up with processor
 - Small (~Kilobytes), so it should be relatively cheap
- Use a large amount of DRAM for **main memory**
 - Can scale up to ~Gigabytes in size



Scaling Up

- What if we access many memory locations, and cache isn't large enough to hold all of them?
- How do we choose what to keep in the cache?
 - How does hardware predict what's most likely to be needed soon?
- Answer: **locality**
 - Temporal locality
 - Spatial locality



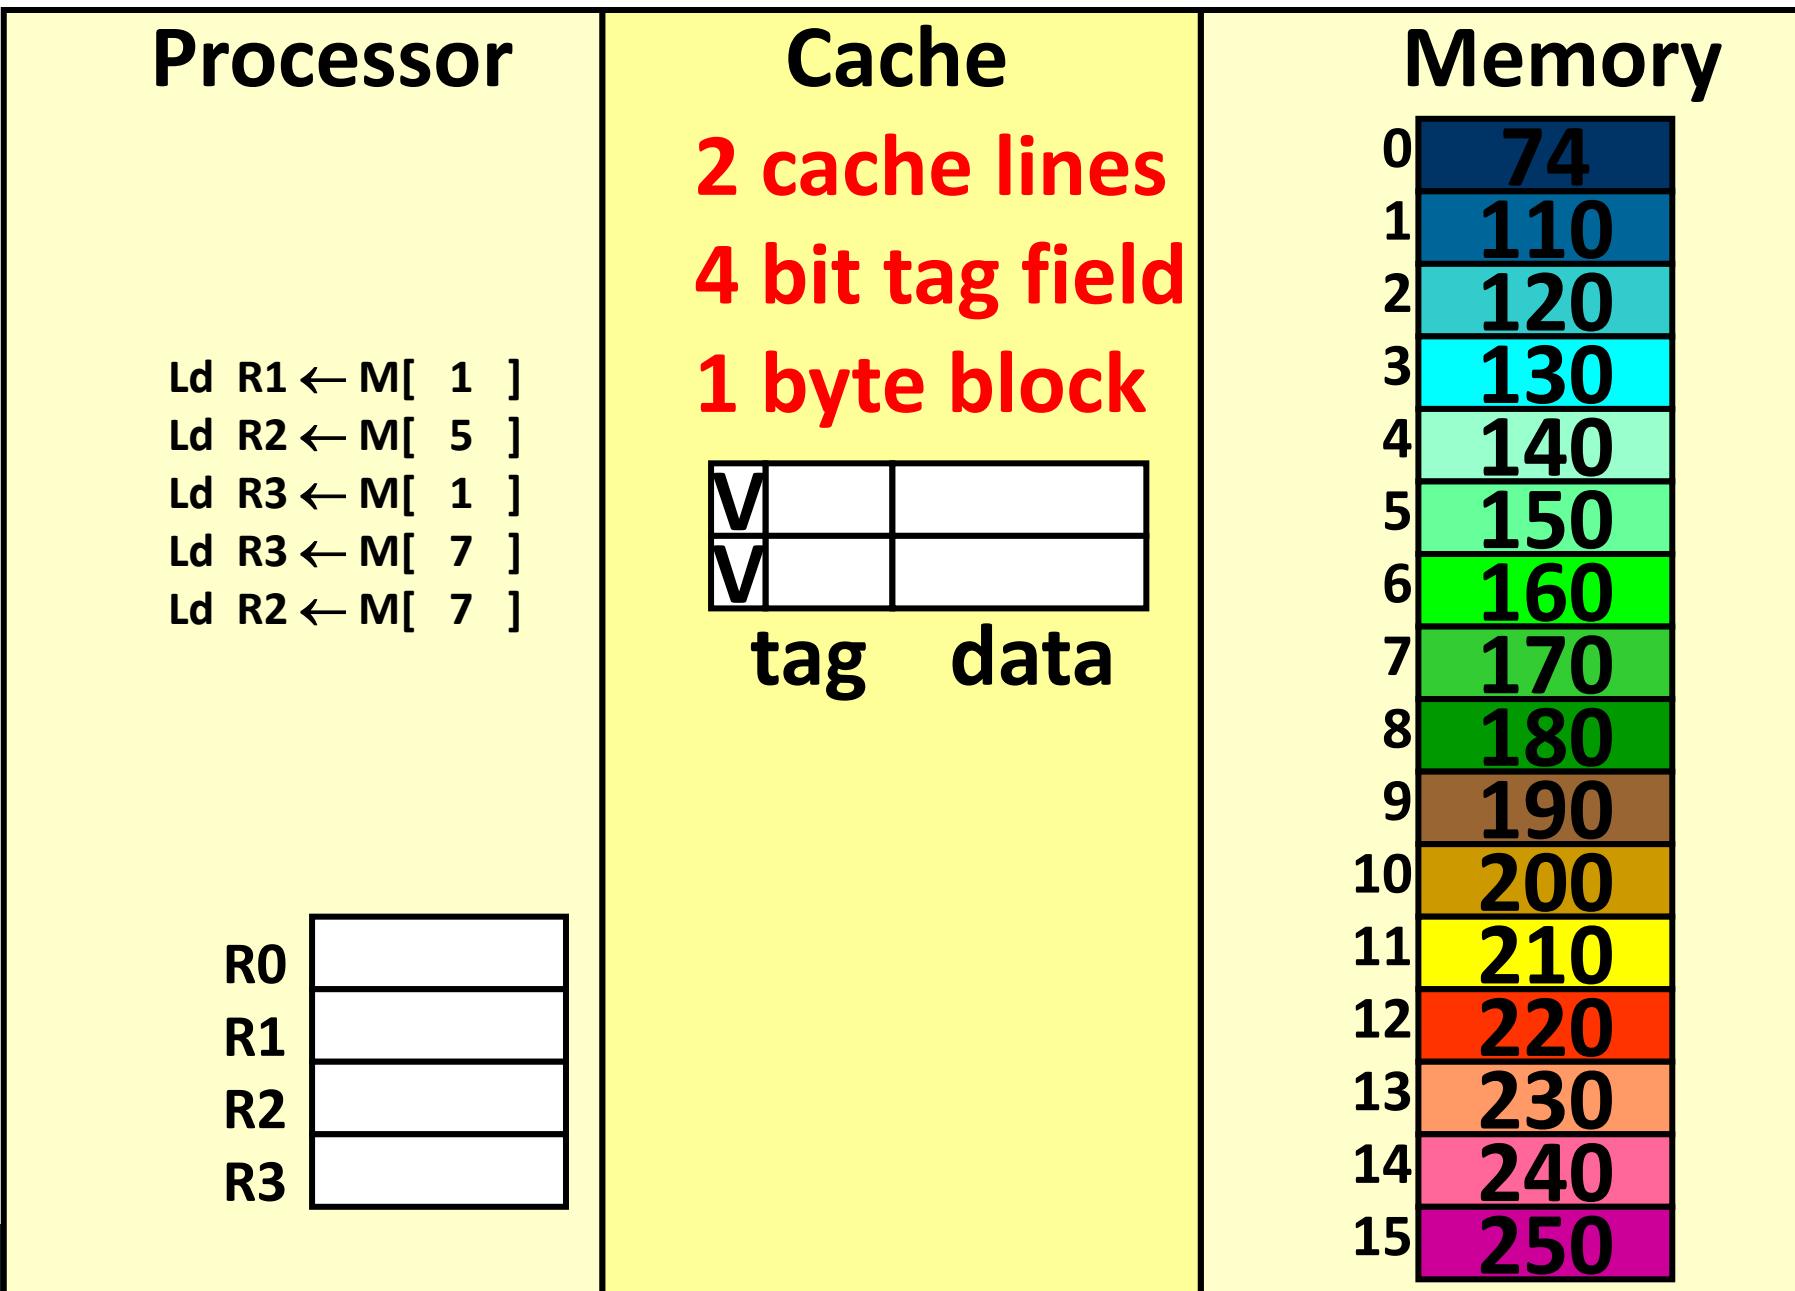
Temporal Locality

- Temporal locality: if a given memory location is referenced now, it will probably be used again in the near future
 - Why? Take a look at code you've written. You tend to use a variable multiple times
 - Corollary: if you haven't used a variable in a while, you probably won't need it very soon either
- Hardware should take advantage of this by:
 - Placing items we just accessed in the cache
 - When we need to evict something, evict whatever data was **least recently used (LRU)**

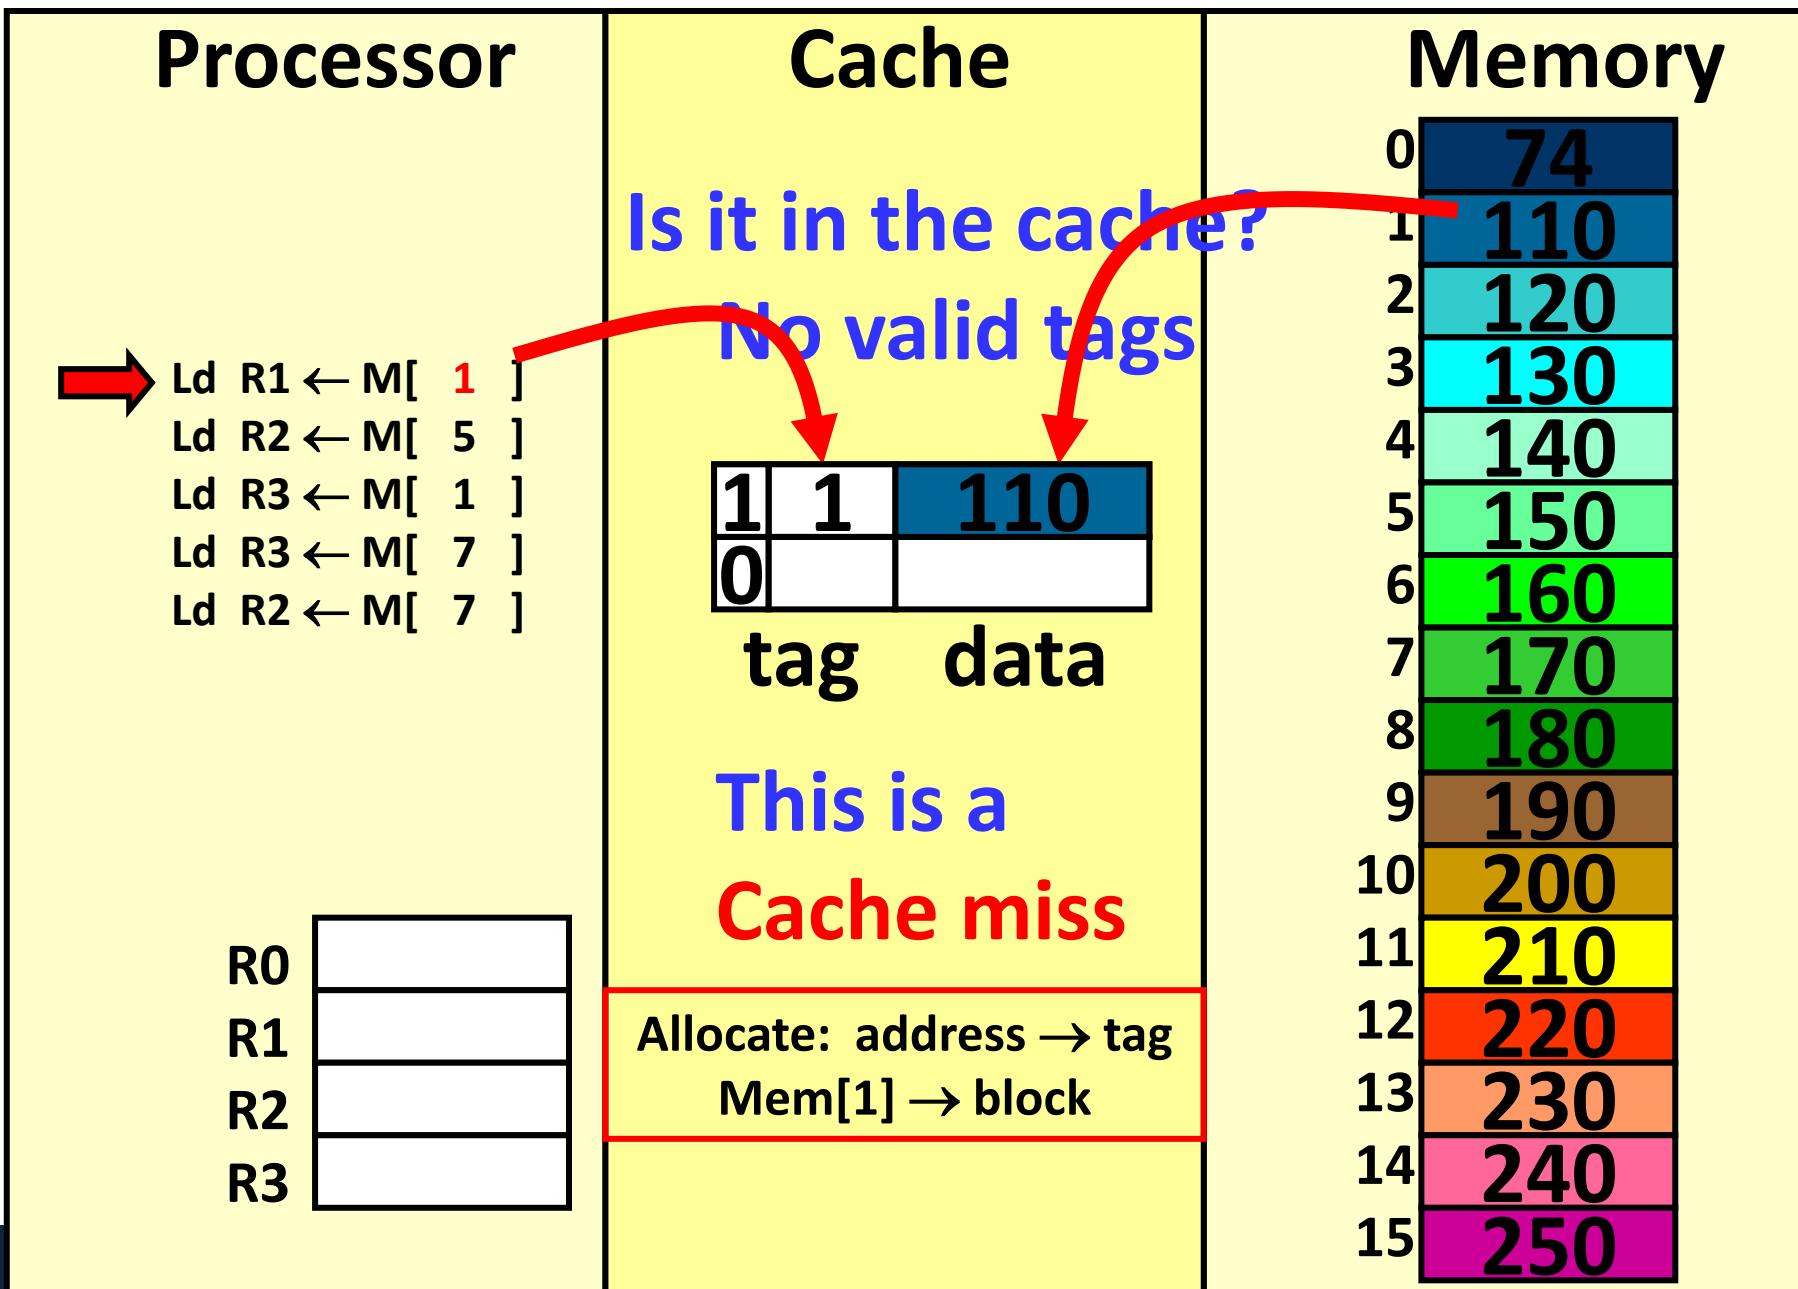
Definitions

- **Hit:** when data for a memory access is found in the cache
- **Miss:** when data for a memory access is not found in the cache
- **Hit/Miss rate:** percentage of memory accesses that hit/miss in the cache

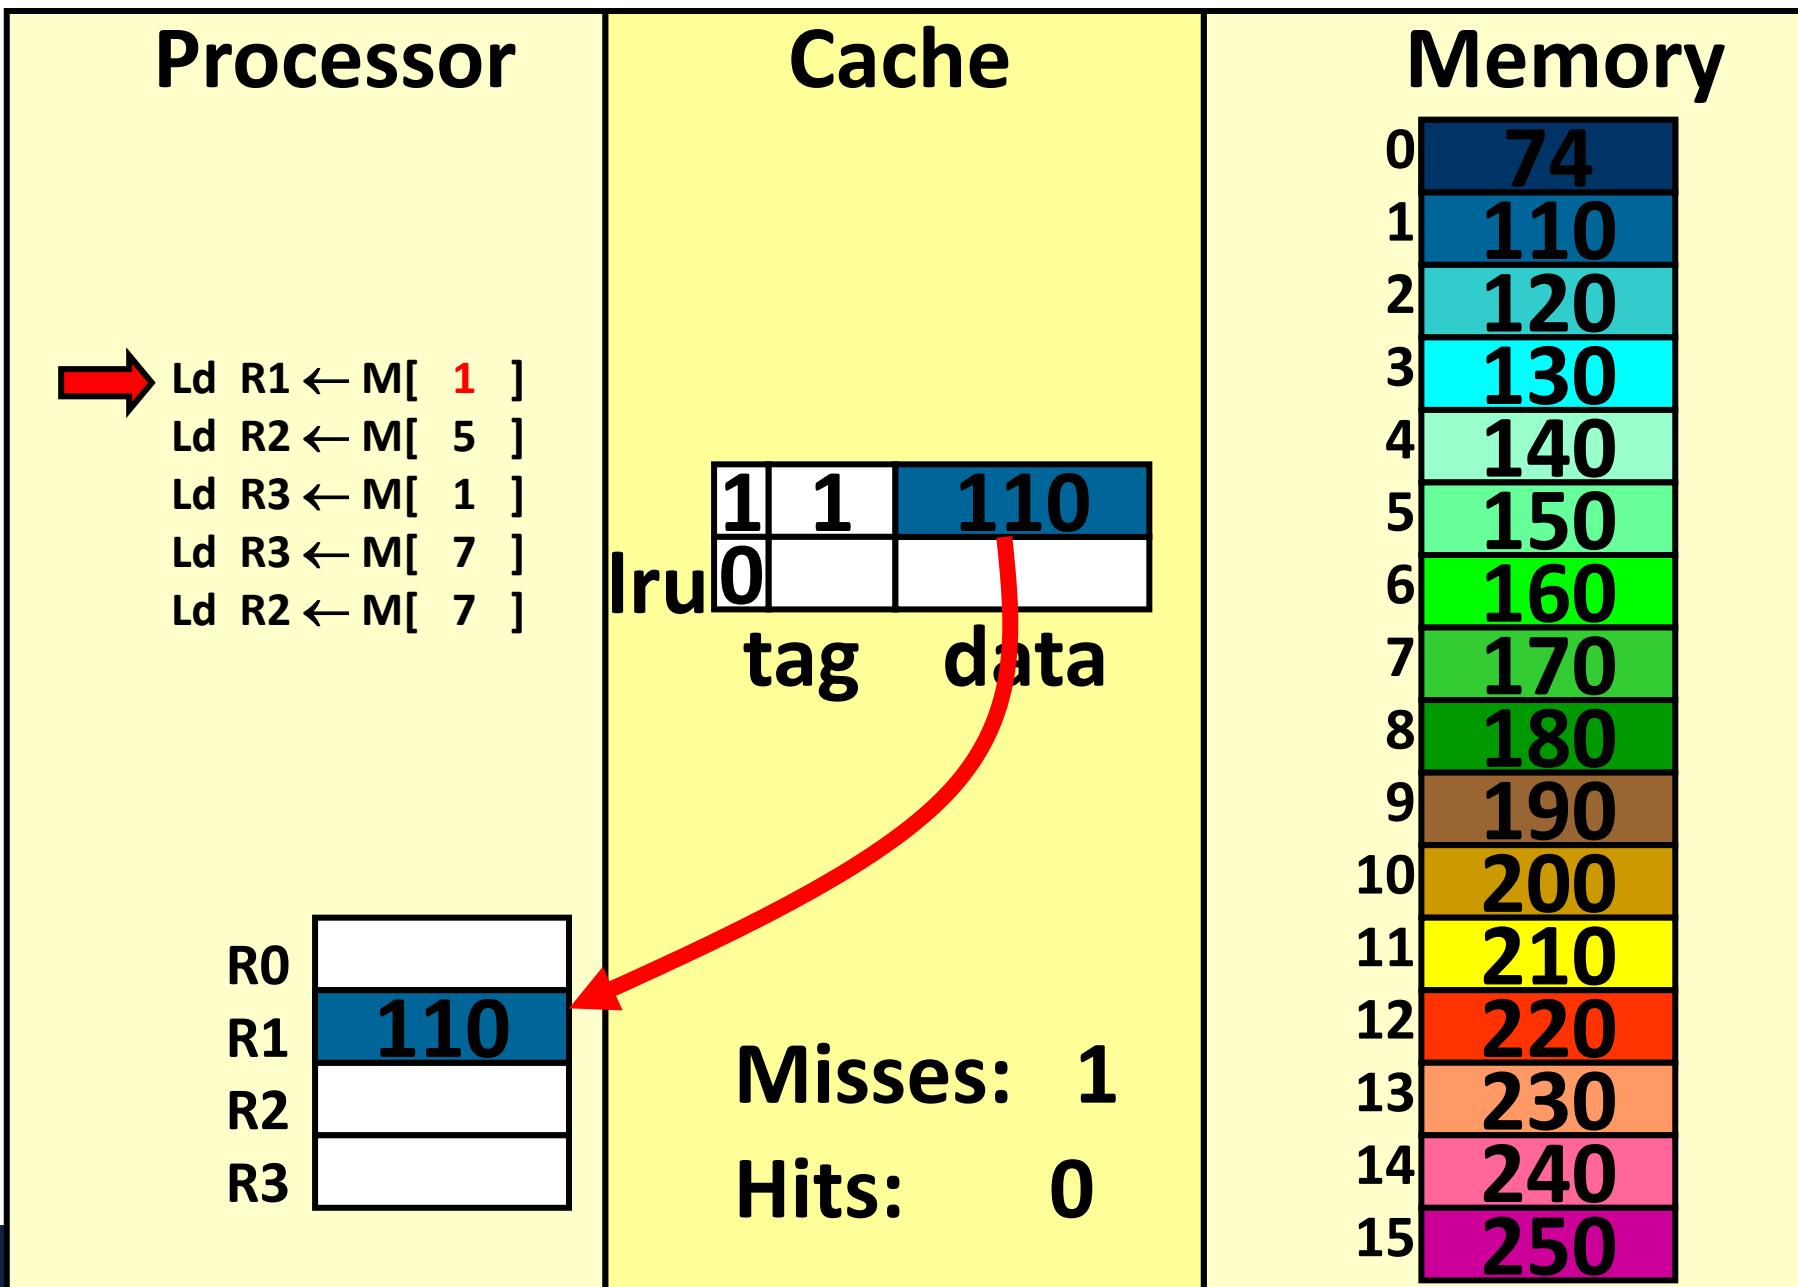
A Very Simple Memory System



A Very Simple Memory System

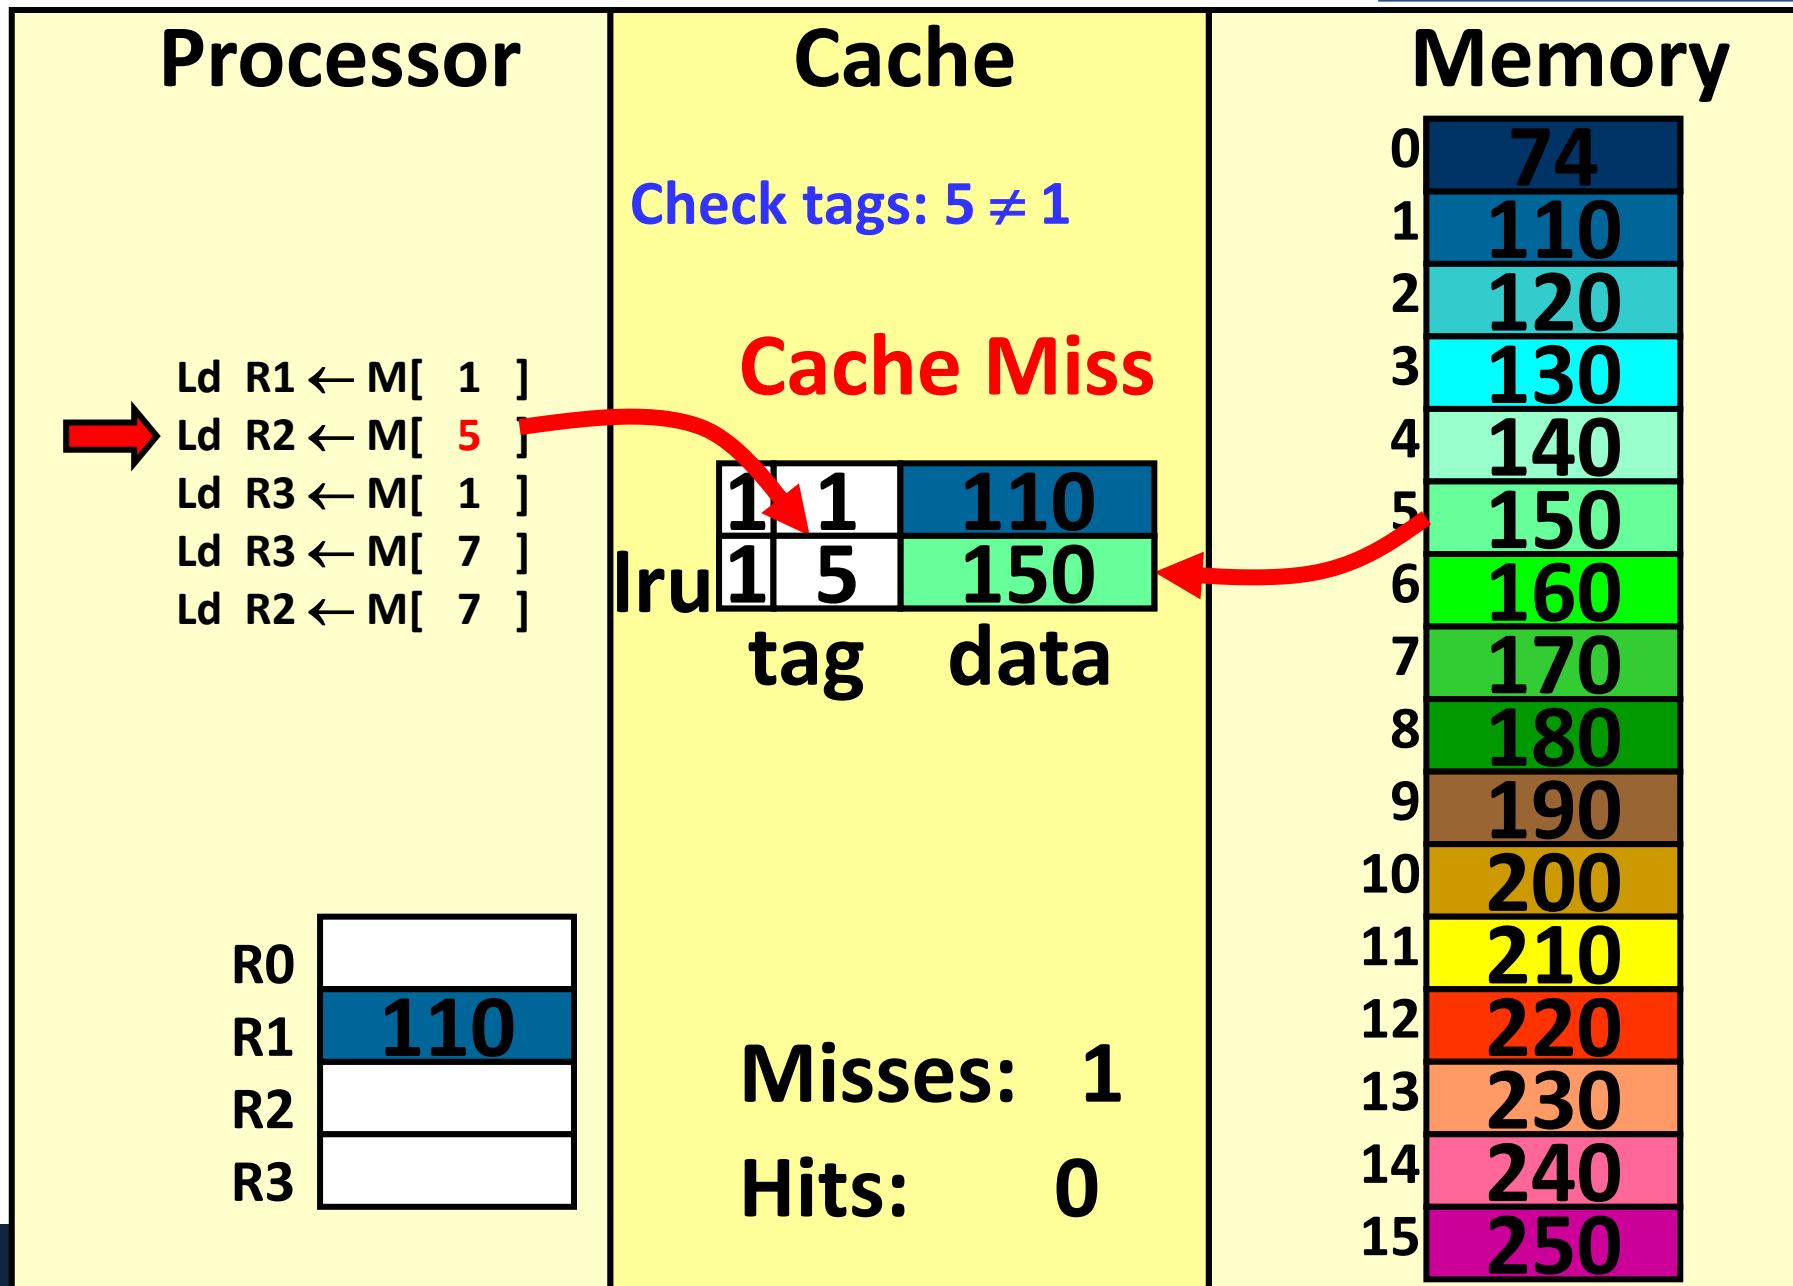


A Very Simple Memory System

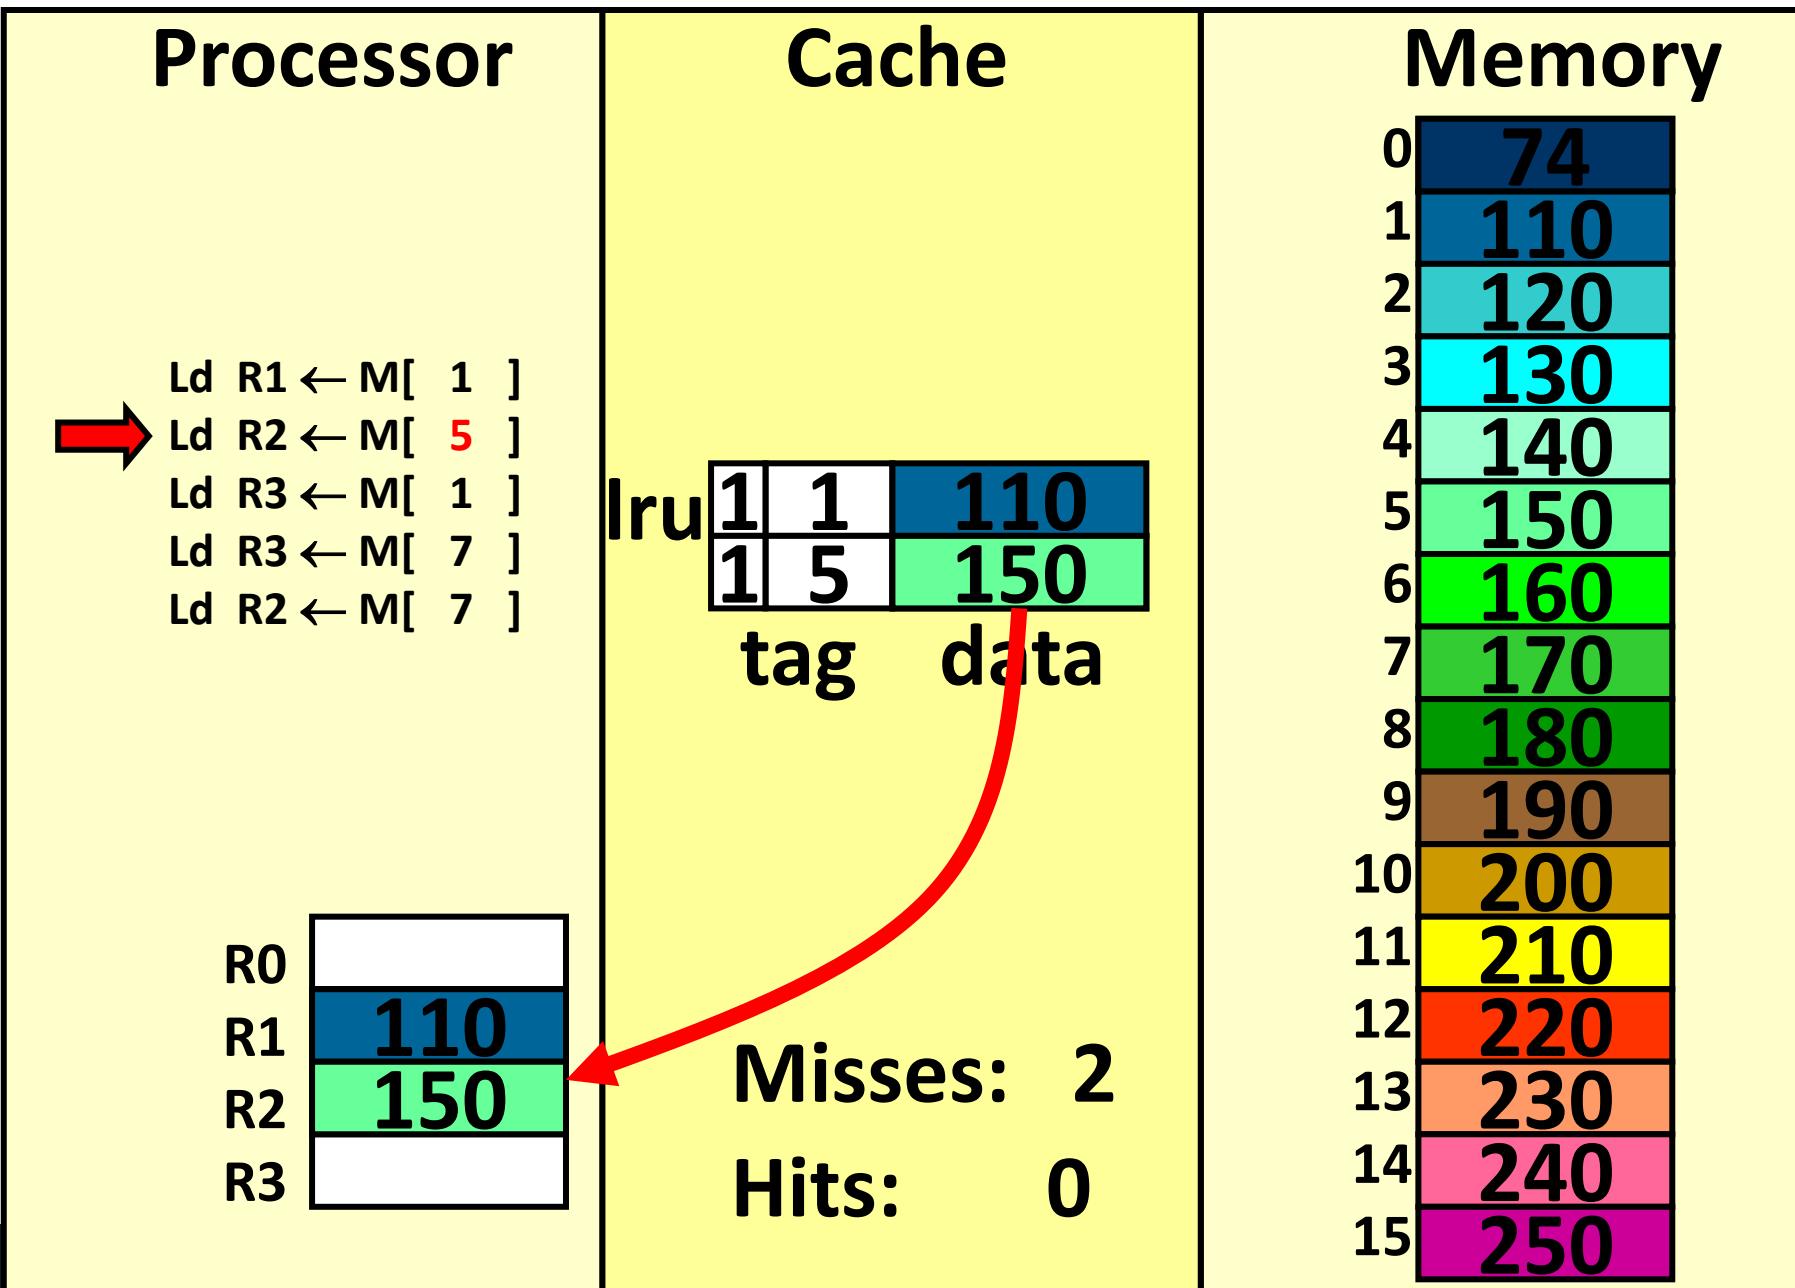


A Very Simple Memory System

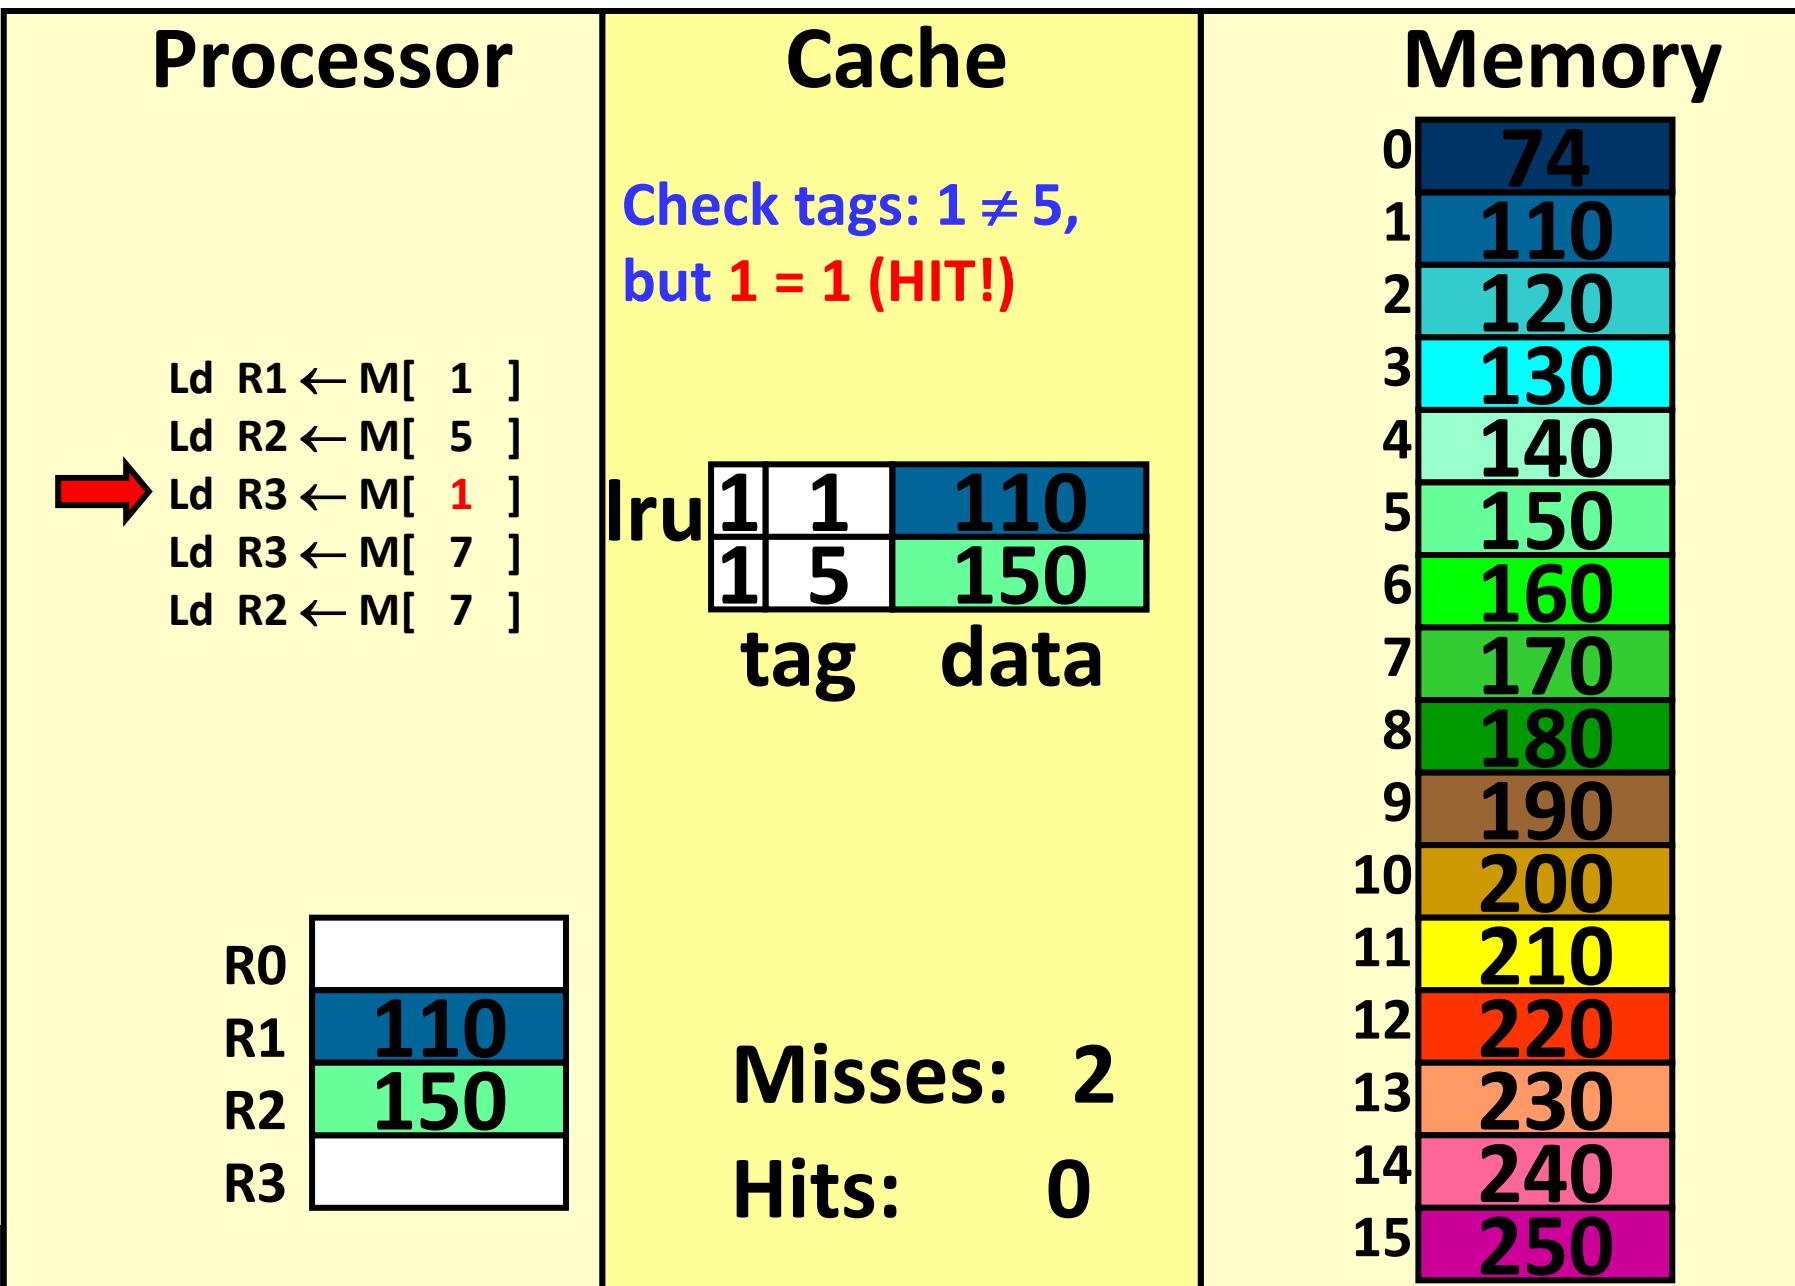
Tag comparison uses hardware called "content-addressable memory (CAM)"



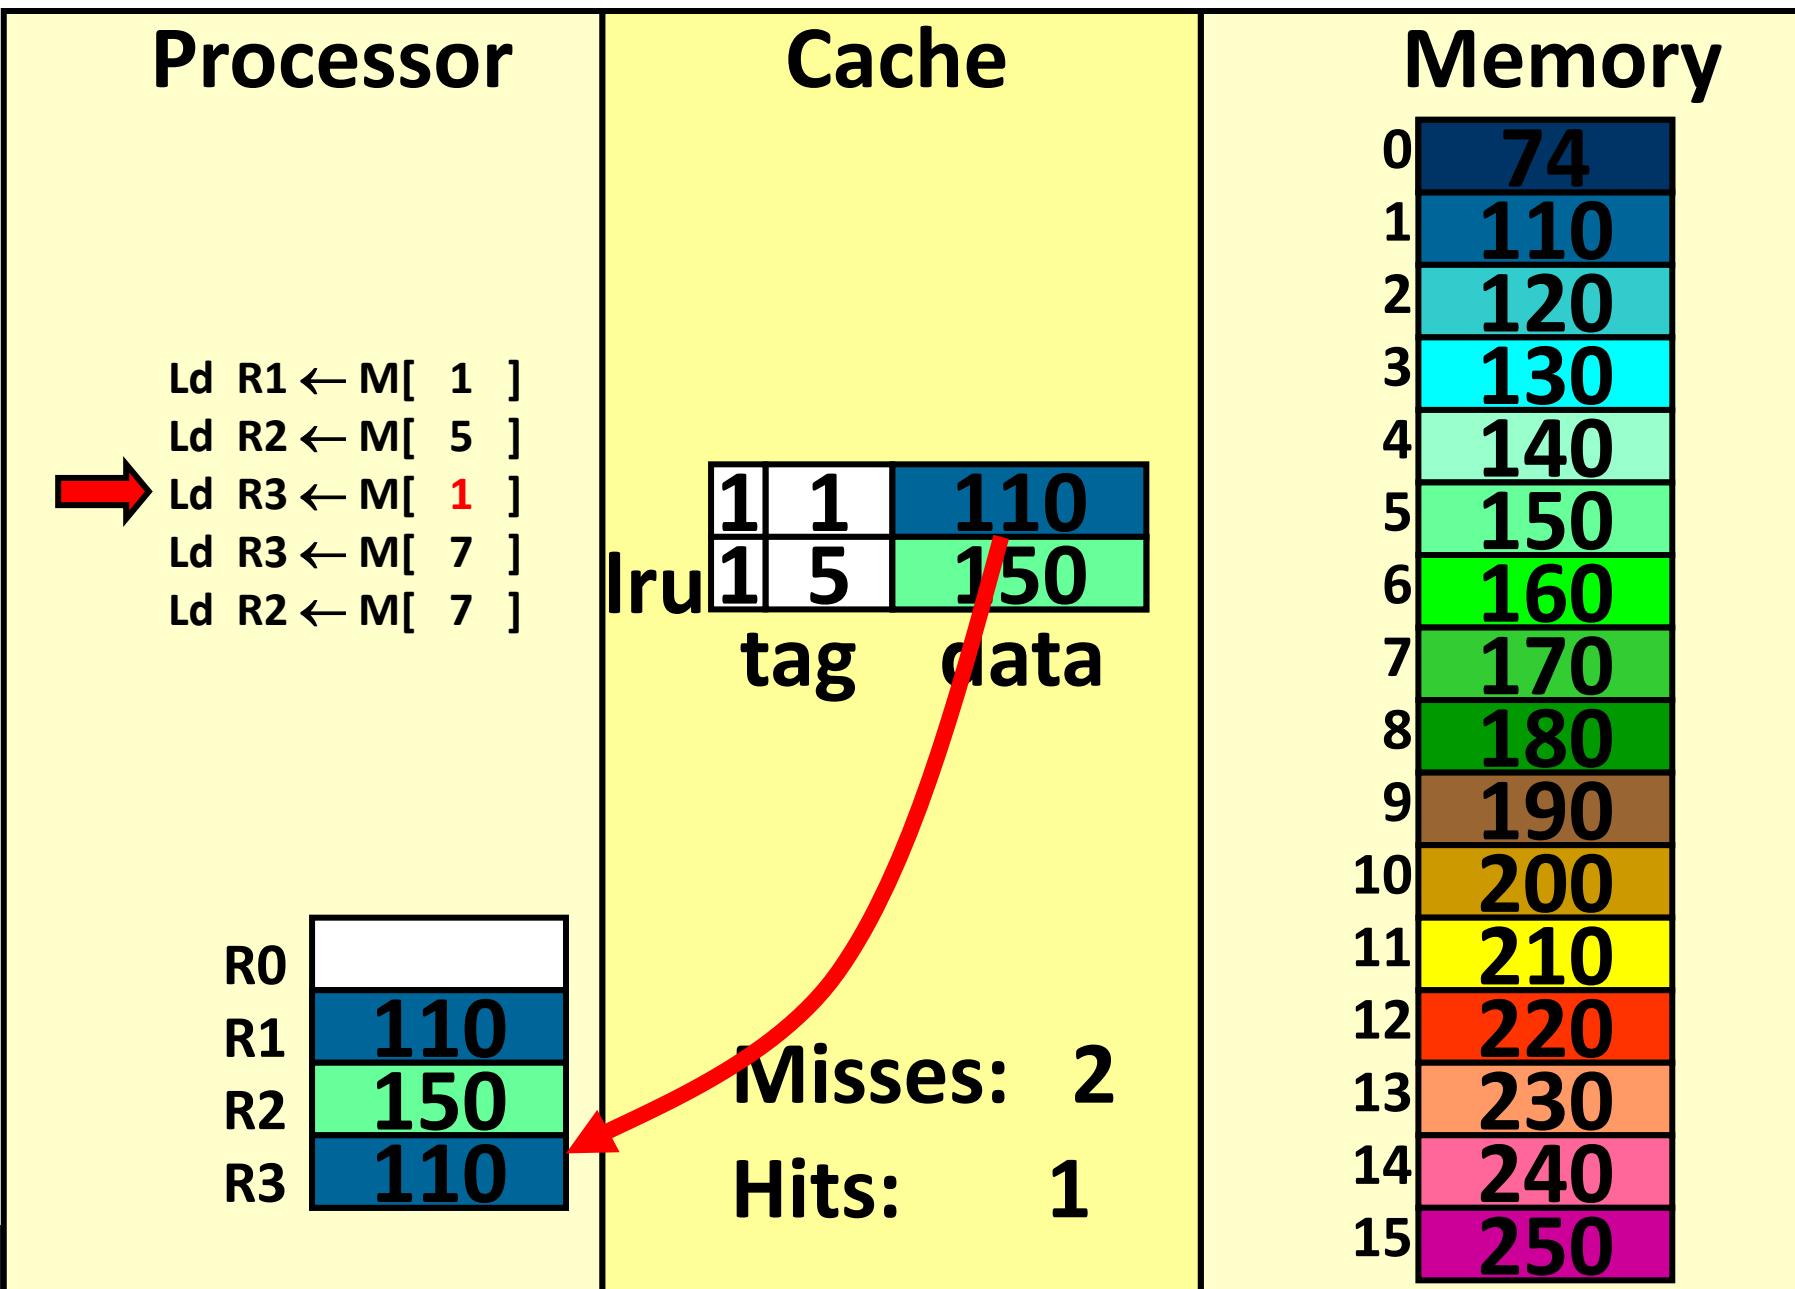
A Very Simple Memory System



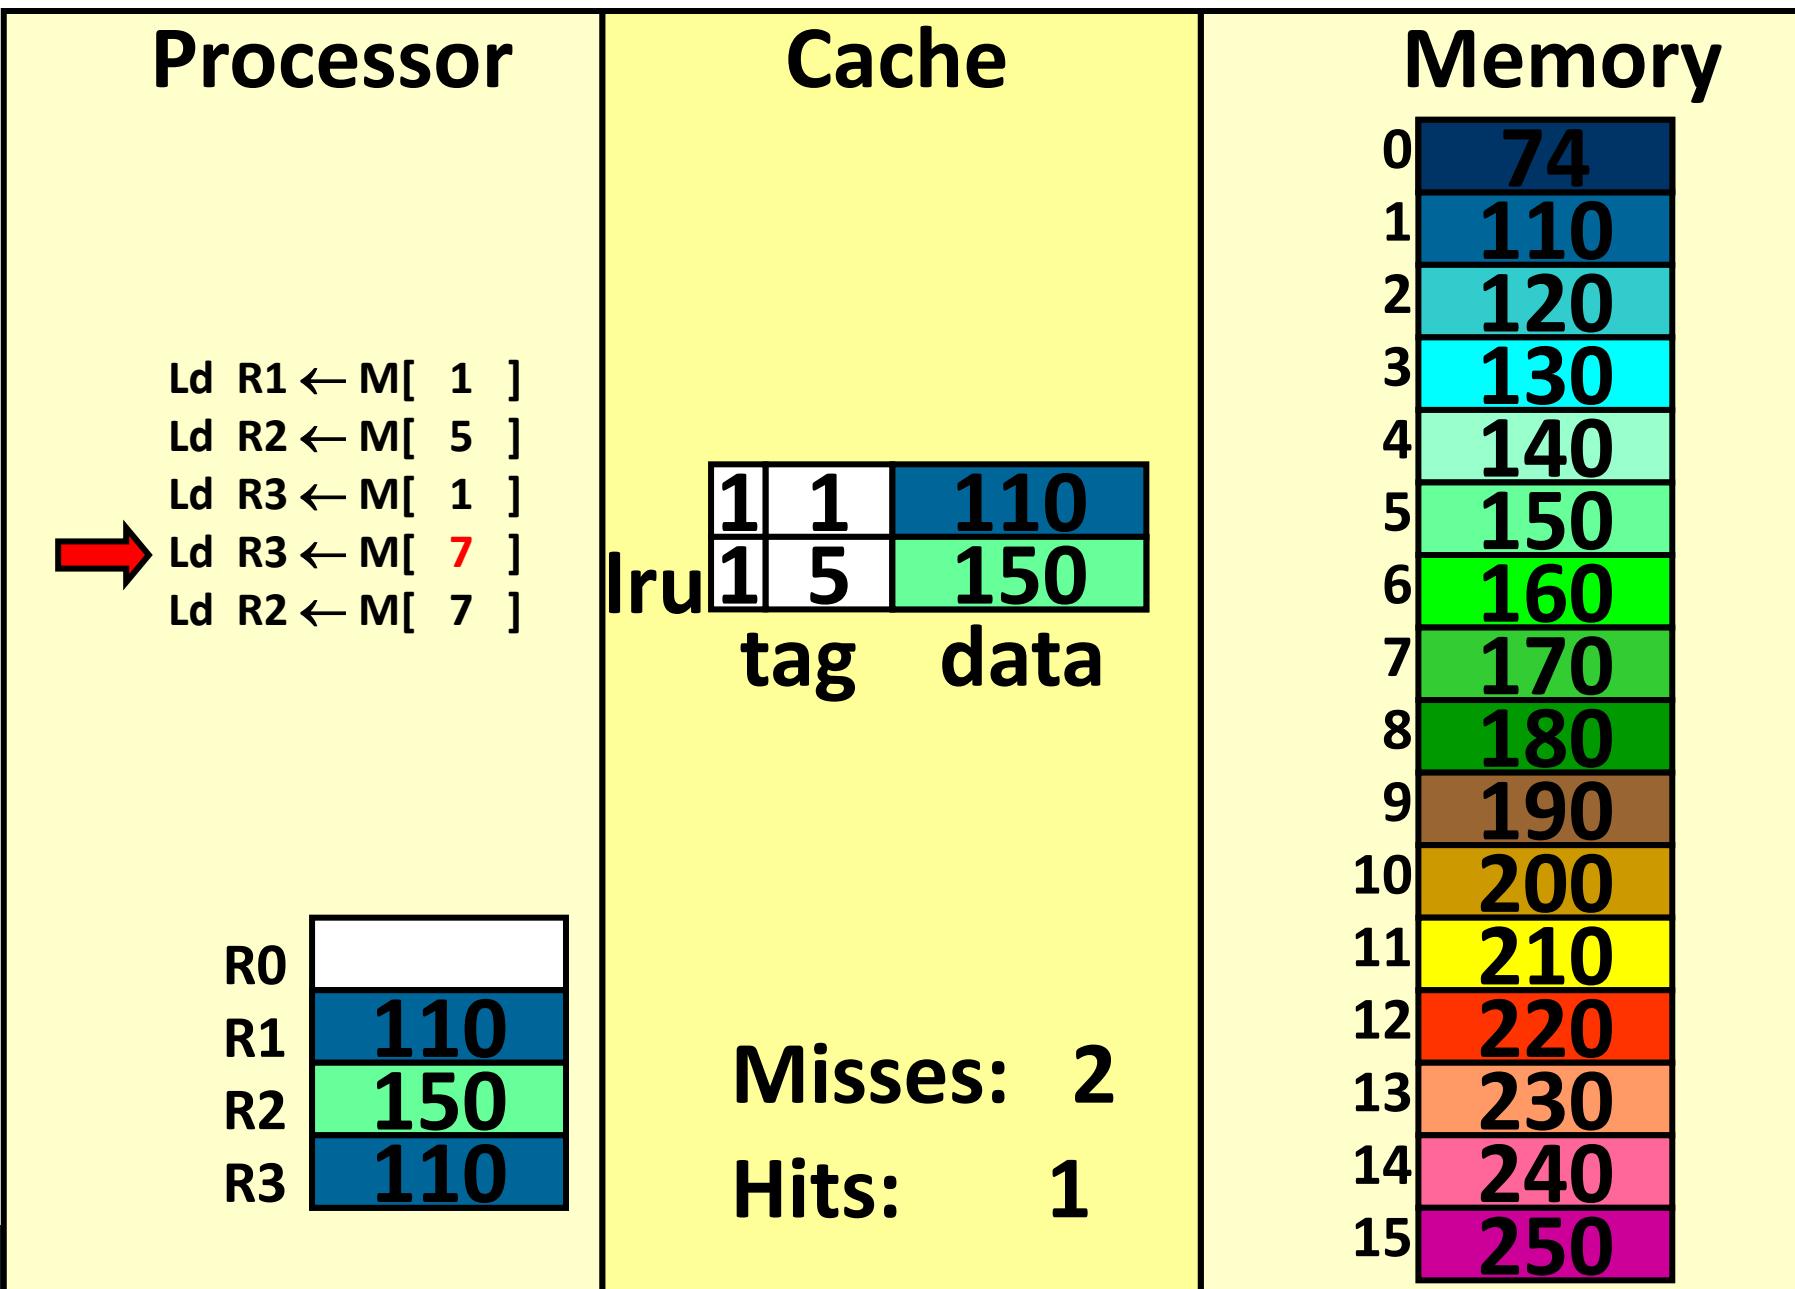
A Very Simple Memory System



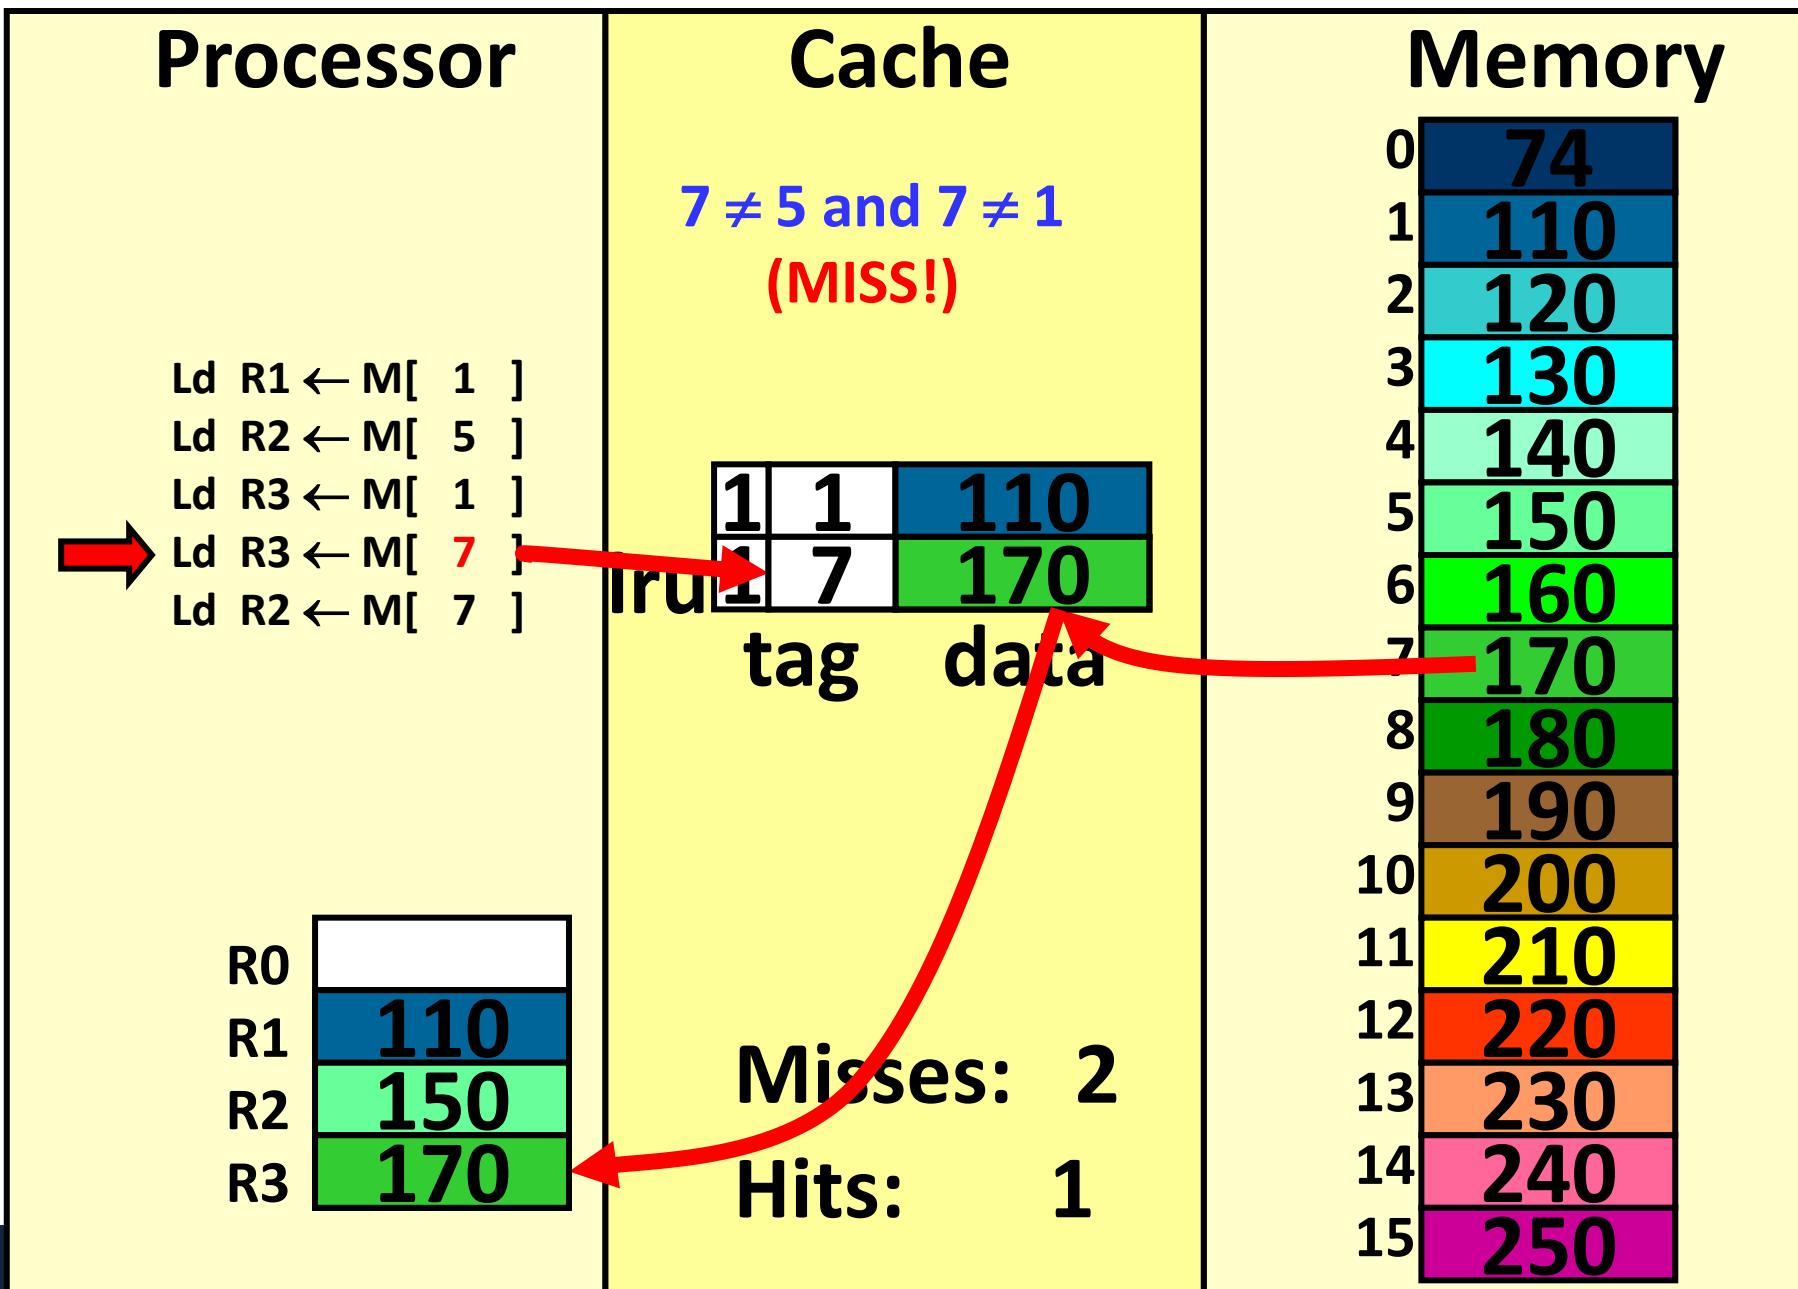
A Very Simple Memory System



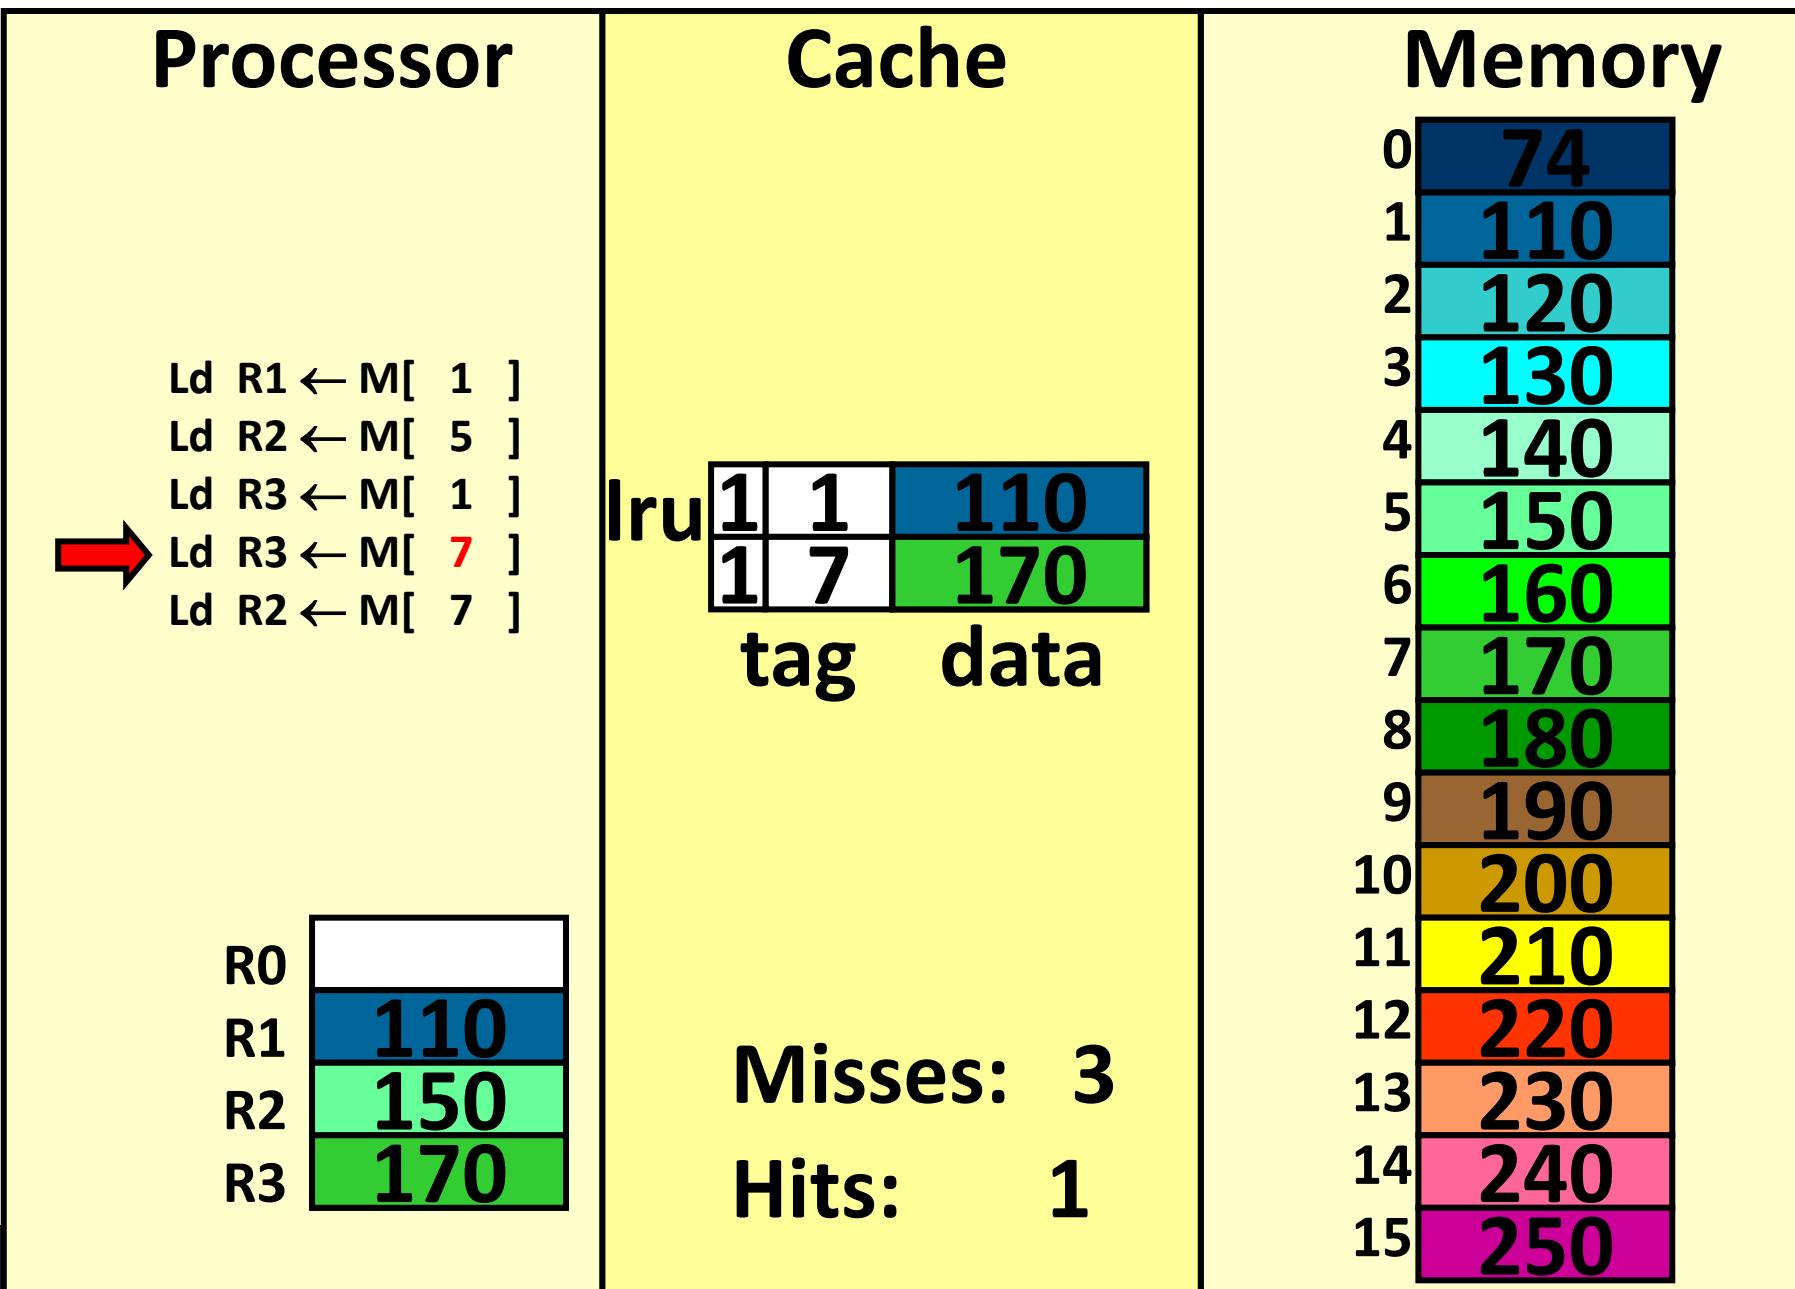
A Very Simple Memory System



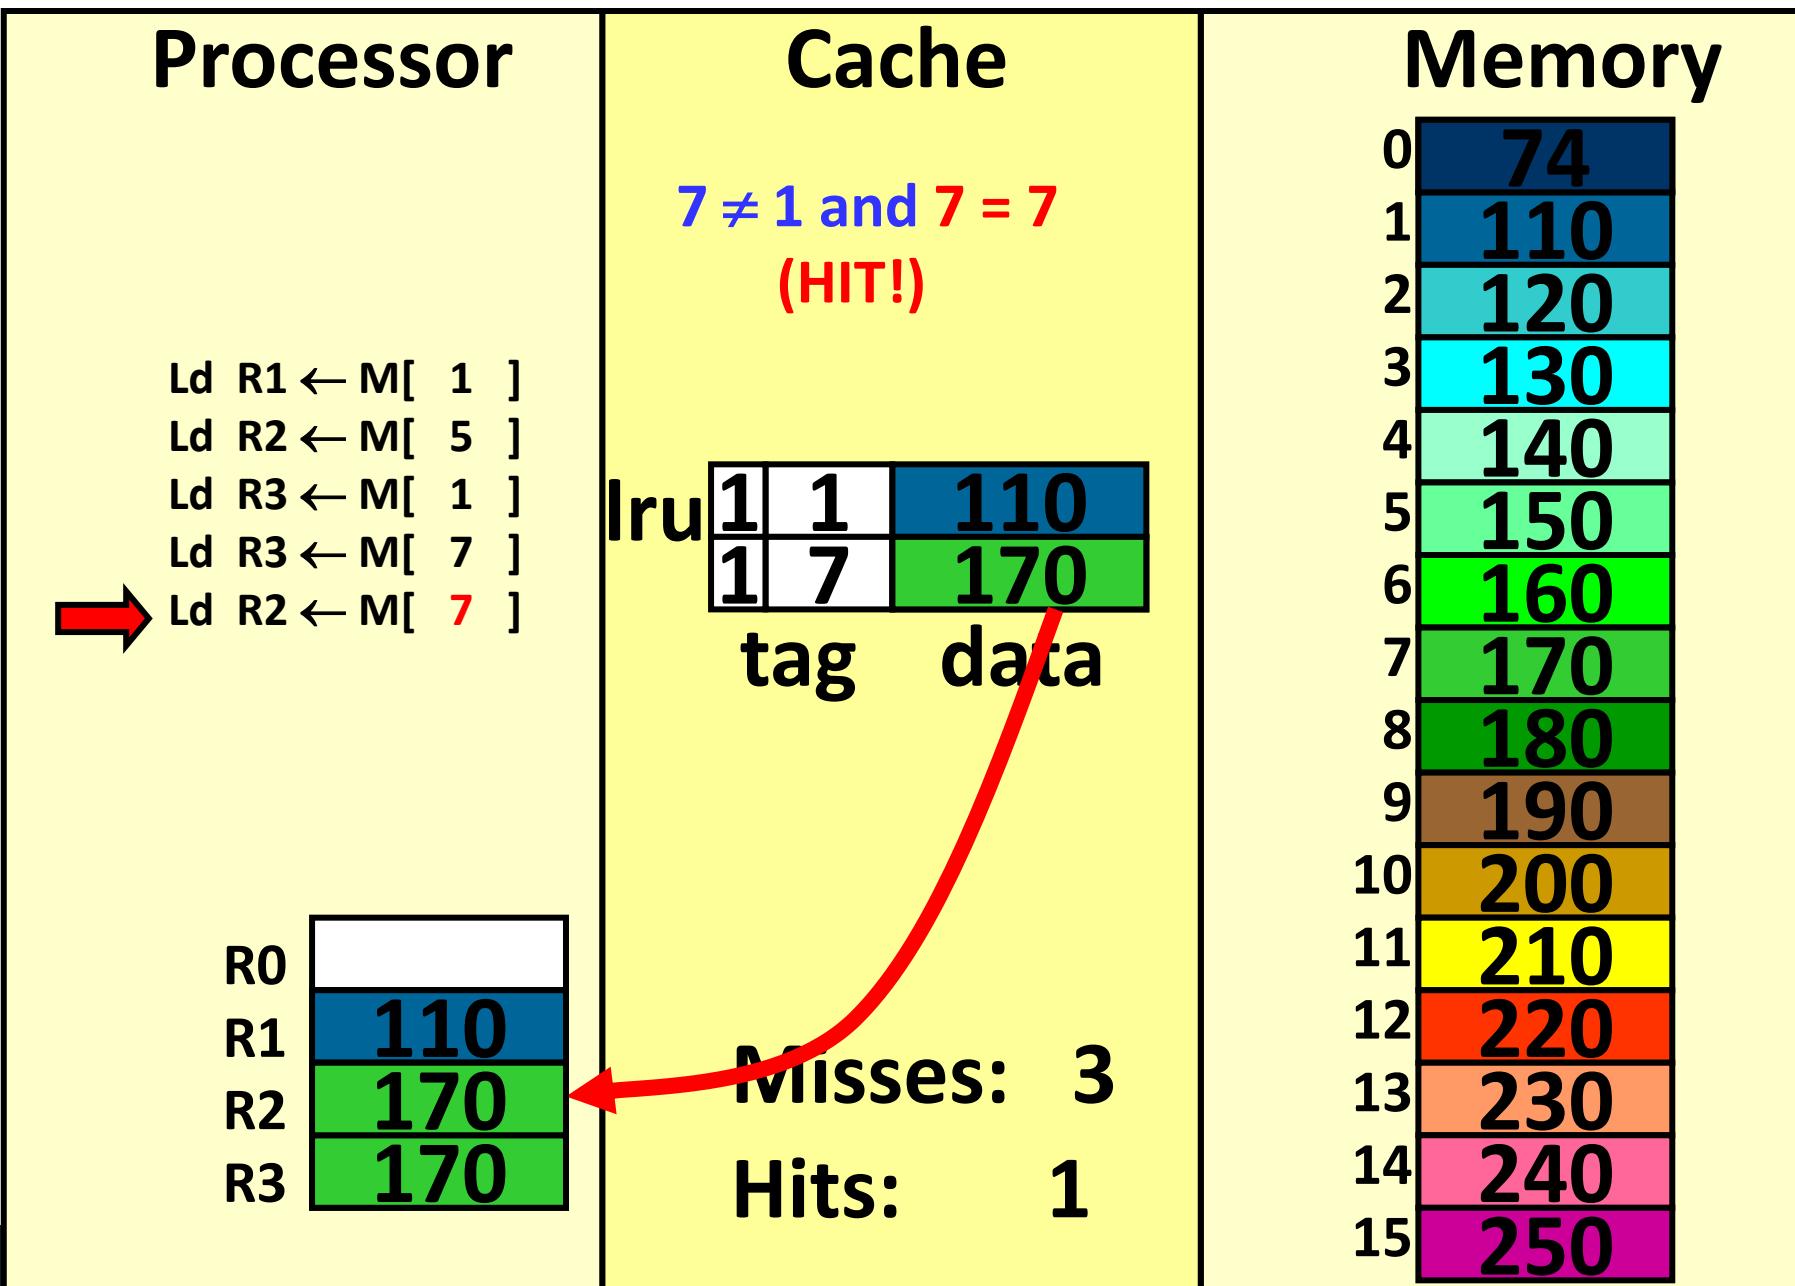
A Very Simple Memory System



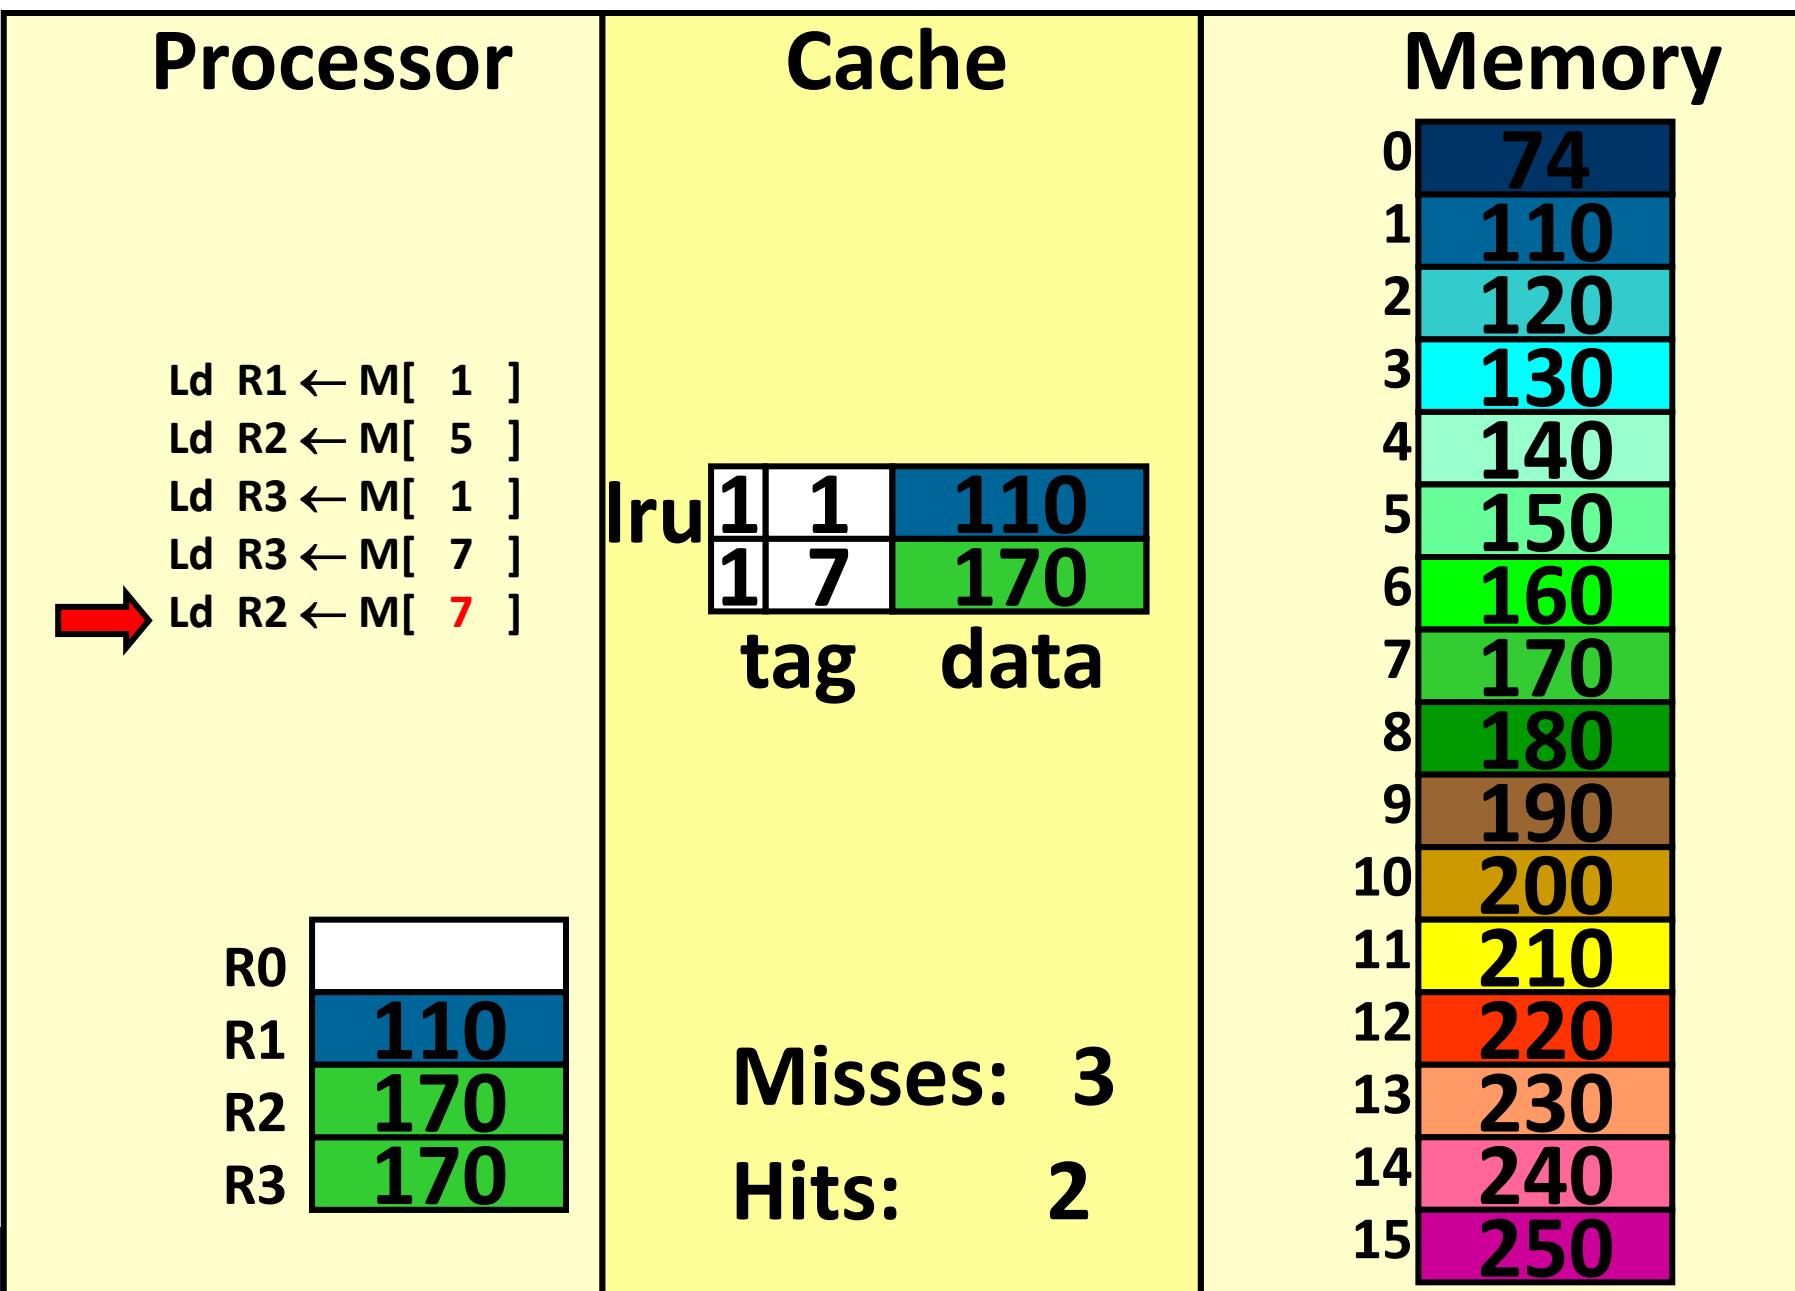
A Very Simple Memory System



A Very Simple Memory System



A Very Simple Memory System



Poll: Average access time?

Example Problem

- Assume the following:
 - Cache has 1 cycle access time
 - If data is not found in cache, main memory is then accessed instead
 - Main memory has 100 cycle access time
- If we have a 90% hit rate in the cache, what is the average memory latency?

$$1 + 0.1 * (100) = 11$$

Calculating Average Access Latency

- Average Latency:
 - $cache\ latency + (memory\ latency \cdot miss\ rate)$
- To improve latency, either:
 - Improve memory access latency, or
 - Improve cache access latency, or
 - Improve cache hit rate



Agenda

- Larger Cache Blocks
- Extra Problems
- LRU with More than Two Blocks
- Write-Through Cache
- Write-Back Cache



Calculating Size

- How many bits is used in cache?
 - Storing data
 - 2 bytes of SRAM
 - Calculate overhead (non-data)
 - This cost is often forgotten for caches, but it drives up the cost of real designs!
 - 2 4-bit tags
 - 2 valid bits
- What is the storage requirement

Poll: Which of the following would reduce tag overhead (as an overall percentage)? (select all that apply)

- a) Increase number of cache entries
- b) Decrease number of cache entries
- c) Use smaller addresses
- d) Store more data in each cache entry



How can we reduce overhead?

- Have a smaller address
 - Impractical, and caches are supposed to be micro-architectural
- Cache bigger units than bytes
 - Each block has a single tag, and blocks can be whatever size we choose.



Increasing Block Size

Case 1:

Block size: 1 bytes

1	0	74
1	6	160

V tag data (block)

How many bits needed per tag?

$$= \log_2(\text{number of blocks in memory}) = \log_2(16)$$

= 4 bits

$$\text{Overhead} = (4+1) / 8 = 62.5\%$$

Case 2:

Block size: 2 bytes

1	0	74	110
1	3	160	170

V tag data (block)

How many bits needed per tag?

$$= \log_2(\text{number of blocks in memory}) = \log_2(8)$$

= 3 bits

$$\text{Overhead} = (3+1) / 16 = 25\%$$

Memory	Tag (case 1)	Tag (case 2)
0	74	0
1	110	1
2	120	1
3	130	1
4	140	2
5	150	2
6	160	3
7	170	3
8	180	4
9	190	4
10	200	5
11	210	5
12	220	6
13	230	6
14	240	7
15	250	7

Figuring out the tag

- If block size is N, what's the pattern for figuring out the tag from the address?

- $tag = \lfloor \frac{addr}{block\ size} \rfloor$

- If block size is power of 2, then this is just everything except the $\log_2(block\ size)$ bits of the address in binary!

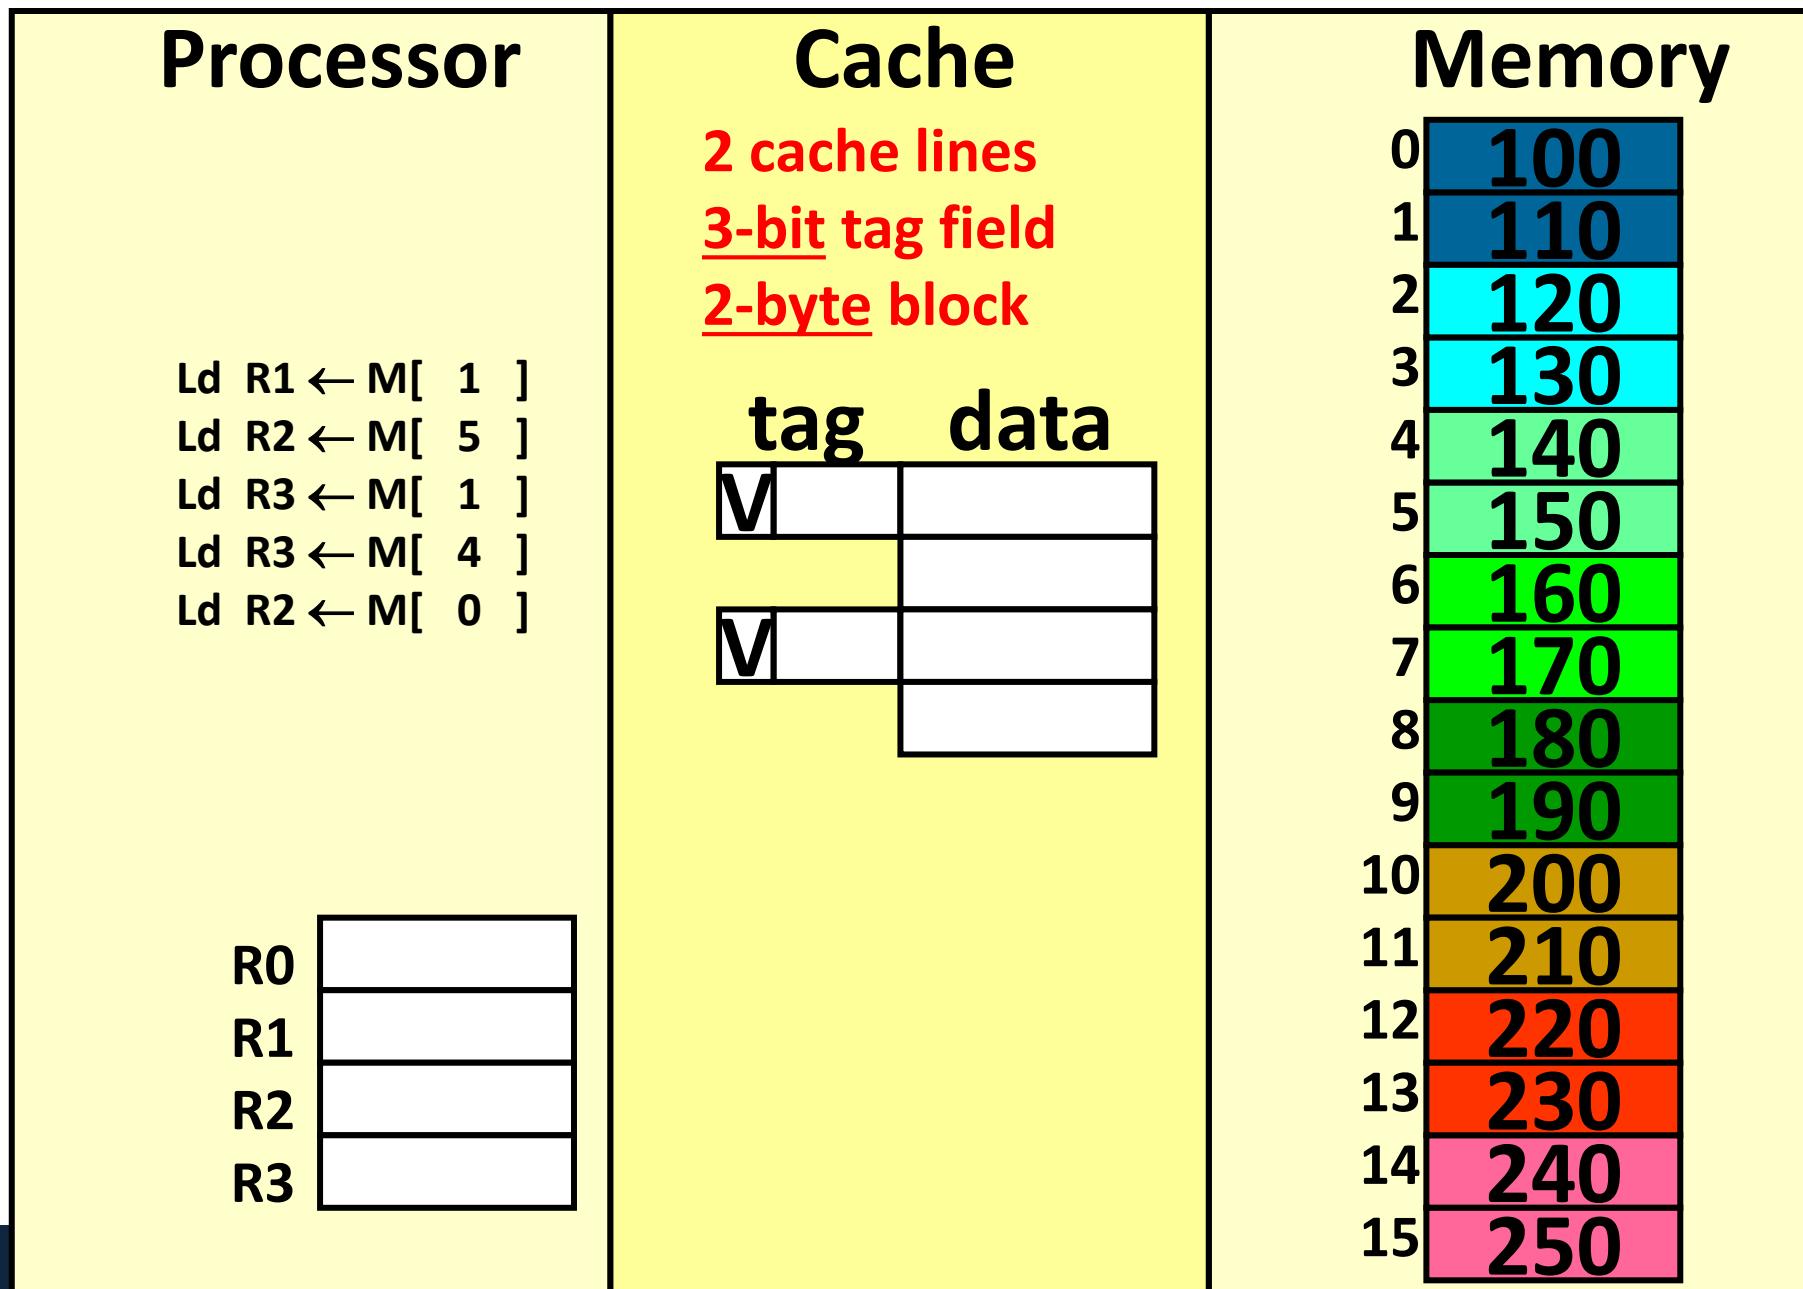
- E.g.

$$\begin{aligned}0d11 &= 0b1011 \\Tag &= 0b101 = 0d5 \\Block\ Offset &= 1\end{aligned}$$

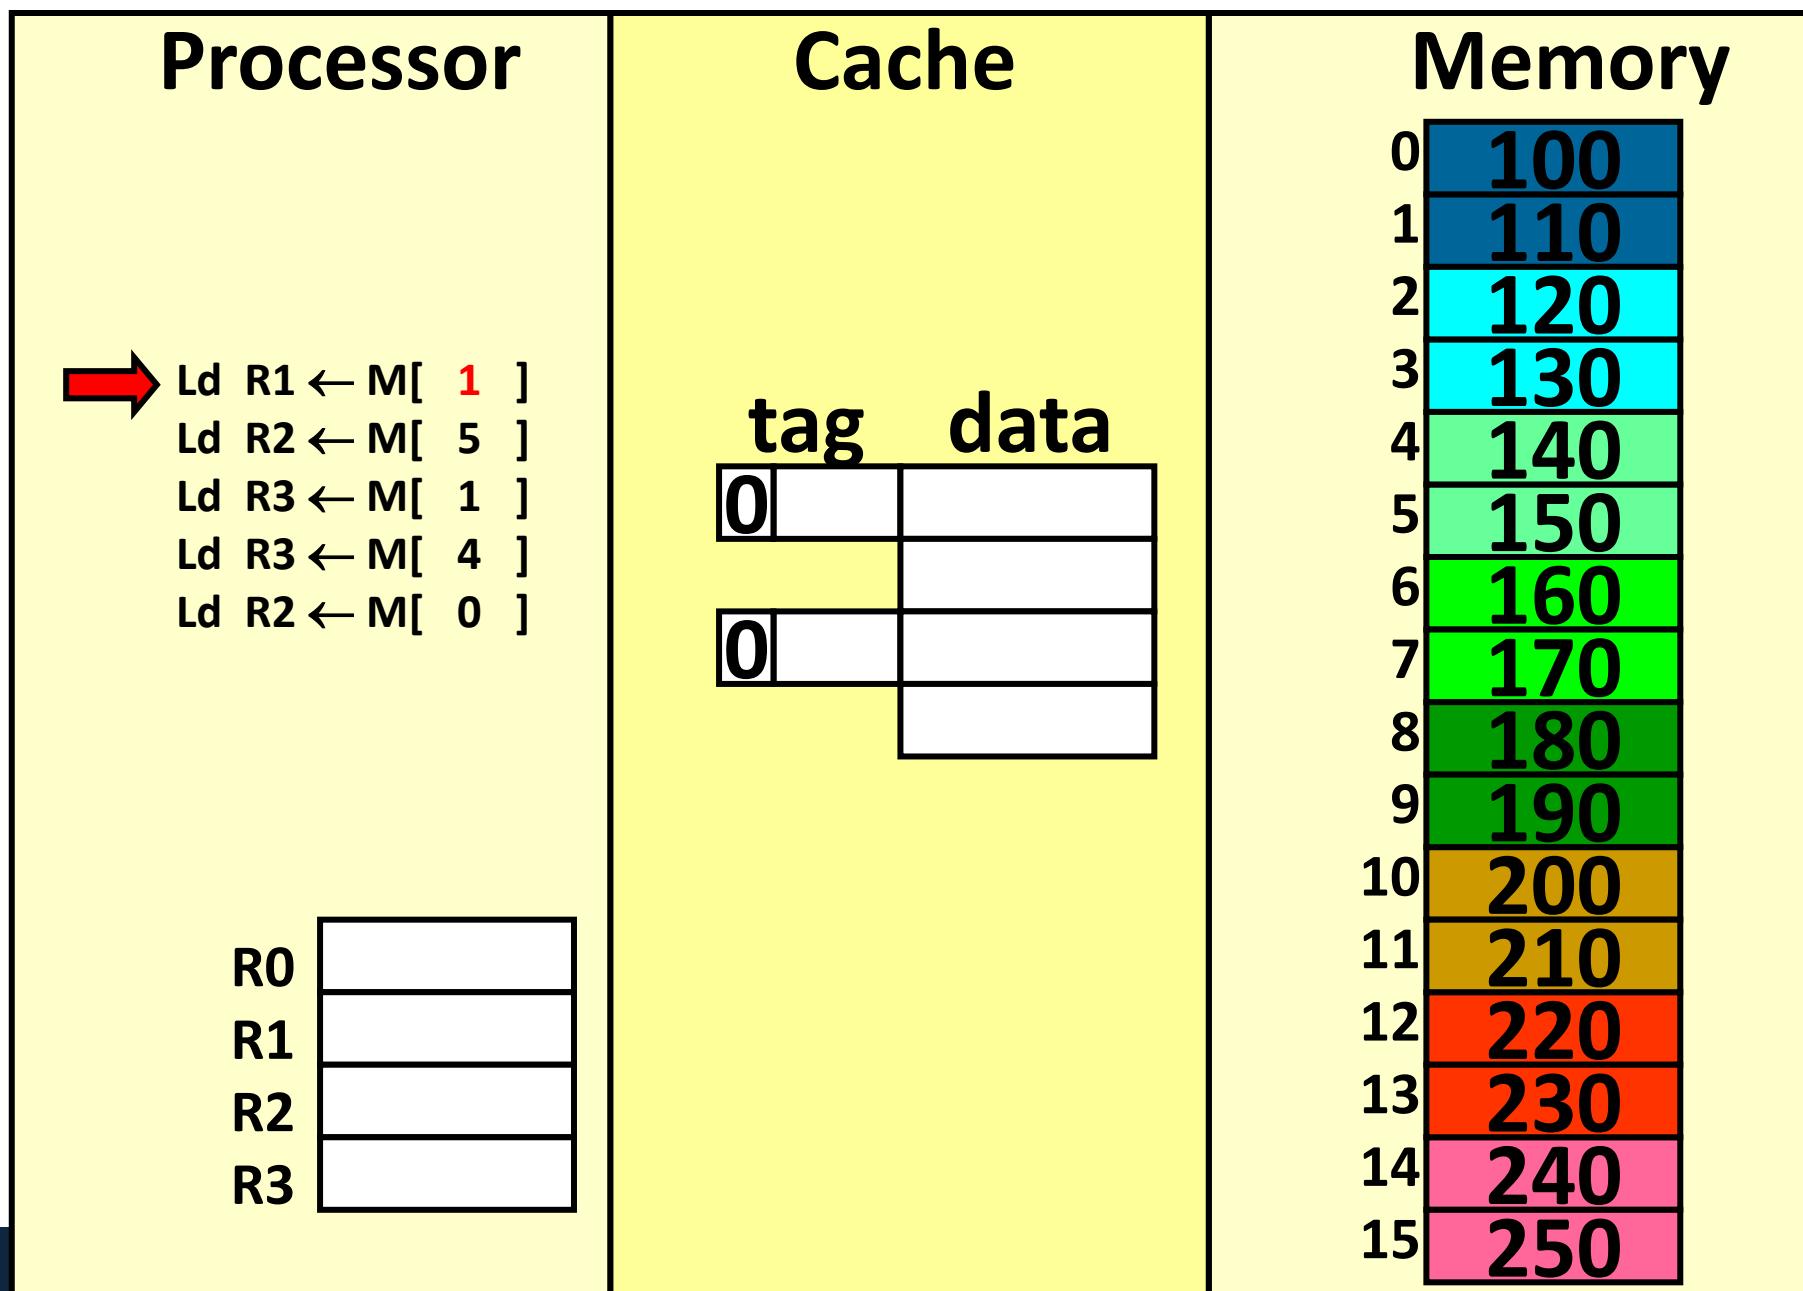
- Remaining bits (block offset) tells us how far into the block the data is

Memory	Tag (case 2)
0 74	0
1 110	0
2 120	1
3 130	1
4 140	2
5 150	2
6 160	3
7 170	3
8 180	4
9 190	4
10 200	5
11 210	5
12 220	6
13 230	6
14 240	7
15 250	7

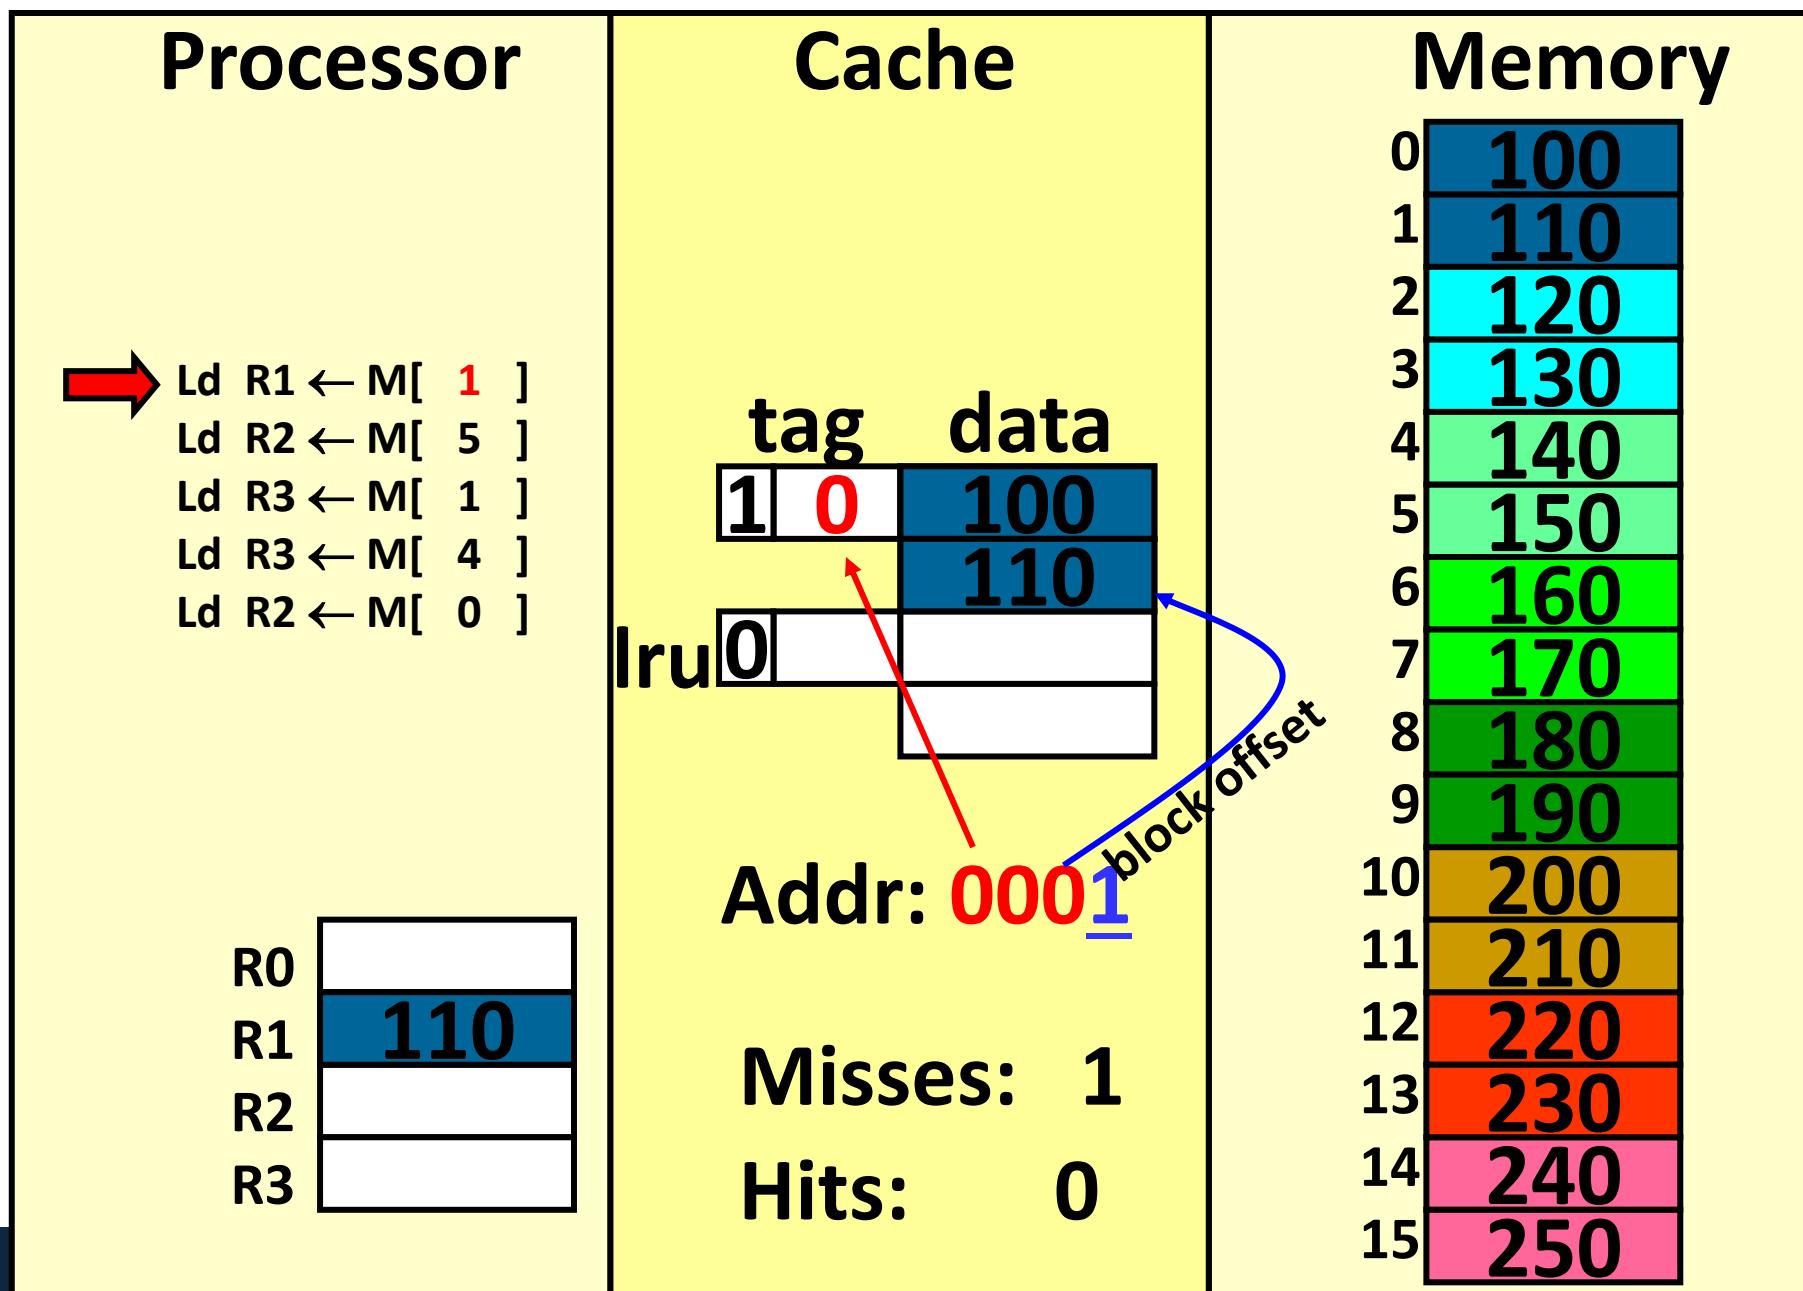
Block size for caches



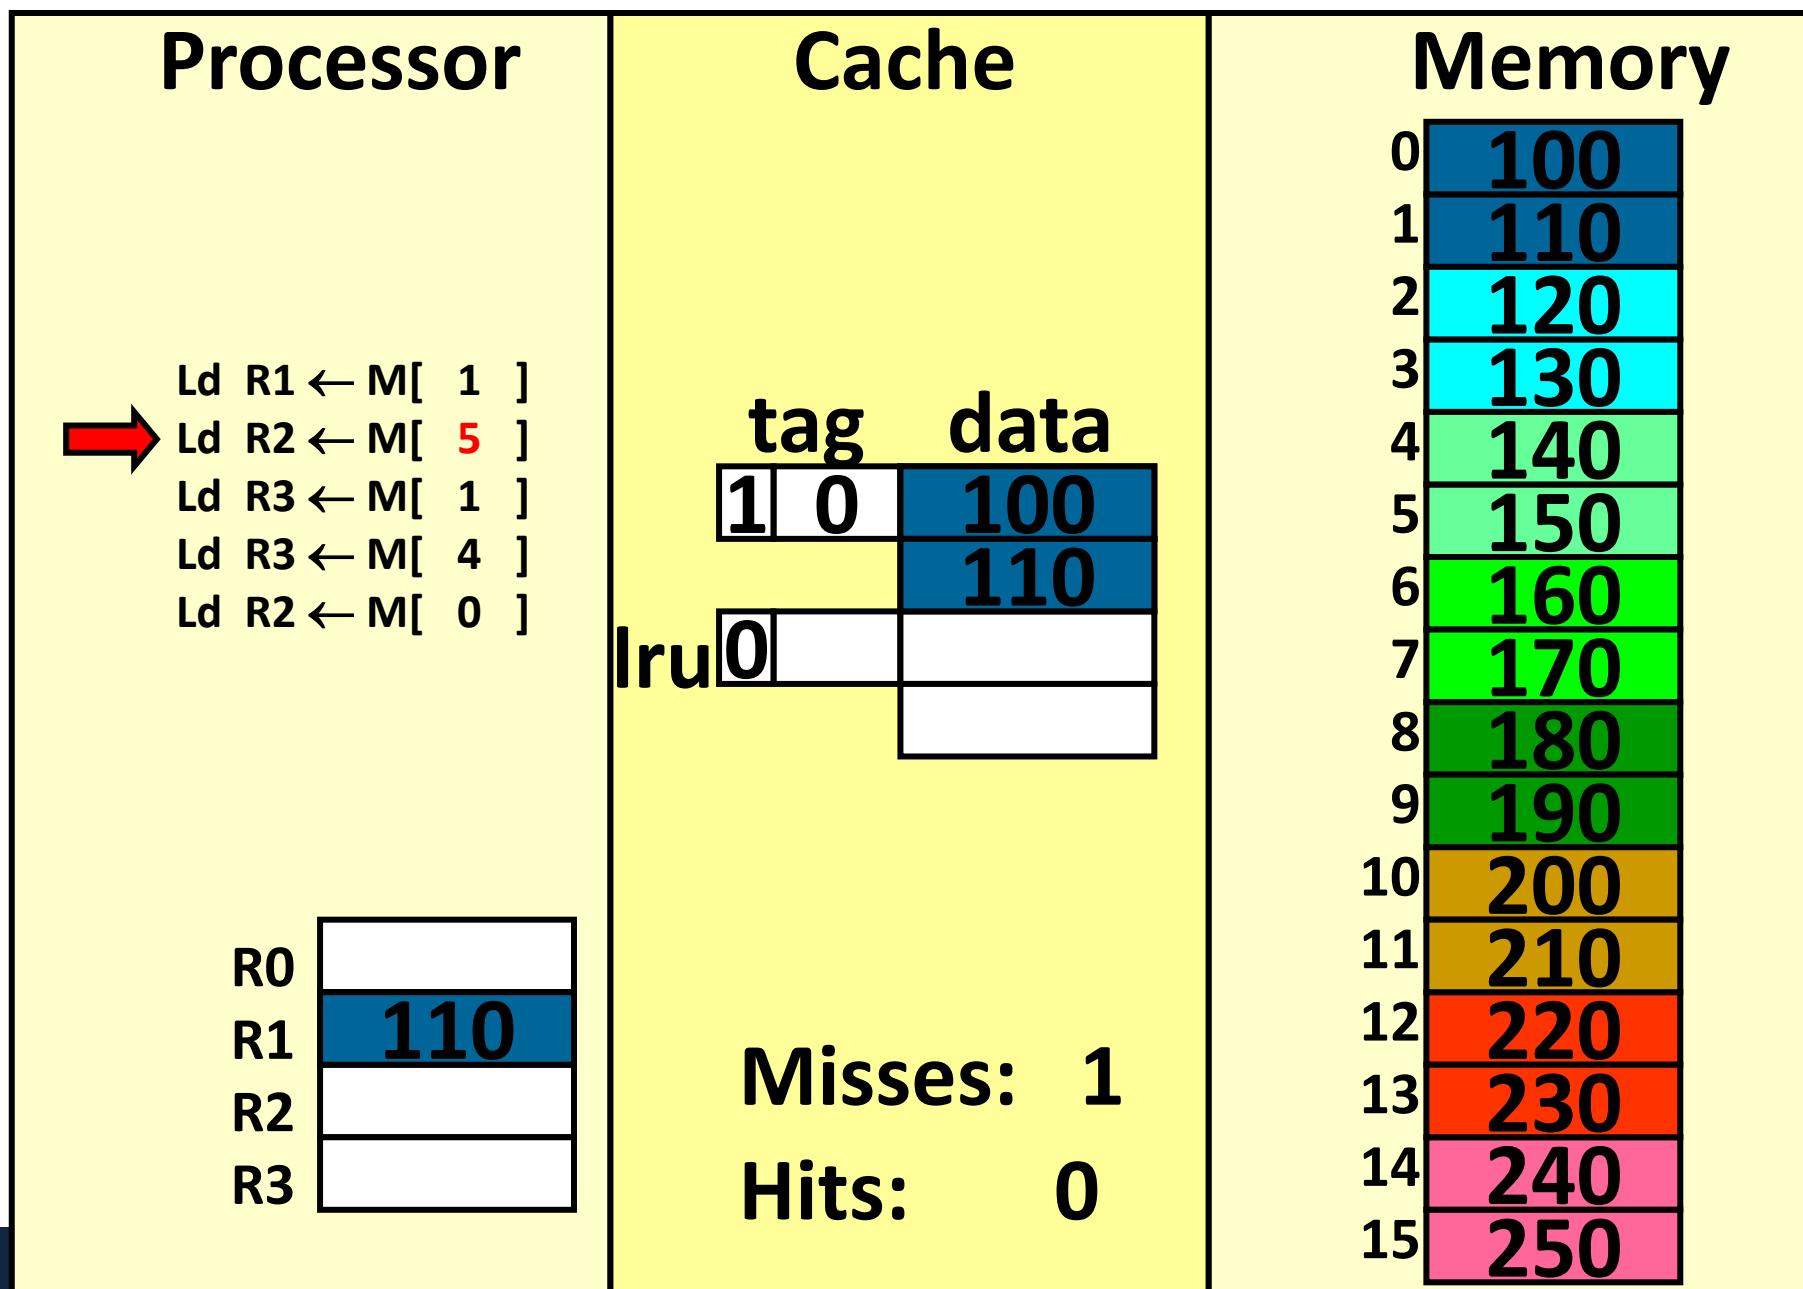
Block size for caches



Block size for caches

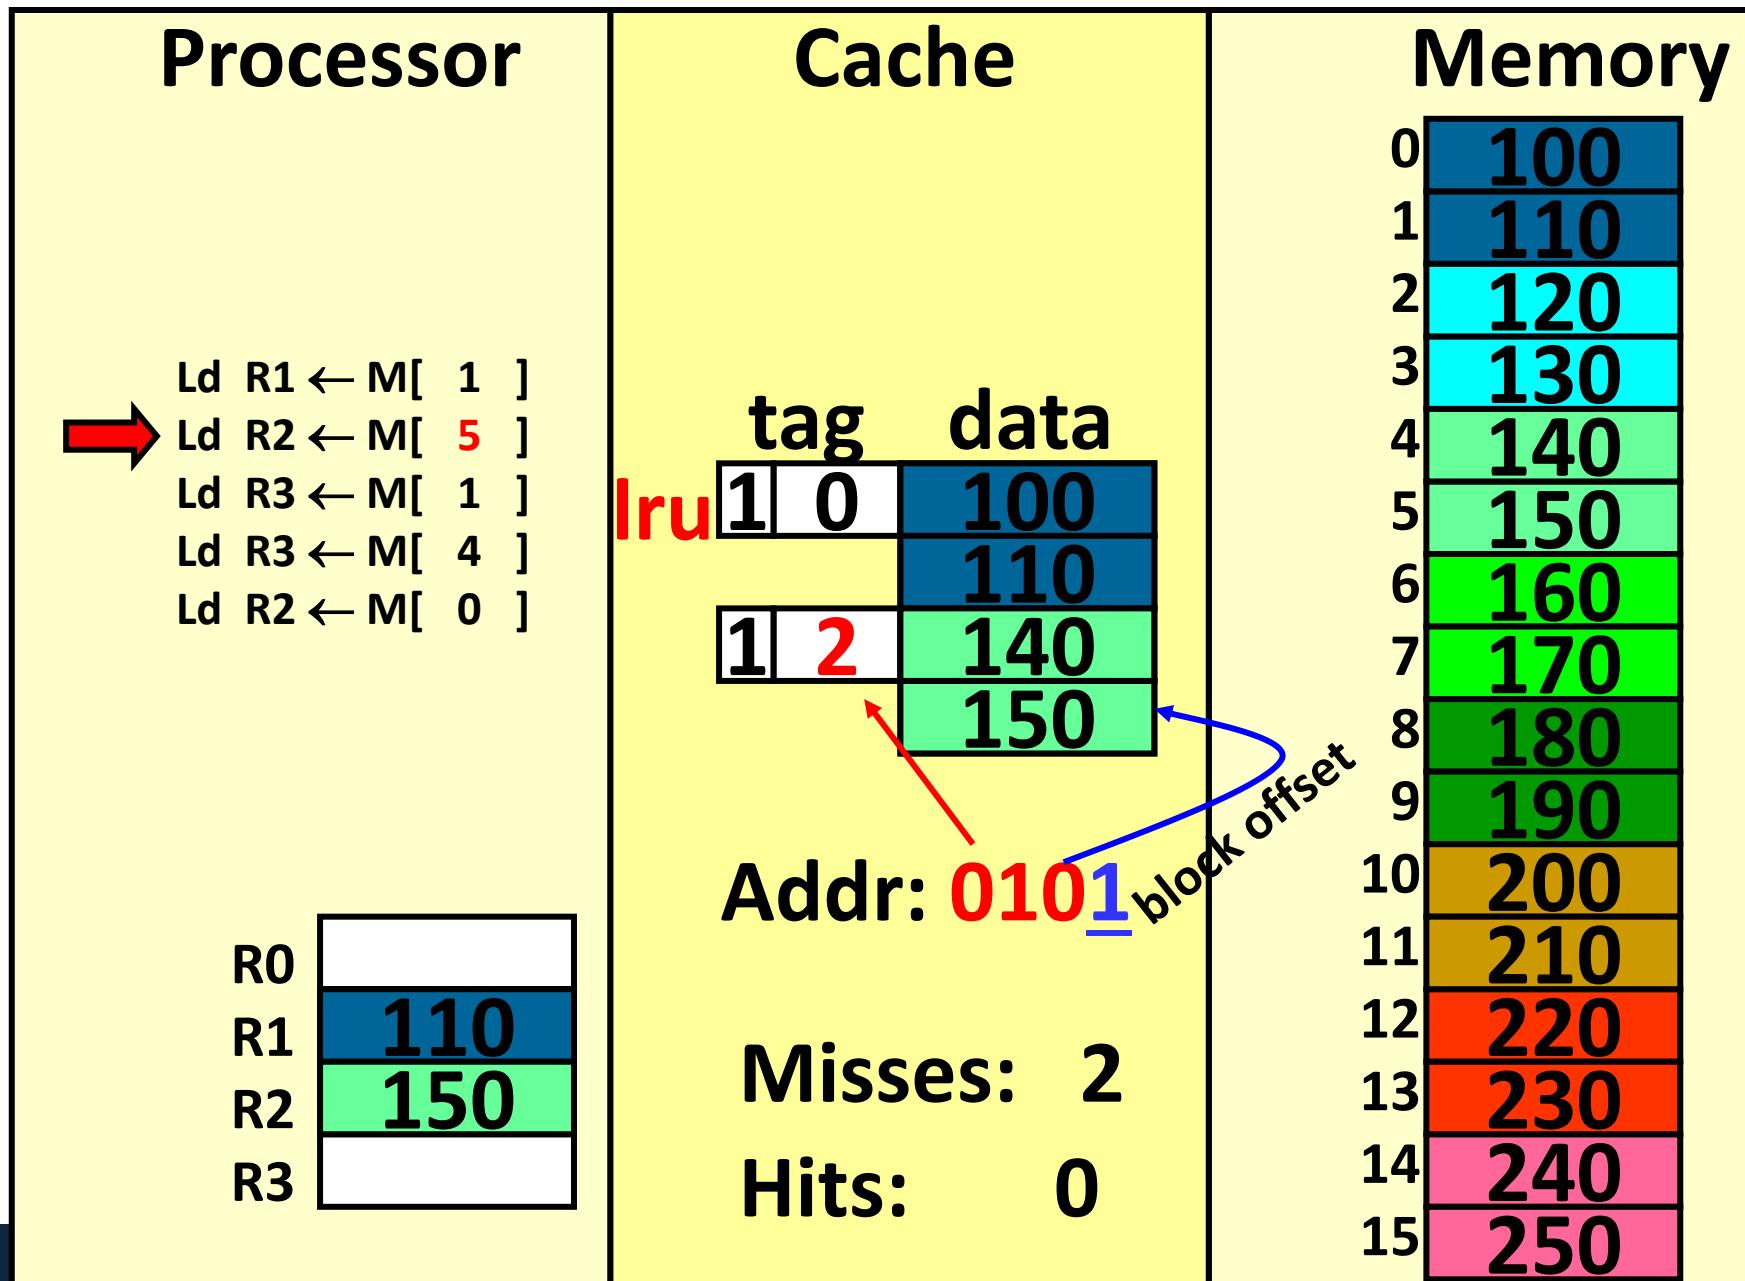


Block size for caches

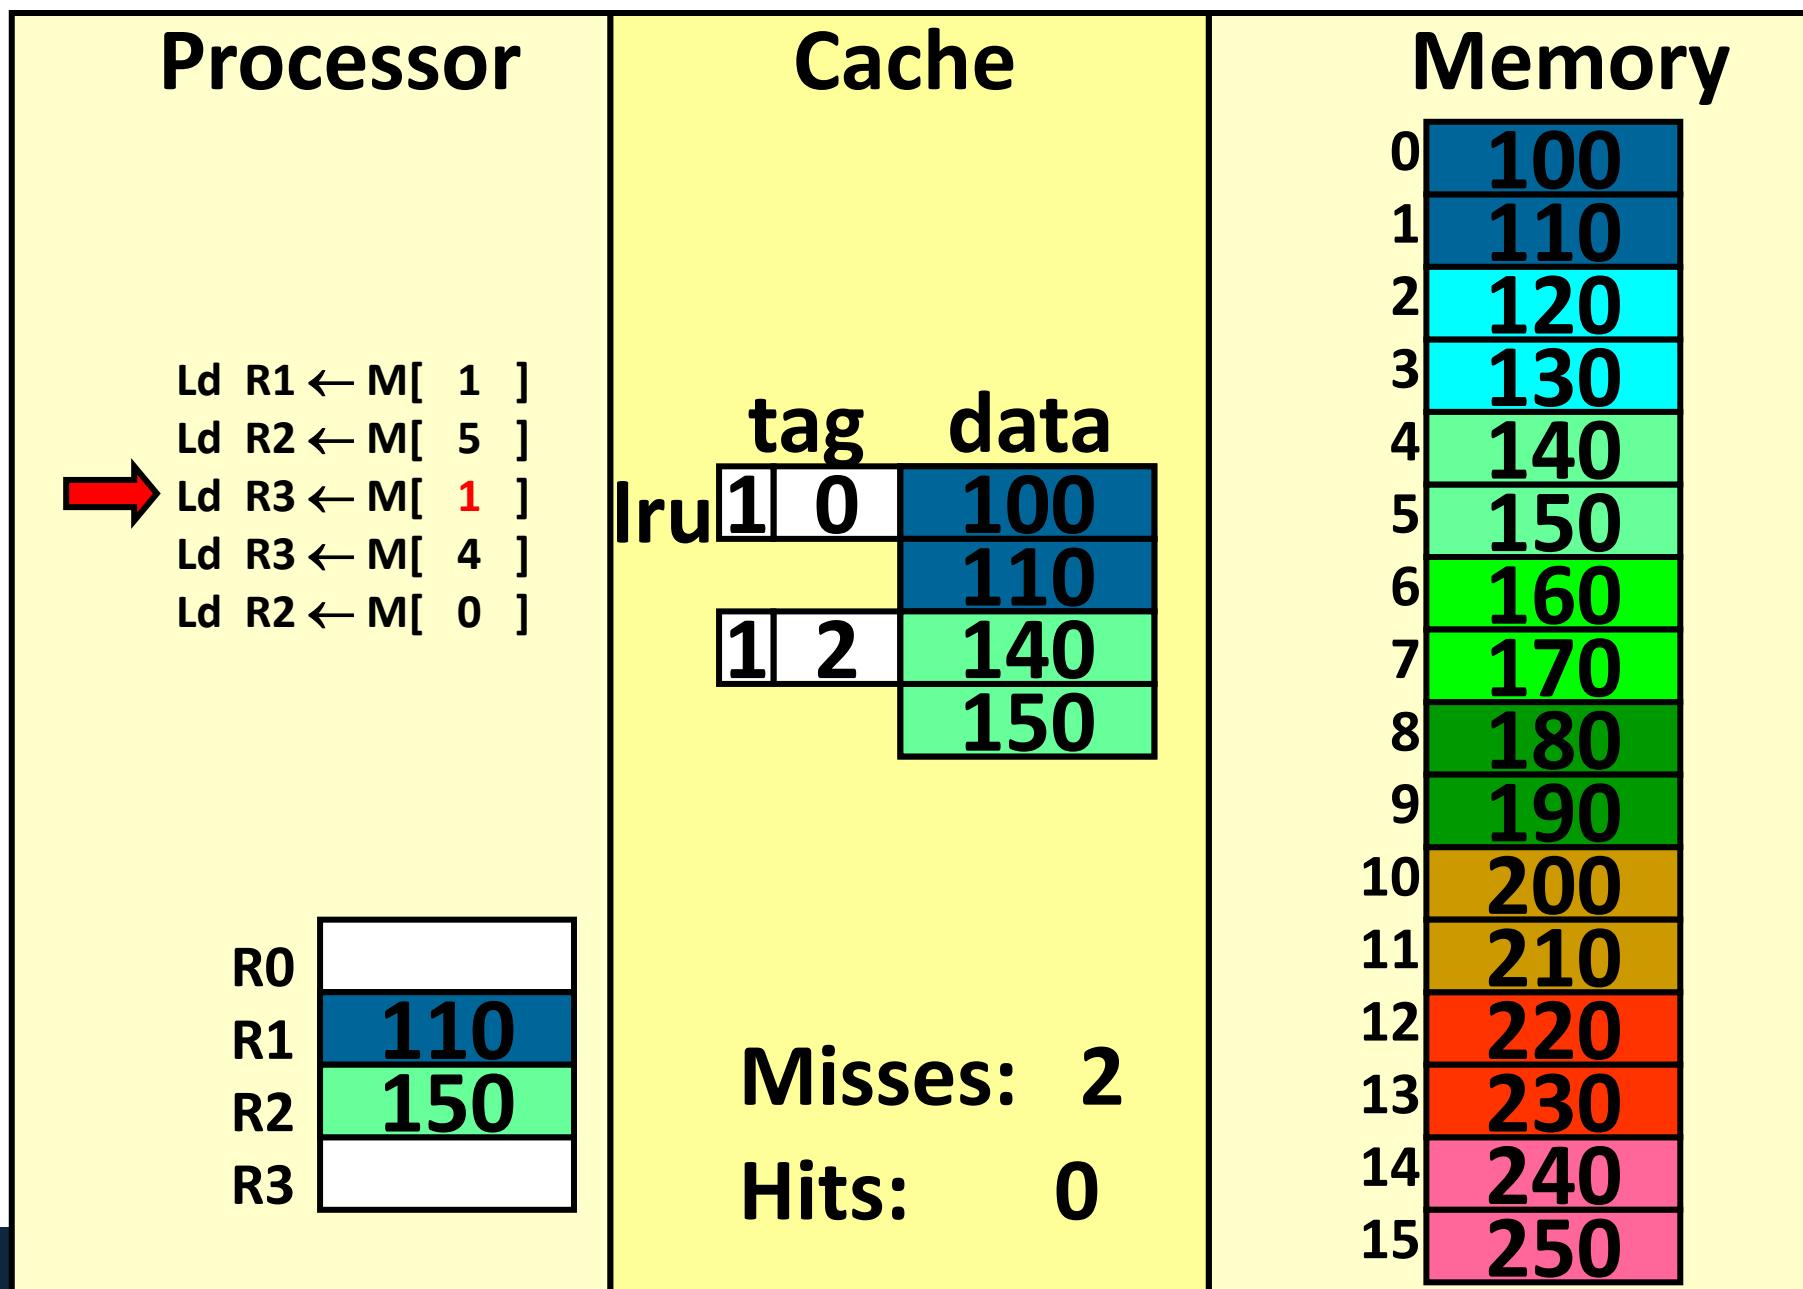


Block size for caches

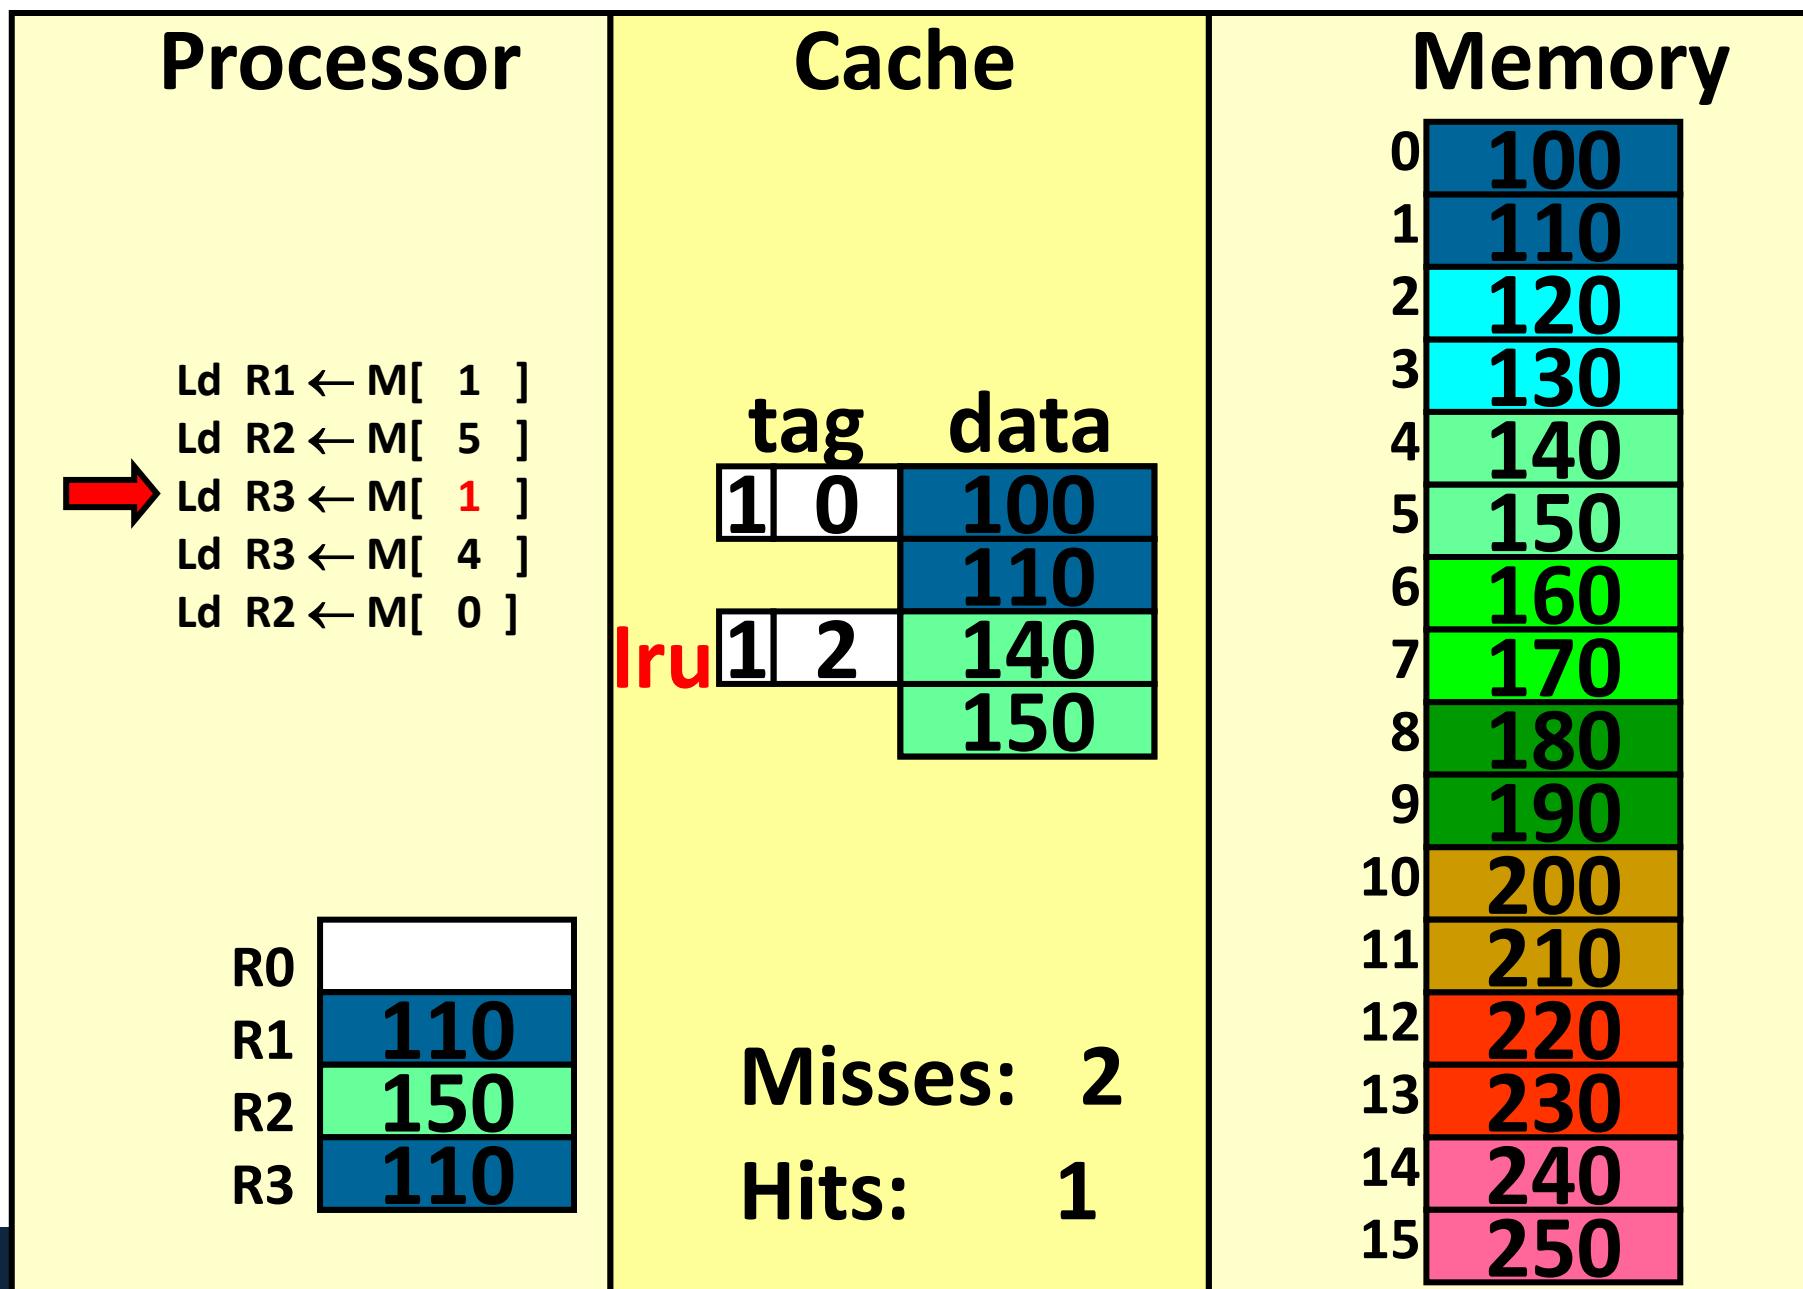
Poll: Complete the last 3 instructions yourself



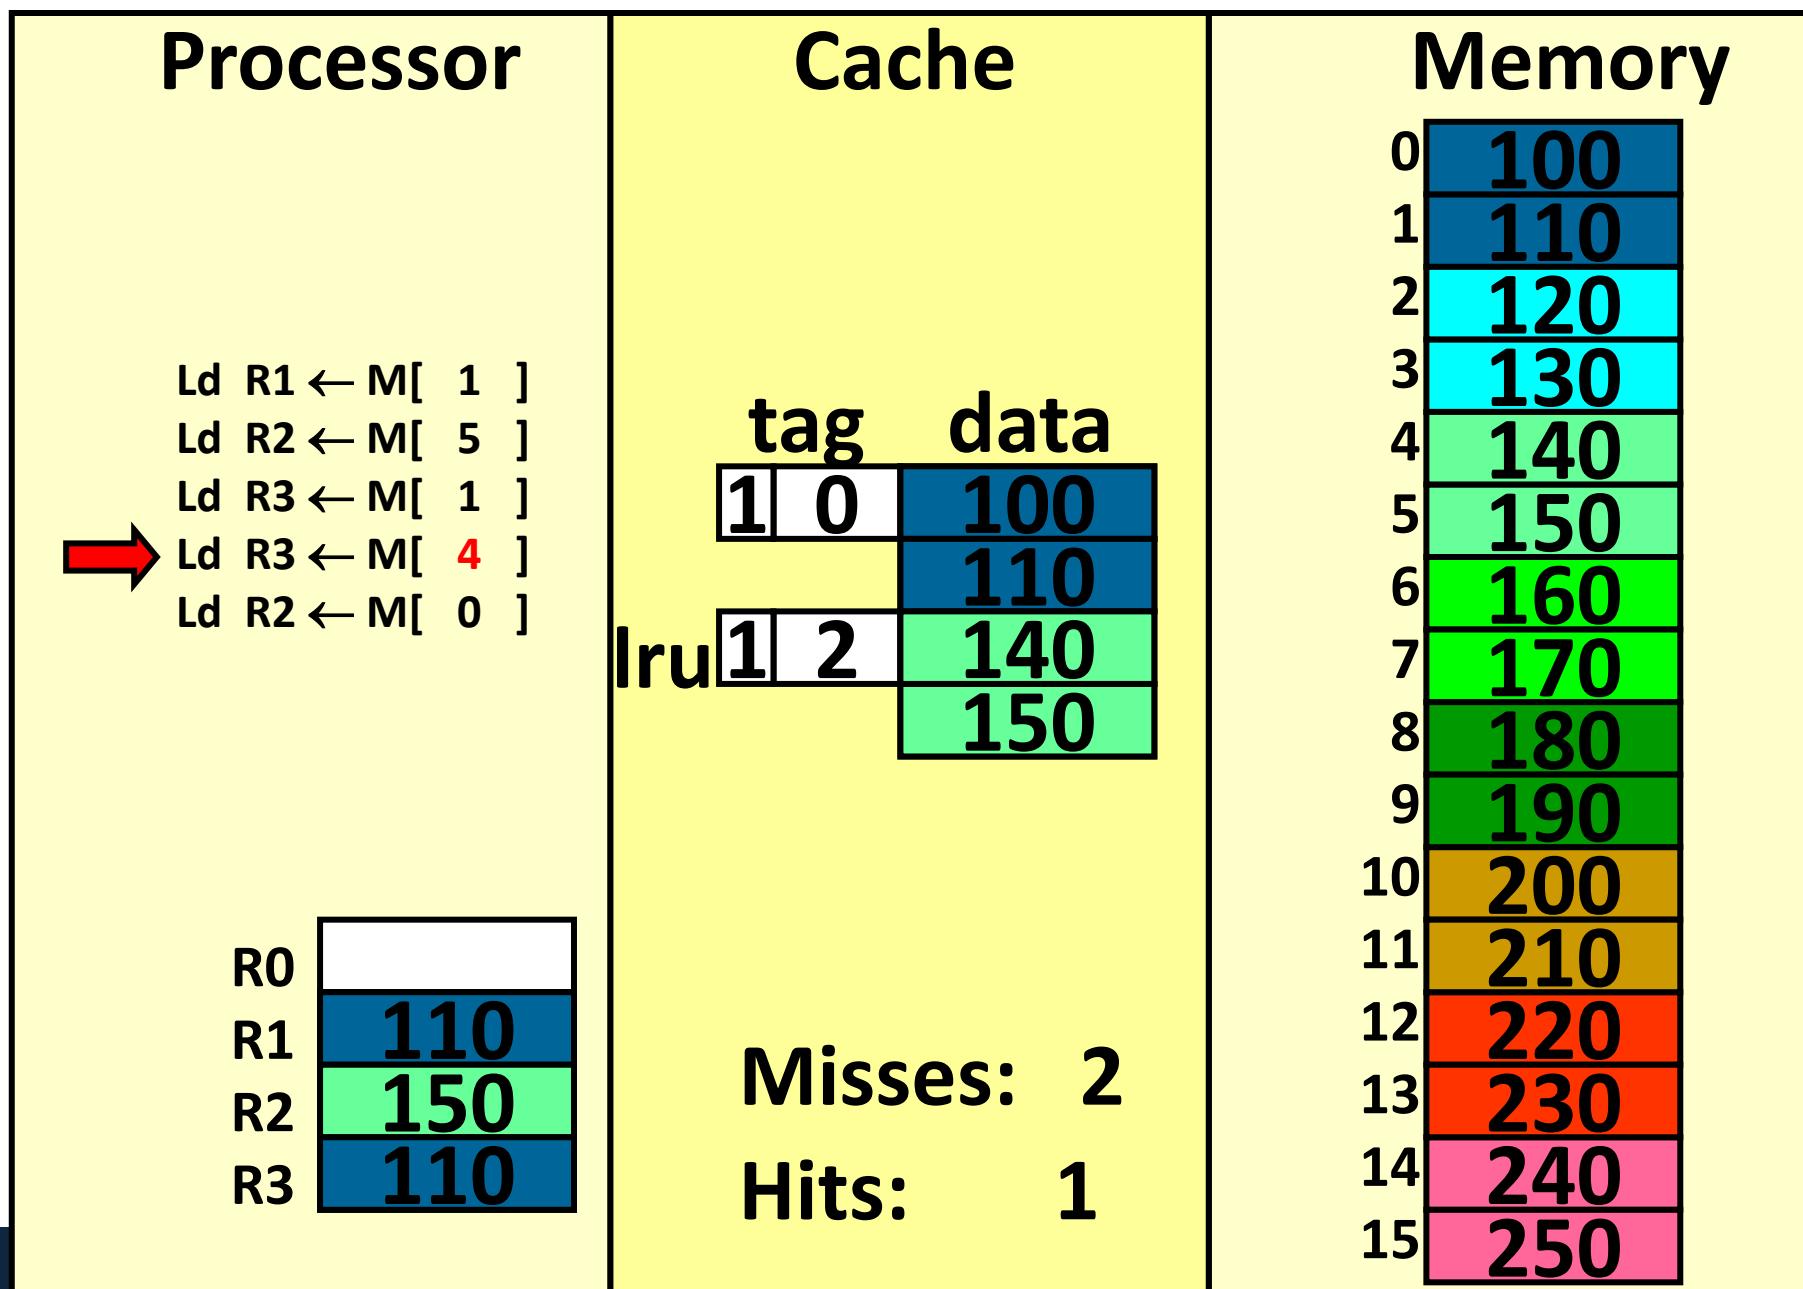
Block size for caches



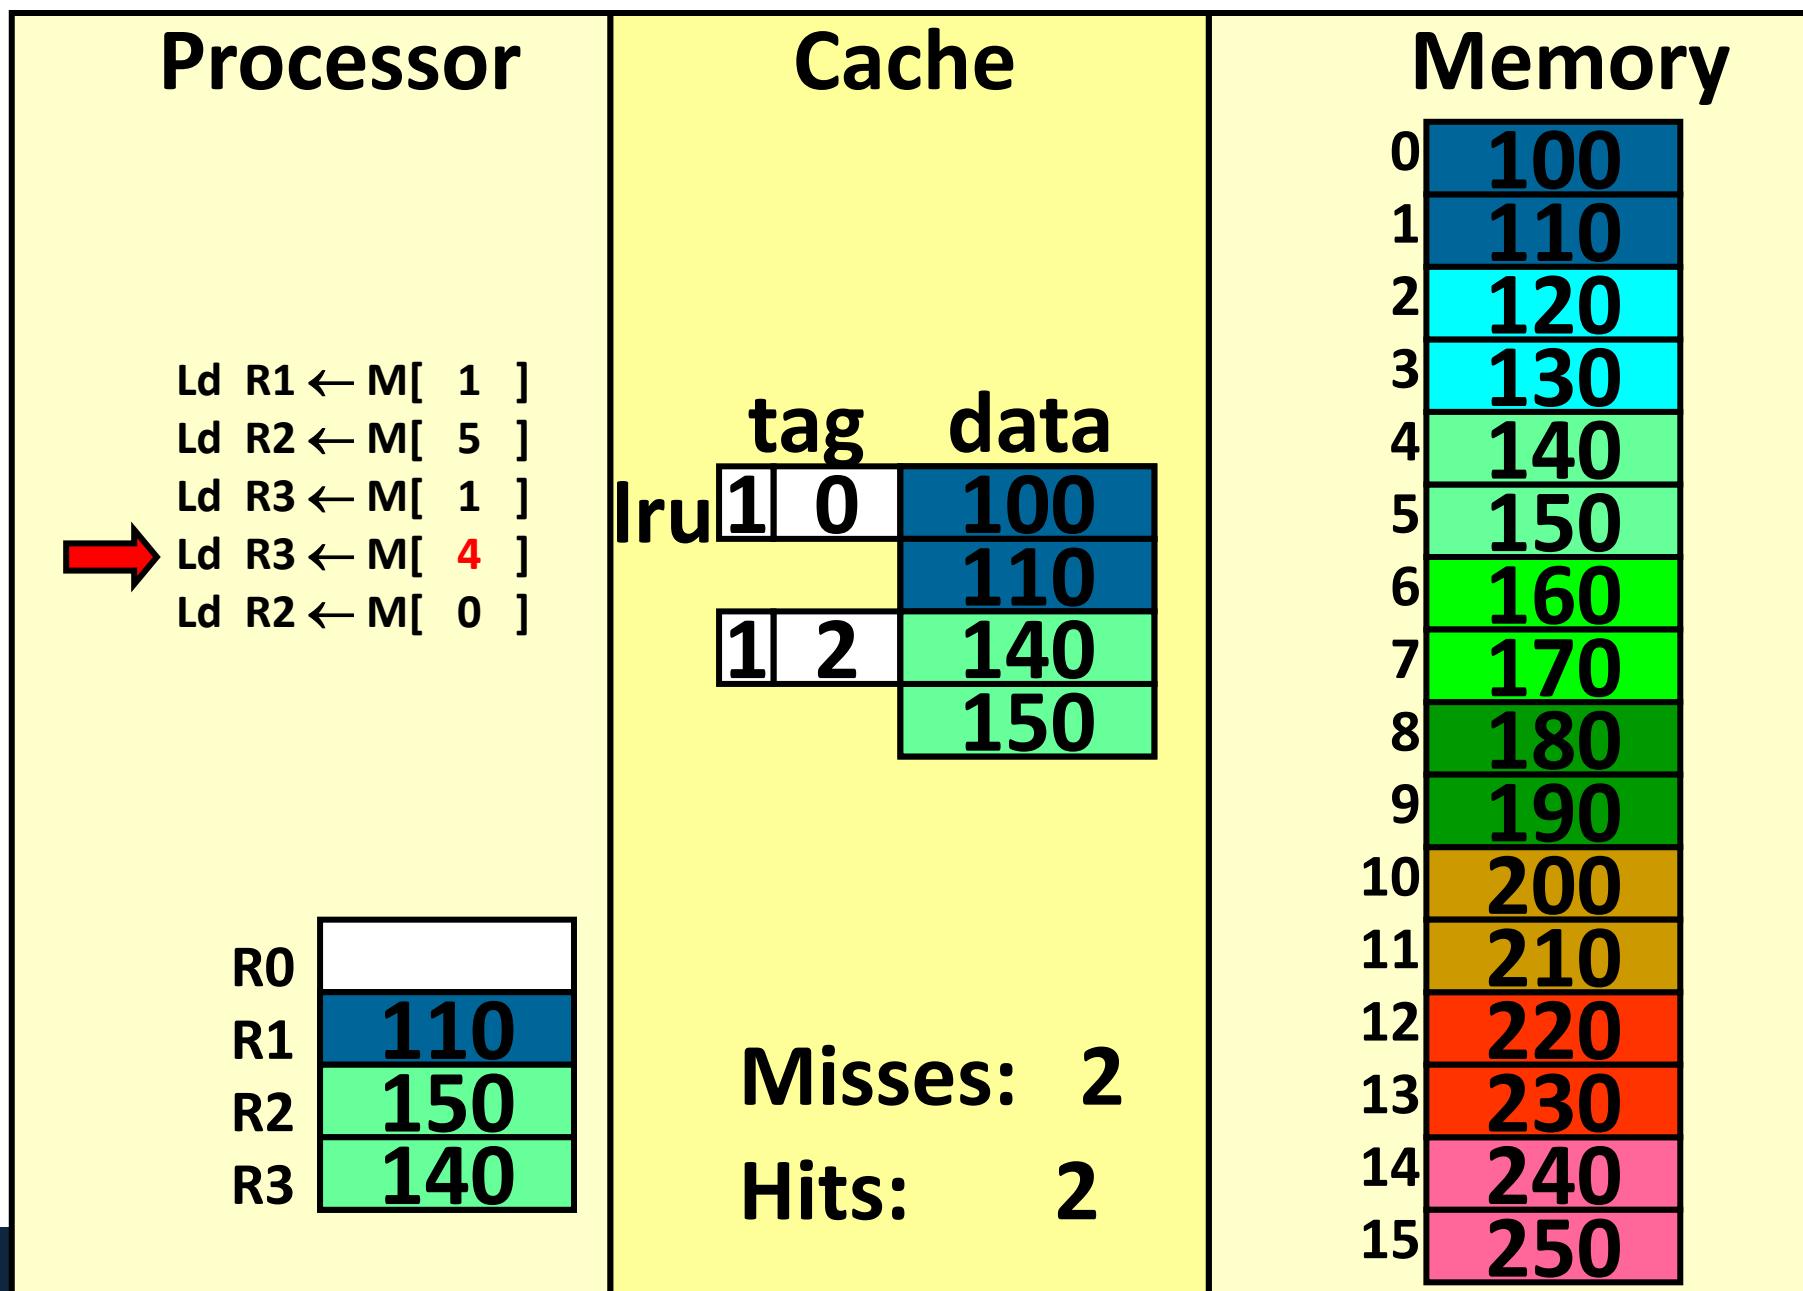
Block size for caches



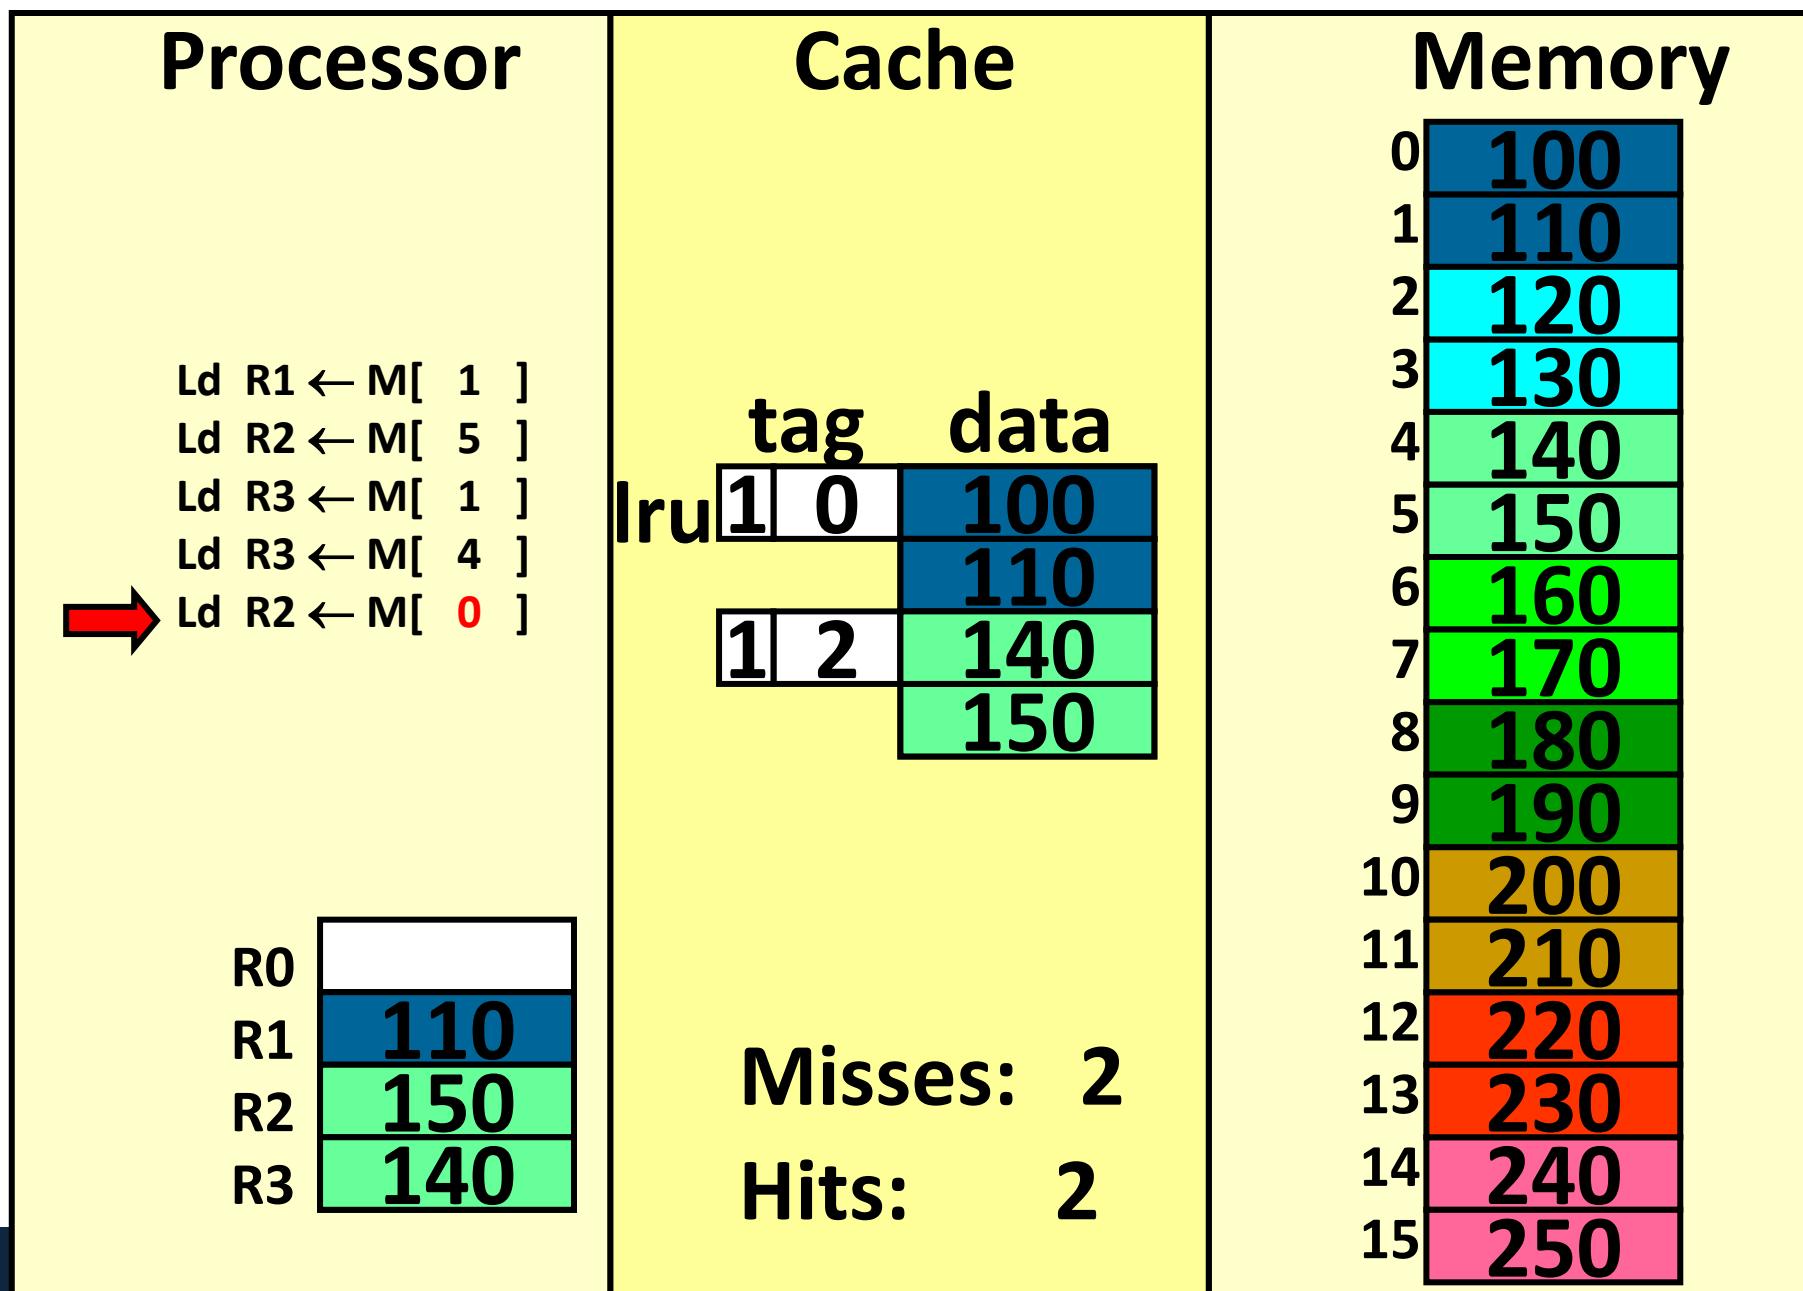
Block size for caches



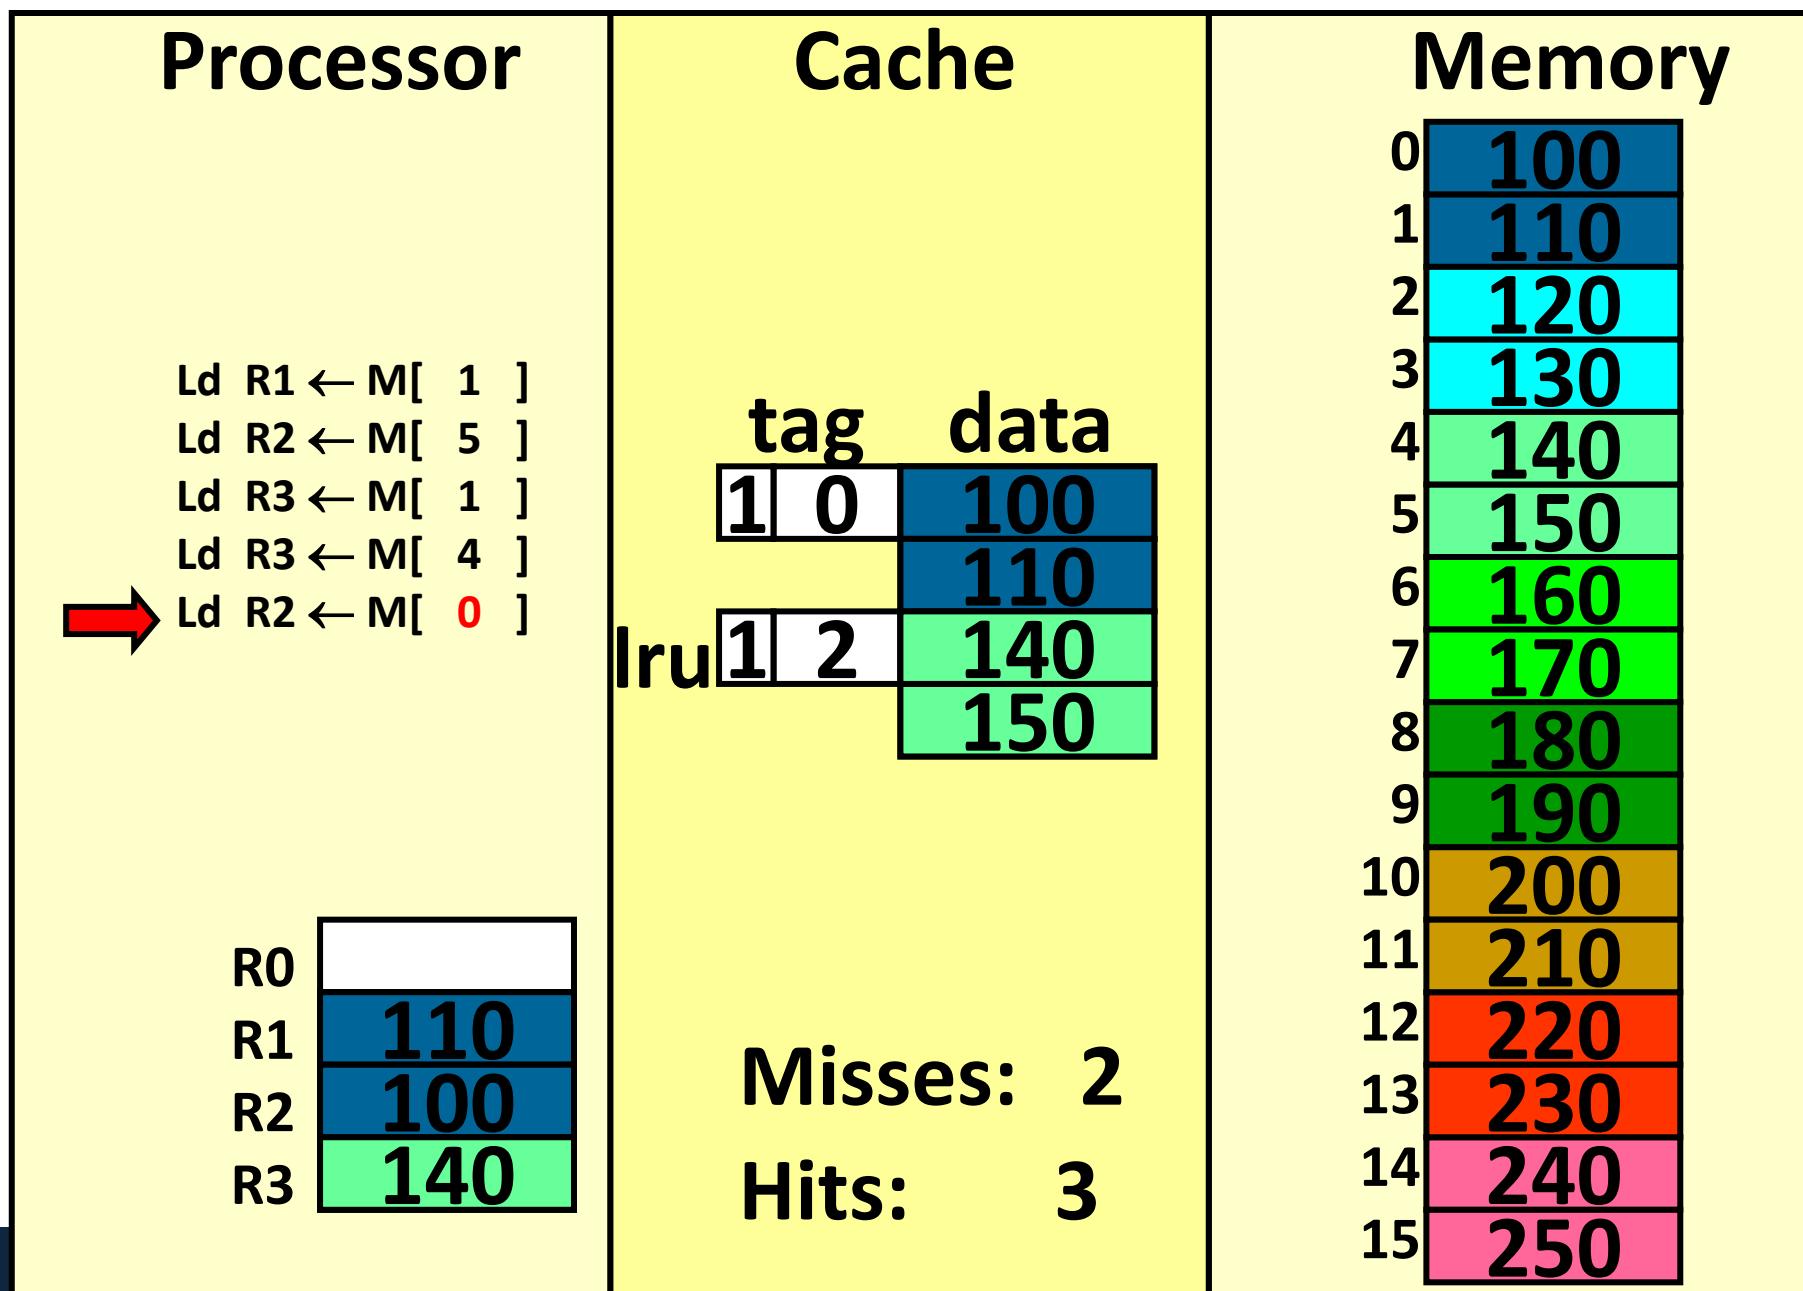
Block size for caches



Block size for caches



Block size for caches



Spatial Locality

- Notice that when we accessed address 1, we also brought in address 0
 - This turned out to be a good thing, since we later referenced address 0 and found it in the cache
- This is taking advantage of **spatial locality**:
 - If we access a memory location (e.g. 1000), we are more likely to access a location near it (e.g. 1001) than some random location
 - Arrays and structs are a big reason for this

```
for(i=0; i< N; i++)
    for(j = 0; j < N; j++ )
    {
        count++;
        arrayInt[i][j] = 10;
    }
```

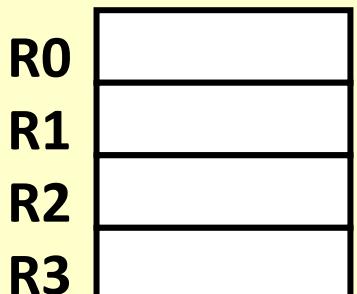
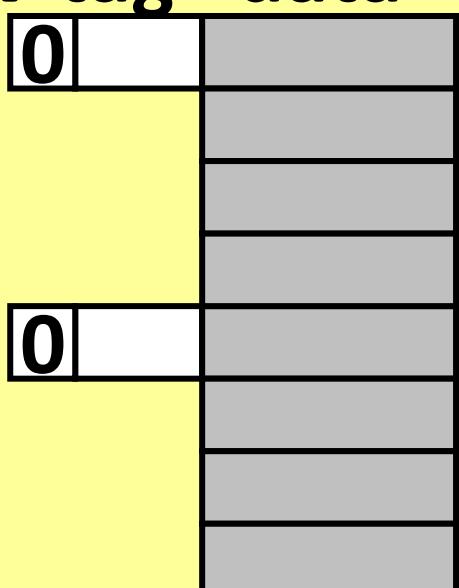


Agenda

- Larger Cache Blocks
- **Extra Problems**
- LRU with More than Two Blocks
- Write-Through Cache
- Write-Back Cache



Extra Practice Problem

Processor	Cache <u>2 cache lines</u> <u>2-bit tag field</u> <u>4-byte block</u>	Memory																																									
<pre>Ld R0 <- M[3] Ld R2 <- M[12] Ld R3 <- M[15] Ld R1 <- M[4] Ld R2 <- M[9]</pre> 	<p>V tag data</p>  <p>Cache contents:</p> <table border="1"><thead><tr><th>Line</th><th>Tag</th><th>Data</th></tr></thead><tbody><tr><td>0</td><td>00</td><td>120</td></tr><tr><td>1</td><td>01</td><td>123</td></tr></tbody></table>	Line	Tag	Data	0	00	120	1	01	123	<table border="1"><tbody><tr><td>0</td><td>78</td></tr><tr><td>1</td><td>29</td></tr><tr><td>2</td><td>120</td></tr><tr><td>3</td><td>123</td></tr><tr><td>4</td><td>71</td></tr><tr><td>5</td><td>150</td></tr><tr><td>6</td><td>162</td></tr><tr><td>7</td><td>173</td></tr><tr><td>8</td><td>18</td></tr><tr><td>9</td><td>21</td></tr><tr><td>10</td><td>33</td></tr><tr><td>11</td><td>28</td></tr><tr><td>12</td><td>19</td></tr><tr><td>13</td><td>200</td></tr><tr><td>14</td><td>210</td></tr><tr><td>15</td><td>225</td></tr></tbody></table>	0	78	1	29	2	120	3	123	4	71	5	150	6	162	7	173	8	18	9	21	10	33	11	28	12	19	13	200	14	210	15	225
Line	Tag	Data																																									
0	00	120																																									
1	01	123																																									
0	78																																										
1	29																																										
2	120																																										
3	123																																										
4	71																																										
5	150																																										
6	162																																										
7	173																																										
8	18																																										
9	21																																										
10	33																																										
11	28																																										
12	19																																										
13	200																																										
14	210																																										
15	225																																										

Solution to Practice Problem

Ld R0 \leftarrow M[3]

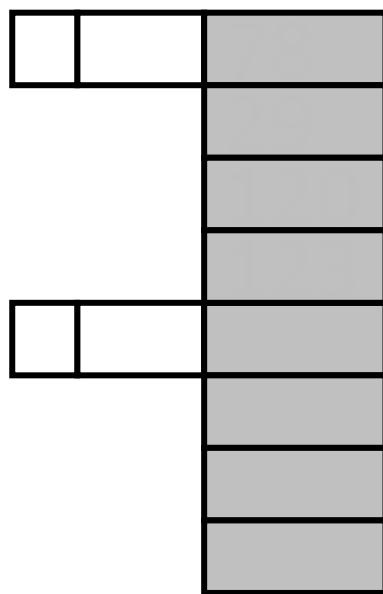
Ld R2 \leftarrow M[12]

Ld R3 \leftarrow M[15]

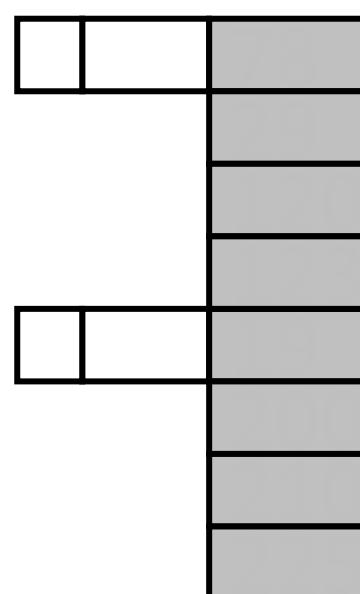
Ld R1 \leftarrow M[4]

Ld R2 \leftarrow M[9]

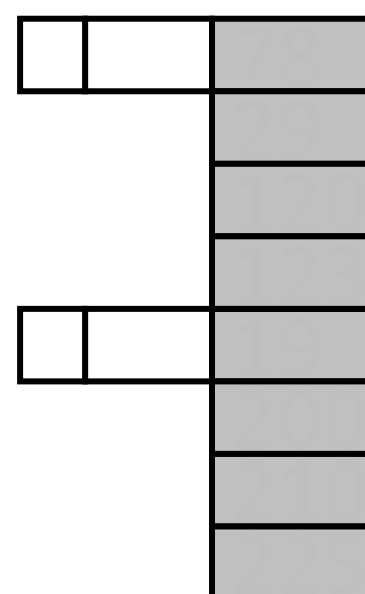
V tag data



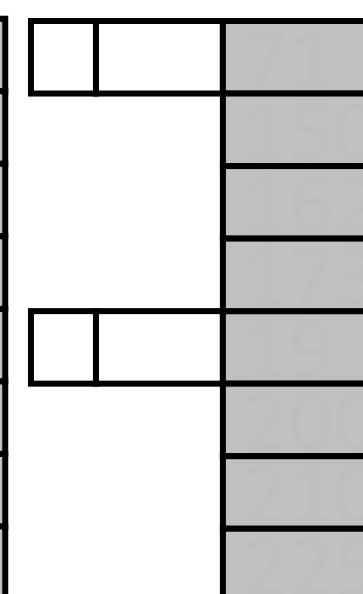
V tag data



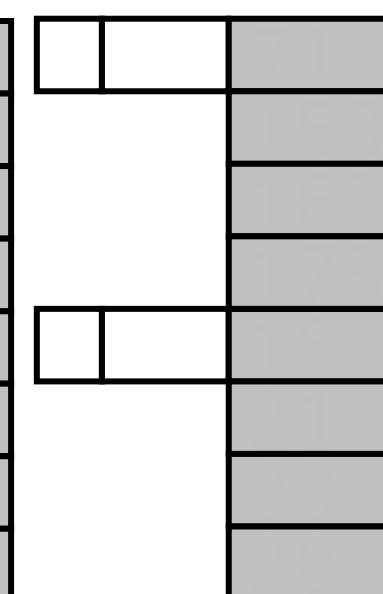
V tag data



V tag data



V tag data



Solution to Practice Problem

Ld R0 \leftarrow M[3]

V tag data

1	0	78
		29
		120
		123
0		
Iru		

miss

Ld R2 \leftarrow M[12]

V tag data

1	0	78
		29
		120
		123
1	3	19
		200
		210
		225

miss

Ld R3 \leftarrow M[15]

V tag data

1	0	78
		29
		120
		123
1	3	19
		200
		210
		225

hit

Ld R1 \leftarrow M[4]

V tag data

1	1	71
		150
		162
		173
1	3	19
		200
		210
		225

miss

Ld R2 \leftarrow M[9]

V tag data

1	1	71
		150
		162
		173
1	2	18
		21
		33
		28

miss

Extra Class Problem

*We'll see later that this is called a "fully-associative cache"

- Given a cache that works as we've described* with the following configuration: total size is 8 bytes, block size is 2 bytes, LRU replacement. The memory address size is 16 bits and is byte addressable.
 - How many bits are for each tag? How many blocks in the cache?
 - For the following reference stream, indicate whether each reference is a hit or miss: 0, 1, 3, 5, 12, 1, 2, 9, 4
 - What is the hit rate?
 - How many bits are needed for storage overhead for each block?

Agenda

- Larger Cache Blocks
- Extra Problems
- **LRU with More than Two Blocks**
- Write-Through Cache
- Write-Back Cache



LRU with more than 2 entries

- If we have more than 2 things we're keeping track of...
 - Can't just track LRU
 - Once we access that element, how do we know which of the other elements are LRU?
 - Must track the *full ordering* of when elements were accessed*
- Each element must store a number $[0-(N-1)] \rightarrow \log_2(N)$ bits
- 0 is LRU, 1 is 2nd LRU... N-1 is most recently used



LRU with more than 2 entries

- If we have more than 2 things we're keeping track of...
 - Can't just track LRU
 - Once we access that element, how do we know which of the other elements are LRU?
 - Must track the *full ordering* of when elements were accessed*
- Each element must store a number $[0-(N-1)] \rightarrow \log_2(N)$ bits
- 0 is LRU, 1 is 2nd LRU... N-1 is most recently used
- When element i is used:

```
X = counter[i]
counter[i] = N-1
for (j=0 to N-1)
    if ((j != i) AND (counter[j]>X)) counter[j]—
```

Initial State				
Element	0	1	2	3
Count	0	1	2	3

LRU with more than 2 entries

- If we have more than 2 things we're keeping track of...
 - Can't just track LRU
 - Once we access that element, how do we know which of the other elements are LRU?
 - Must track the *full ordering* of when elements were accessed*
- Each element must store a number $[0-(N-1)] \rightarrow \log_2(N)$ bits
- 0 is LRU, 1 is 2nd LRU... N-1 is most recently used
- When element i is used:

```
X = counter[i]
counter[i] = N-1
for (j=0 to N-1)
    if ((j != i) AND (counter[j]>X)) counter[j]—
```

Initial State				
Element	0	1	2	3
Count	0	1	2	3

Access Element 2				
Element	0	1	2	3
Count	0	1	3	2

LRU with more than 2 entries

- If we have more than 2 things we're keeping track of...
 - Can't just track LRU
 - Once we access that element, how do we know which of the other elements are LRU?
 - Must track the *full ordering* of when elements were accessed*
- Each element must store a number $[0-(N-1)] \rightarrow \log_2(N)$ bits
- 0 is LRU, 1 is 2nd LRU... N-1 is most recently used
- When element i is used:

```
X = counter[i]
counter[i] = N-1
for (j=0 to N-1)
    if ((j != i) AND (counter[j]>X)) counter[j]—
```
- Evict element with counter = 0 when needed
- Gets expensive for moderate to large N

Initial State				
Element	0	1	2	3
Count	0	1	2	3

Access Element 2				
Element	0	1	2	3
Count	0	1	3	2

Access Element 0				
Element	0	1	2	3
Count	3	0	2	1