


EECS 370

Classifying Cache Misses



Announcements


- Don't forget pre-lab quiz tonight!
- P3 due tomorrow
- P4 released in the next day
 - checkout cache simulator on website!
- HW 3 due Monday
- Bonus lecture Thursday



Cache Organization Comparison

Block size = 2 bytes, total cache size = 8 bytes for all caches


1. Fully associative (4-way associative)




2. Direct mapped




3. 2-way associative




Set-associative cache (REF 4)




Set-associative cache (REF 4)




Set-associative cache (REF 5)




Set-associative cache (REF 5)




Set-associative cache (REF 6)



Set-associative cache (REF 6)



Set-associative cache (REF 6)



Agenda

- Set-associativity overview
- Example
- **Class problem**
- Integrating caches into our processor



Class Problem 1

- For a 32-bit address and 16KB cache with 64-byte blocks, show the breakdown of the address for the following cache configuration:

A) fully associative cache

B) 4-way set associative cache

C) Direct-mapped cache



Class Problem 1

- For a 32-bit address and 16KB cache with 64-byte blocks, show the breakdown of the address for the following cache configuration:

A) fully associative cache

Block Offset = $\log_2(64)=6$ bits

Tag = $32 - 6 = 26$ bits

C) Direct-mapped cache

Block Offset = 6 bits

#lines = 256 Line Index = 8 bits

Tag = $32 - 6 - 8 = 18$ bits

B) 4-way set associative cache

Block Offset = 6 bits

#sets = #lines / ways = 64

Set Index = 6 bits

Tag = $32 - 6 - 6 = 20$ bits

Agenda

- Set-associativity overview
- Example
- Class problem
- **Integrating caches into our processor**



Multi-Level Caches

- We've been considering proc -> cache -> memory
- This works well if working data set is \leq size of cache
- But if data set is a little larger than cache, performance can plummet



Multi-Level Caches

- This is the motivation of multiple levels of caches
- L1 – smallest, fastest, closest to processor
- LN – biggest, slowest, closest to memory
- Allows for gradual performance degradation as data set size increases
- 3 levels of cache is pretty common in today's systems

What about cache for instructions

- We've been focusing on caching loads and stores (i.e. data)
- Instructions should be cached as well
- We have two choices:
 1. Treat instruction fetches as normal data and allocate cache lines when fetched
 2. Create a second cache (called the **instruction cache** or **ICache**) which caches instructions only
 - More common in practice

How do you know which cache to use?

What are advantages of a separate ICache?

Integrating Caches into Pipeline

- How are caches integrated into a pipelined implementation?
 - Replace instruction memory with Icache
 - Replace data memory with Dcache
- Issues:
 - Memory accesses now have variable latency
 - Both caches may miss at the same time



Agenda

- **Motivation**
- Example
- How to optimize cache design
- Practice Problem 1
- Practice Problem 2
- Practice Problem 3
- Practice Problem 4



Improving our Caches

- If our cache is getting a lot of misses, how do we improve it?
 - Depends on why the misses occurring
 - Is the cache too small? Is the associativity too restrictive? Something else?
- A decent first step is to **classify** the types of missing we are observing



Classifying Cache Misses

- Cache misses happen for 3* reasons
 - The 3C's of Cache misses:
- **Compulsory miss**
 - We've never accessed this data before
- **Capacity miss**
 - Cache is not large enough to hold all the data
 - May have been avoided if we used a bigger cache
- **Conflict miss**
 - Cache is large enough to hold data, but was replaced due to overly restrictive associativity
 - May have been avoided if we used a higher-associative cache

*On multi-core systems, there's a 4th C – take EECS 470/570 to learn more



Classifying Cache Misses

- Scenario: run given program on system with N-way cache of size M
 - Identify each miss
- We can classify each miss in a program by simulating on 3 different caches
 - If miss still occurs in cache where size \geq memory size: **compulsory miss**
 - Else, if miss occurs in fully associative cache of size M: **capacity miss**
 - Else, if miss occurs in N-way cache of size M (original cache): **conflict miss**

Agenda

- Motivation
- **Example**
- How to optimize cache design
- Practice Problem 1
- Practice Problem 2
- Practice Problem 3
- Practice Problem 4



3C's Sample Problem

Consider a cache with the following configuration: write-allocate, total size is 64 bytes, block size is 16 bytes, and 2-way associative. The memory address size is 16 bits and byte-addressable. The replacement policy is LRU. The cache is empty at the start.

For the following memory accesses, indicate whether the reference is a hit or miss, and the type of a miss (compulsory, conflict, capacity)



3 C's Practice Problem – 3 C's

64 bytes total, 16 byte
blocks, 2-way, 2 sets

Address	Infinite	FA	SA	3Cs
0x00				
0x14				
0x27				
0x08				
0x38				
0x4A				
0x18				
0x27				
0x0F				
0x40				



3 C's Practice Problem – 3 C's

Poll: How many blocks will be in a 64 byte FA cache?

Address	Infinite	FA	SA	3Cs
0x00	M			
0x14	M			
0x27	M			
0x08	H			
0x38	M			
0x4A	M			
0x18	H			
0x27	H			
0x0F	H			
0x40	H			



3 C's Practice Problem – 3 C's

64 bytes total, 16 byte
blocks, 2-way, 2 sets

Address	Infinite	FA	SA	3Cs
0x00	M	M		
0x14	M	M		
0x27	M	M		
0x08	H	H		
0x38	M	M		
0x4A	M	M		
0x18	H	M		
0x27	H	M		
0x0F	H	M		
0x40	H	H		



3 C's Practice Problem – 3 C's

64 bytes total, 16 byte blocks, 2-way, 2 sets

Address	Infinite	FA	SA	3Cs
0x00	M	M	M	Compulsory
0x14	M	M	M	Compulsory
0x27	M	M	M	Compulsory
0x08	H	H	H	---
0x38	M	M	M	Compulsory
0x4A	M	M	M	Compulsory
0x18	H	M	H	---
0x27	H	M	M	Capacity
0x0F	H	M	M	Capacity
0x40	H	H	M	Conflict



Agenda

- Motivation
- Example
- **How to optimize cache design**
- Practice Problem 1
- Practice Problem 2
- Practice Problem 3
- Practice Problem 4



How to reduce cache misses


- Compulsory miss
 - Reduce by **increasing cache block size**
 - Reduces total number of blocks for given cache size 😞
 - Or by using prefetching (guess we'll need data based on previous memory patterns - discussed more in EECS 470)
- Capacity miss
 - Reduce by **building a bigger cache**
 - Increase access latency 😞
- Conflict miss
 - Reduce by **increasing associativity**
 - Increase access latency / overheads 😞

Cache Performance

- How does changing these parameters affect performance?
 - Cache size
 - Block size
 - Associativity


Cache Size

- Cache size is the total data (not including tag) capacity
 - bigger can exploit temporal locality better
 - not ALWAYS better
- Too large a cache adversely affects hit & miss latency
 - smaller is faster => bigger is slower
 - access time may degrade critical path
- Too small a cache
 - doesn't exploit temporal locality well
 - useful data replaced often
- **Working set**: the whole set of data executing application references
 - **Within a time interval**




Block size

- Block size is the data that is associated with an address tag
 - Sub-blocking: A block divided into multiple pieces (each with V bit)
 - Can improve “write” performance
 - Take 470 to learn more
- Too small blocks
 - don’t exploit spatial locality well
 - have larger tag overhead
- Too large blocks
 - too few total # of blocks
 - likely-useless data transferred
 - Extra bandwidth/energy consumed



Associativity

- How many blocks can map to the same index (or set)?
- Larger associativity
 - lower miss rate, less variation among programs
 - diminishing returns
- Smaller associativity
 - lower cost
 - faster hit time
 - Especially important for L1 caches



Agenda

- Motivation
- Example
- How to optimize cache design
- **Practice Problem 1**
- Practice Problem 2
- Practice Problem 3
- Practice Problem 4



Poll: What is the CPI increase over base (1) due to hazards?

Practice Problem 1: CPI with caches

The *blaster* application run on the LC2k with full data forwarding and all branches predicted not-taken has the following instruction frequencies:

45% R-type 20% Branches 15% Loads 20% Stores

In *blaster*, 40% of branches are taken and 50% of LWs are followed by an immediate use.

The I-cache has a miss rate of 3% and the D-cache has a miss rate of 6% (no overlapping of misses). On a miss, the main memory is accessed and has a latency of 100 ns. The clock frequency is 500 MHz.



Problem 1 Solution

The *blaster* application run on the LC2k with full data forwarding and all branches predicted not-taken has the following instruction frequencies:

45% R-type 20% Branches 15% Loads 20% Stores

In *blaster*, 40% of branches are taken and 50% of LWs are followed by an immediate use.

The I-cache has a miss rate of 3% and the D-cache has a miss rate of 6% (no overlapping of misses). On a miss, the main memory is accessed and has a latency of 100 ns. The clock frequency is 500 MHz.

What is the CPI of *blaster* on the LC2k?

Stalls per cache miss = $100 \text{ ns} / 2\text{ns} = 50 \text{ cycles}$ (500 Mhz \rightarrow 2ns cycle time)

CPI = 1 + data hazard stalls + control hazard stalls + icache stalls + dcache stalls

CPI = 1 + $0.15 * 0.50 * 1$ + $0.20 * 0.40 * 3$ + $1 * 0.03 * 50$ + $0.35 * 0.06 * 50$

Agenda

- Motivation
- Example
- How to optimize cache design
- Practice Problem 1
- **Practice Problem 2**
- Practice Problem 3
- Practice Problem 4



Poll: How many bytes are read/written if we have no cache?

Practice Problem 2: Memory Usage

- Say you have the following:
 - A program that generates 2 Billion loads and 1 Billion stores, each 4 bytes in size.
 - A cache with a 32-byte block which gets a 95% hit rate on that program.
- How many bytes of memory would be read and written if:
 - We had no cache?
 - We had a write-through cache with a no-write allocate policy?
 - We had a write-back cache with a write-allocate policy? (Assume 25% of all misses result in a dirty eviction)



Practice Problem 2: Memory Usage

- Say you have the following:
 - A program that generates 2 Billion loads and 1 Billion stores, each 4 bytes in size.
 - A cache with a 32-byte block which gets a 95% hit rate on that program.
- Let's start with the no-cache case.
 - All stores go to memory and are 4 bytes each
 - Writes: 1 billion stores * 4 bytes = 4 billion bytes
 - All loads go to memory and are 4 bytes each.
 - Reads: 2 billion loads * 4 bytes = 8 billion bytes
 - Write-through with no write-allocate?

Practice Problem 2: Memory Usage

- Say you have the following:
 - A program that generates 2 Billion loads and 1 Billion stores, each 4 bytes in size.
 - A cache with a 32-byte block which gets a 95% hit rate on that program.
- Write-through, no allocate.
 - All stores still go to memory and are still 4 bytes each.
 - Writes: $1 \text{ billion stores} * 4 \text{ bytes} = 4 \text{ billion bytes}$
 - Only loads that miss in the cache go to memory. But they read the full cache block.
 - Reads: $2 \text{ billion loads} * 0.05 * 32 \text{ bytes} = 3.2 \text{ billion bytes}$

Practice Problem 2: Memory Usage

- Say you have the following:
 - A program that generates 2 Billion loads and 1 Billion stores, each 4 bytes in size.
 - A cache with a 32-byte block which gets a 95% hit rate on that program.
- Write-through, no allocate.



Practice Problem 2: Memory Usage

- Say you have the following:
 - A program that generates 2 Billion loads and 1 Billion stores, each 4 bytes in size.
 - A cache with a 32-byte block which gets a 95% hit rate on that program.
- Write-back, write-allocate (25% of all misses result in a dirty eviction)

Practice Problem 2: Memory Usage

- Say you have the following:
 - A program that generates 2 Billion loads and 1 Billion stores, each 4 bytes in size.
 - A cache with a 32-byte block which gets a 95% hit rate on that program.
- Write-back, write-allocate (25% of all misses result in a dirty eviction)
 - *Store* misses result in a cache block being read.
 - Reads: $1 \text{ billion stores} * 0.05 * 32 \text{ bytes} = 1.6 \text{ billion bytes}$
 - *Load* misses result in a cache block being read.
 - Reads: $2 \text{ billion loads} * 0.05 * 32 \text{ bytes} = 3.2 \text{ billion bytes}$
 - So that is 4.8 billion bytes of data read.

Practice Problem 2: Memory Usage

- Say you have the following:
 - A program that generates 2 Billion loads and 1 Billion stores, each 4 bytes in size.
 - A cache with a 32-byte block which gets a 95% hit rate on that program.
- Write-back, write-allocate (25% of all misses result in a dirty eviction)
 - *Store* misses result in dirty eviction 1/4 of the time.
 - Reads: $1 \text{ billion stores} * 0.05 * 32 \text{ bytes} * (.25) = 0.4 \text{ billion bytes}$
 - *Load* misses result in a cache block being read.
 - Reads: $2 \text{ billion loads} * 0.05 * 32 \text{ bytes} * (.25) = 0.8 \text{ billion bytes}$

Agenda

- Motivation
- Example
- How to optimize cache design
- Practice Problem 1
- Practice Problem 2
- **Practice Problem 3**
- Practice Problem 4



Practice Problem 3: CPI w/ Caches 2

- Given a 200 MHz processor with 8KB instruction and data caches and a memory access latency of 20 cycles. Both caches are 2-way associative. A program running on this processor has a 95% icache hit rate and a 90% dcache hit rate. On average, 30% of the instructions are loads or stores. The CPI of this system, if caches were ideal would be 1.
- Suppose you have 2 options for the next generation processor, which do you pick?
 - Option 1:** Double the clock frequency—assume this will increase your memory latency to 40 cycles. Also assume a base CPI of 1 can still be achieved after this change.
 - Option 2:** Double the size of your caches, this will increase the instruction cache hit rate to 98% and the data cache hit rate to 95%. Assume the hit latency is still 1 cycle.

Practice Problem 3: Solution

Option 1: (double clock freq, base cycle time is 5 ns, so new cycle time is 2.5 ns)

$$\text{CPI} = \text{baseCPI} + \text{IcacheStallCPI} + \text{DcacheStallCPI}$$

$$\text{CPI} = 1.0 + 0.05 * 40 + 0.3 * 0.1 * 40 = 4.2$$

Execution time = 4.2 * Ninstrs * 2.5ns = 10.5ns * Ninstrs

Option 2 (icache/dcache miss rates lowered to 2% and 5%)

$$\text{CPI} = \text{baseCPI} + \text{IcacheStallCPI} + \text{DcacheStallCPI}$$

$$\text{CPI} = 1.0 + 0.02 * 20 + 0.3 * 0.05 * 20 = 1.7$$

Execution time = 1.7 * Ninstrs * 5ns = 8.5ns * Ninstrs

Therefore, Option 2 is the better choice

Agenda

- Motivation
- Example
- How to optimize cache design
- Practice Problem 1
- Practice Problem 2
- Practice Problem 3
- **Practice Problem 4**



Practice Problem 4: Guess that cache!

Here is the series of address references (in hex) to a cache of size 512 bytes. You are asked to **determine the configuration of the cache**. Assume 12-bit addresses

0x310 - Miss

0x30f - Miss

0x510 - Miss

0x31f - Hit

0x72d - Miss

0x72f - Hit

0x320 - Miss

0x520 - Miss

0x720 - Miss

Block size: ?

Associativity: ?

Number of sets: ?



Practice Problem 4: Guess that cache!

Similar to homework! Here is the series of address references (in hex) to a cache of size 512 bytes. You are asked to **determine the configuration of the cache**. Assume 12-bit addresses

0x310 - Miss

0x30f - Miss

0x510 - Miss

0x31f - Hit

0x72d - Miss

0x72f - Hit

0x320 - Miss

0x520 - Miss

0x720 - Miss



Practice Problem 4: Guess that cache!

Similar to homework! Here is the series of address references (in hex) to a cache of size 512 bytes. You are asked to **determine the configuration of the cache**. Assume 12-bit addresses

0x310 - Miss

0x30f - Miss

0x510 - Miss

0x31f - Hit

0x72d - Miss

0x72f - Hit

0x320 - Miss

0x520 - Miss

0x720 - Miss



Practice Problem 4: Guess that cache!

Here is the series of address references (in hex) to a cache of size 512 bytes. You are asked to **determine the configuration of the cache**. Assume 12-bit addresses

0x310 – Miss



Determine block size

0x30f – Miss



First hit must be brought in by another miss

0x510 – Miss



Take closest address: 0x310, so know block size must be at least 16 bytes so 0x31f brought in when 0x310 miss occurs

0x31f – Hit

0x72d – Miss

0x72f – Hit

0x320 – Miss

0x520 – Miss

0x720 - Miss

Now, is the block size larger? Know that 0x30f was a miss, thus 0x310 and 0x30f not in the same block. Thus, block size must be ≤ 16 bytes

Thus Block Size = 16 bytes

Practice Problem 4: Guess that cache!

Here is the series of address references (in hex) to a cache of size 512 bytes. You are asked to **determine the configuration of the cache**. Assume 12-bit addresses

Determine associativity

- 0x310 – Miss 
- 0x30f – Miss
- 0x510 – Miss 
- 0x31f – Hit 
- 0x72d – Miss
- 0x72f – Hit 
- 0x320 – Miss 
- 0x520 – Miss 
- 0x720 - Miss 

Assume direct mapped: 3-bit tag, 5-bit index, 4-bit offset.

If DM, 0x310 and 0x510 would both map to index 17,
Thus 0x31f could not be a hit. So, not direct mapped.

Assume 2-way associative: 4-bit tag, 4-bit index, 4-bit offset
This fixes the green accesses, and allows 0x31f to be a hit.

What about > 2-way associative?

Now we also know that 0x720 is a miss even though 3 accesses earlier 0x72f was a hit, and thus it is in the cache. The intervening 2 accesses must kick it out, 0x320 and 0x520. Both go to set 2. If the associativity was > 2, then 0x720 would be a hit. So, must conclude that cache is 2-way associative.

Lastly, number of sets = $512 / (2 * 16) = 16$

Next time

- Completing the hierarchy: virtual memory